



Booth A



Booth B



Booth A

242

104 (40 %)



Booth B

Walked by

260

Stopped to taste

145 (60 %)



Booth A

242

104 (40 %)

31 (30%)



Booth B

260

145 (60 %)

4 (3 %)

Walked by

Stopped to taste

Bought jam



Recommendation Systems

An Overview

Marie Roald

PhD Defence Trial lecture
2023-09-04



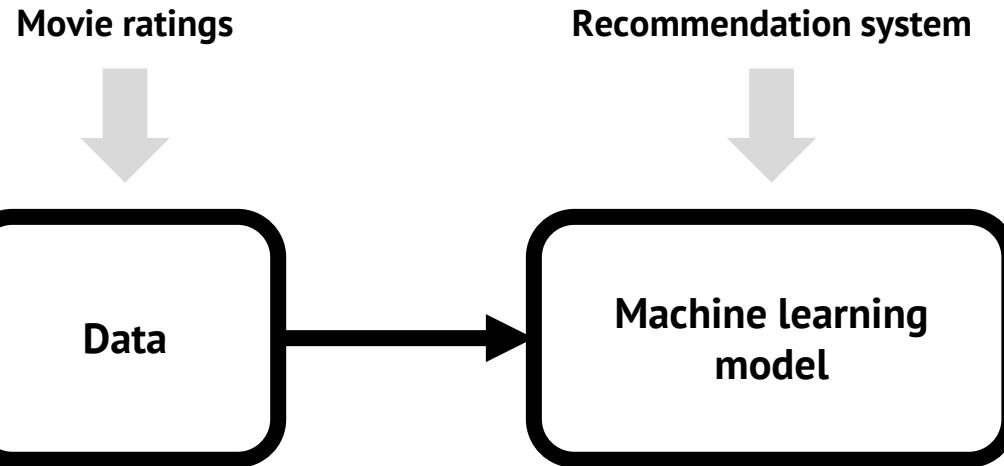
Recommendation engines are machine learning models that predict preferences based on data

Movie ratings

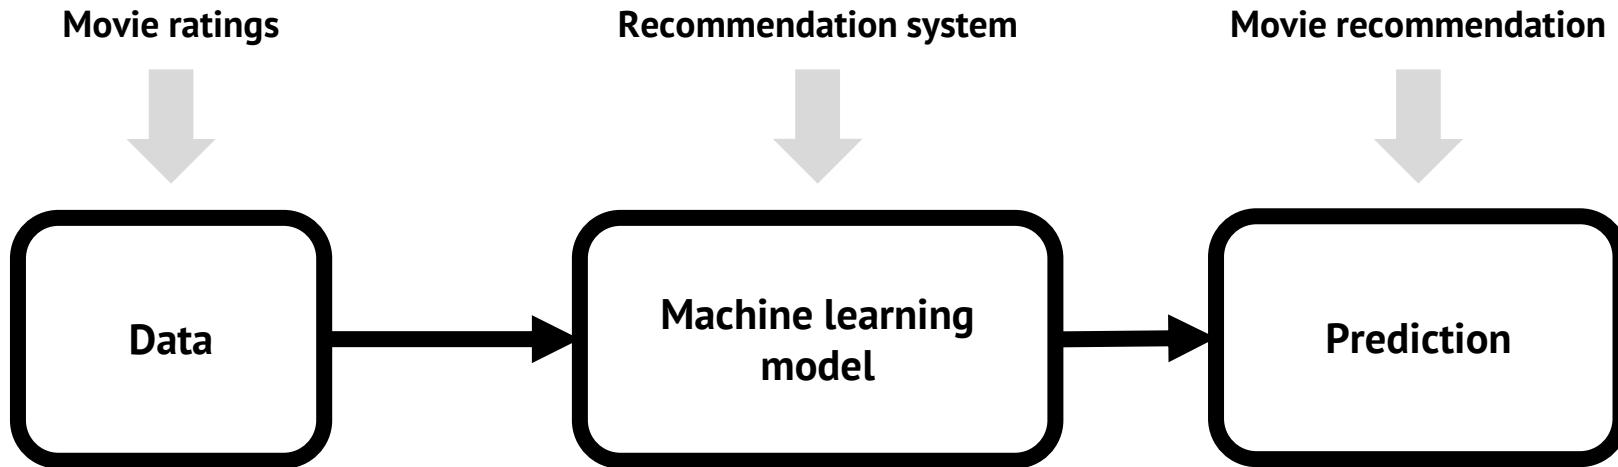


Data

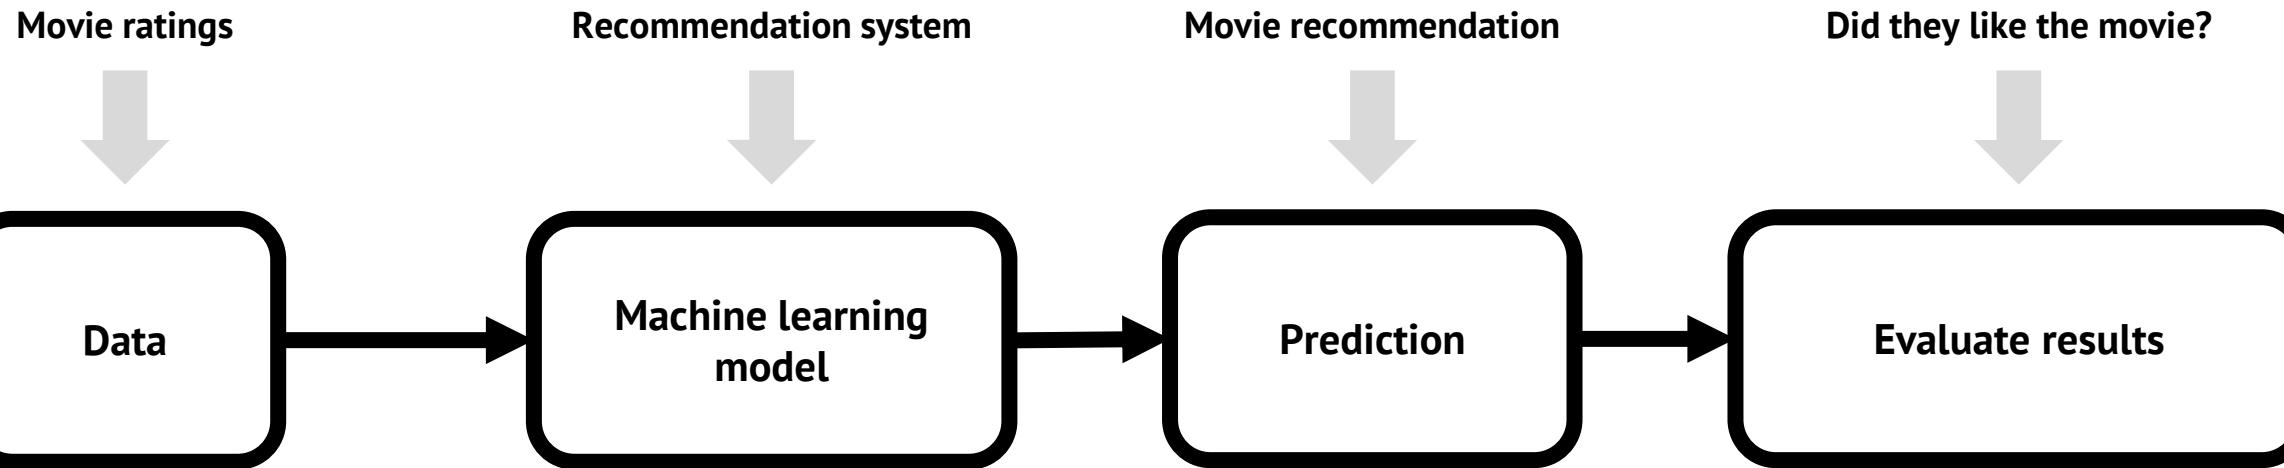
Recommendation engines are machine learning models that predict preferences based on data



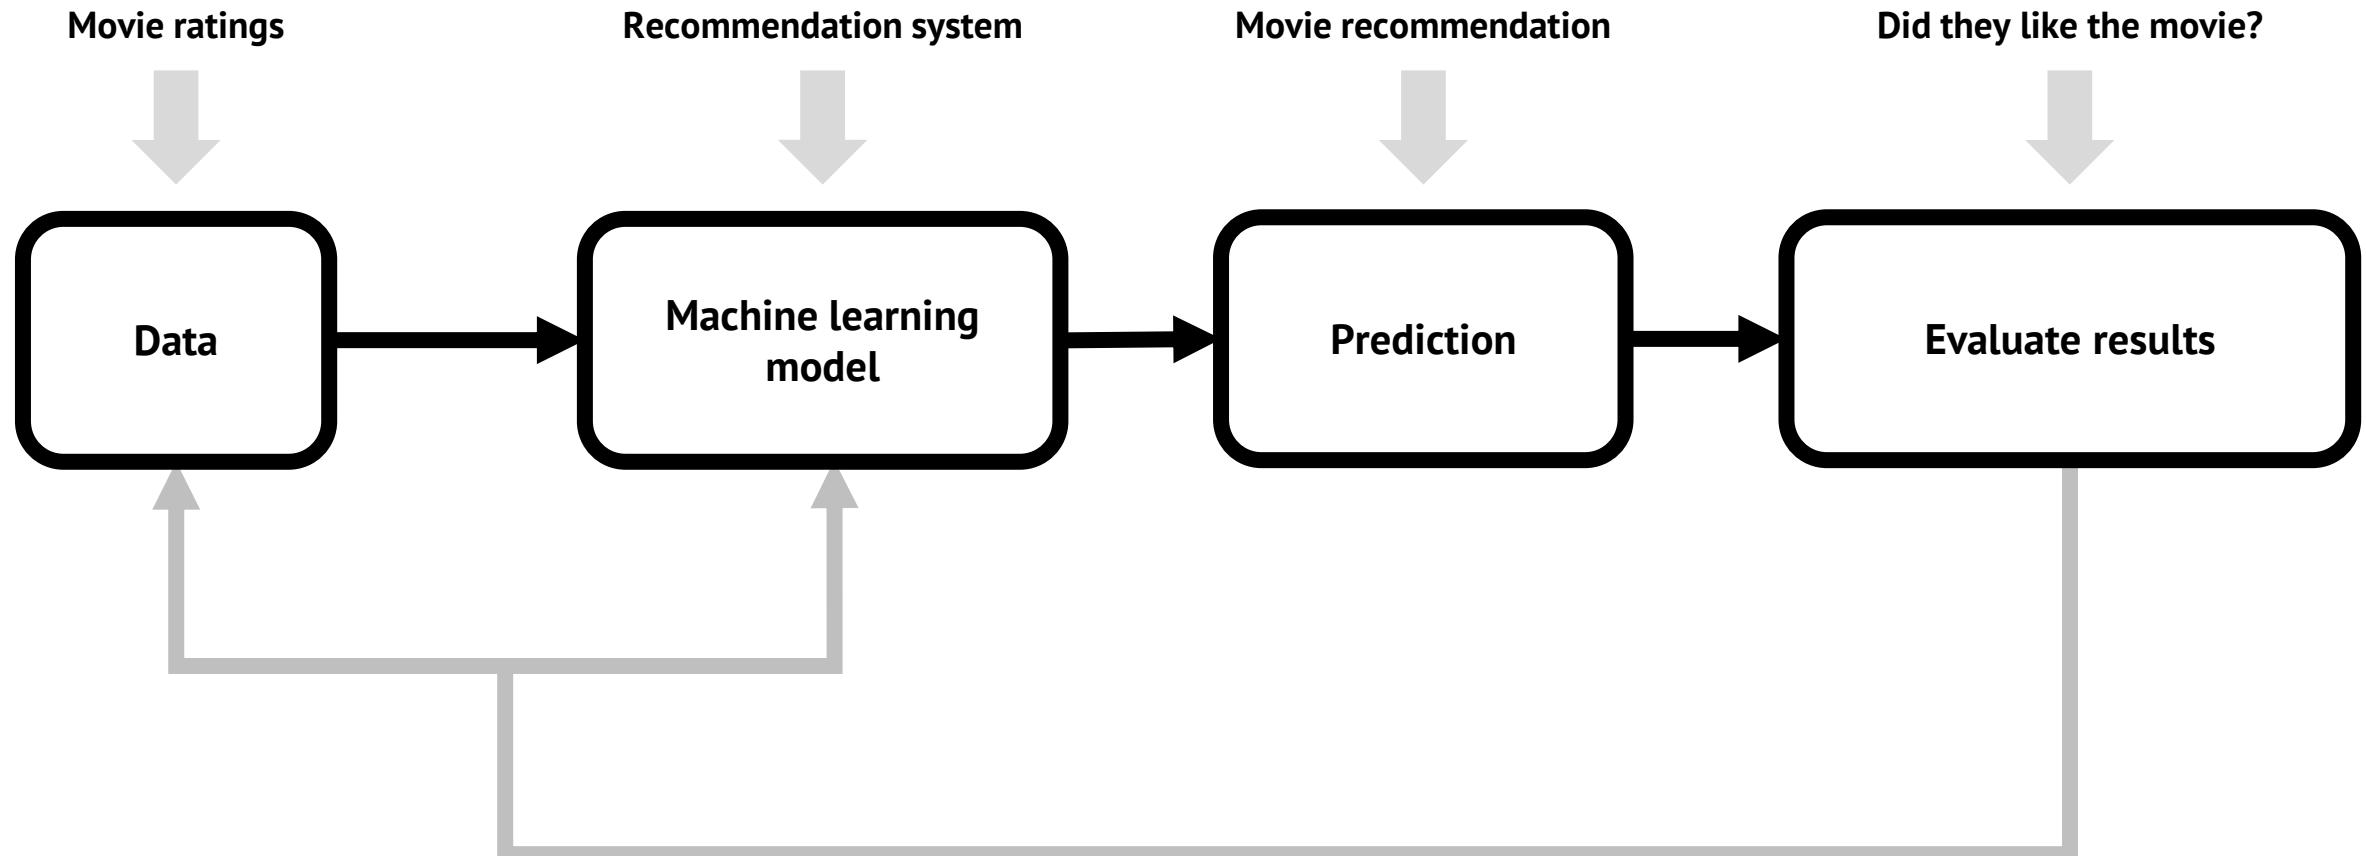
Recommendation engines are machine learning models that predict preferences based on data



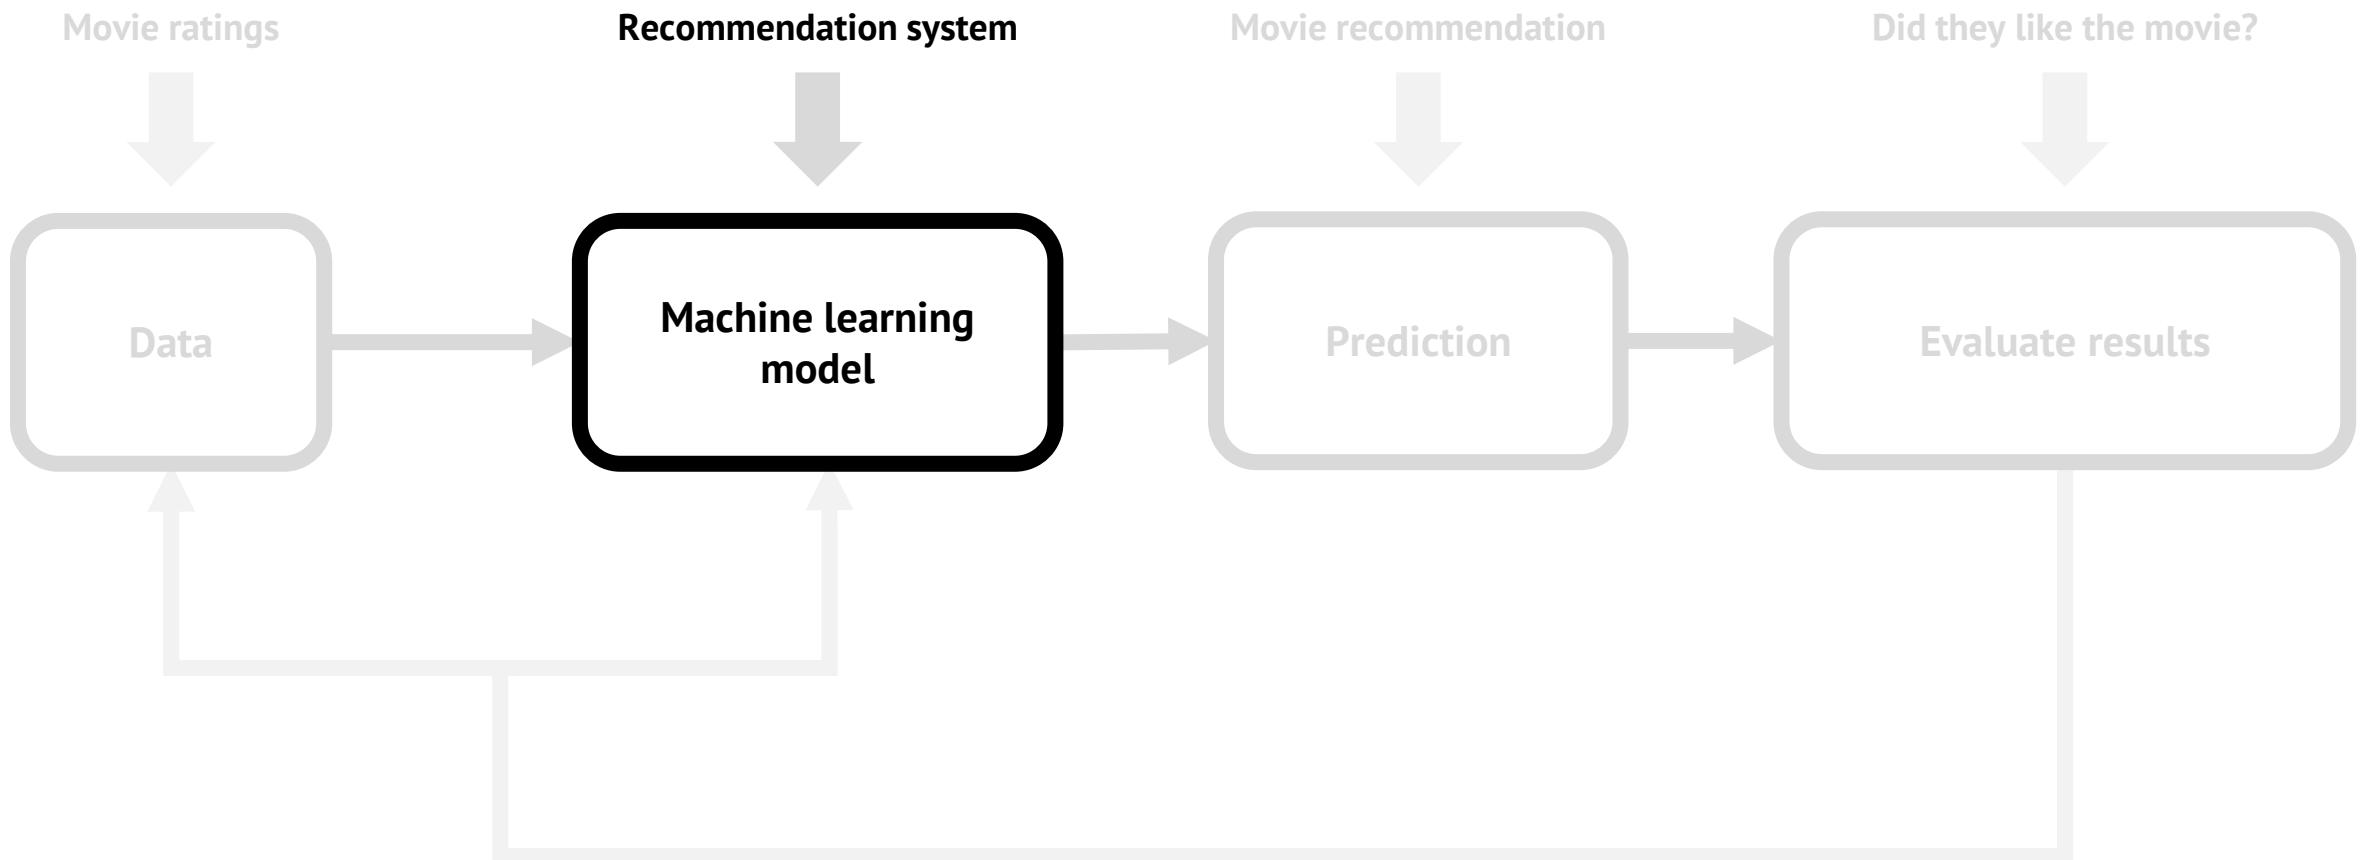
Recommendation engines are machine learning models that predict preferences based on data

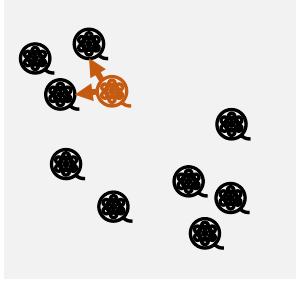


Recommendation engines are machine learning models that predict preferences based on data



Recommendation engines are machine learning models that predict preferences based on data





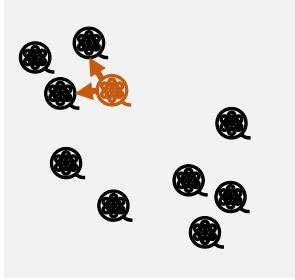
Content-based filtering

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	?	3	10	10
2	2	3	2	5	6

Collaborative filtering



Session-based filtering



Content-based filtering

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	?	3	10	10
2	2	3	2	5	6

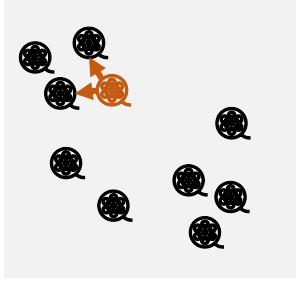
Collaborative filtering



Session-based filtering



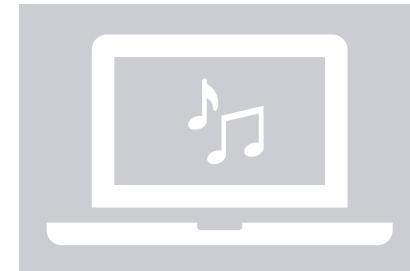
Concluding remarks



Content-based filtering

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	?	3	10	10
2	2	3	2	5	6

Collaborative filtering



Session-based filtering



Concluding remarks

Recommendation systems can for example predict how well you will like something

Predicting ratings

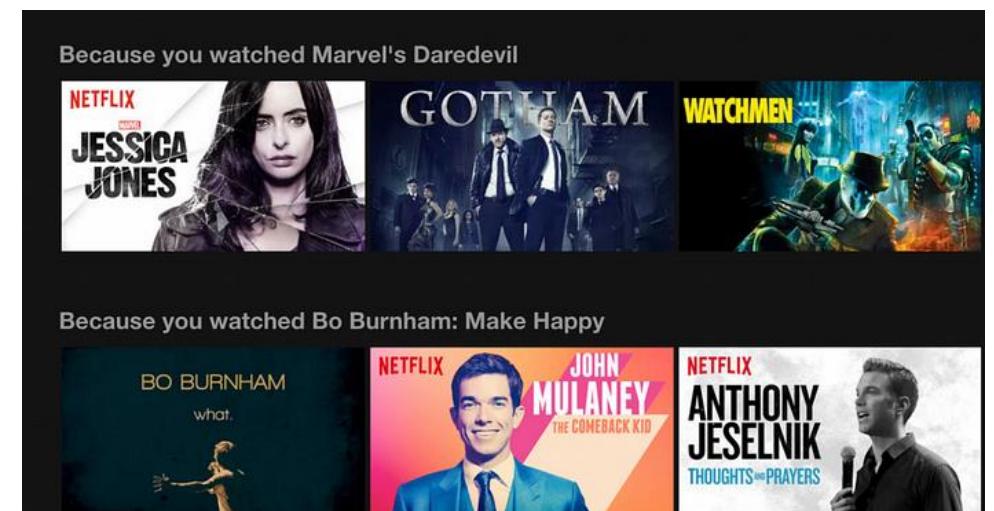
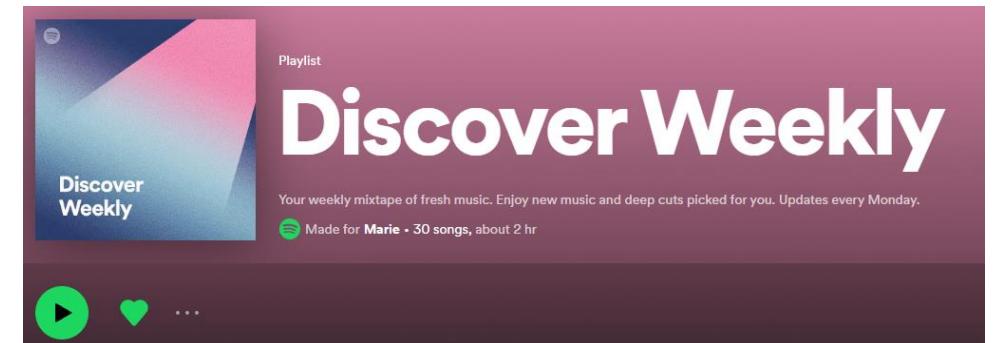


Recommendation systems can for example predict how well you will like something, or give you a list of suggestions for things you might like

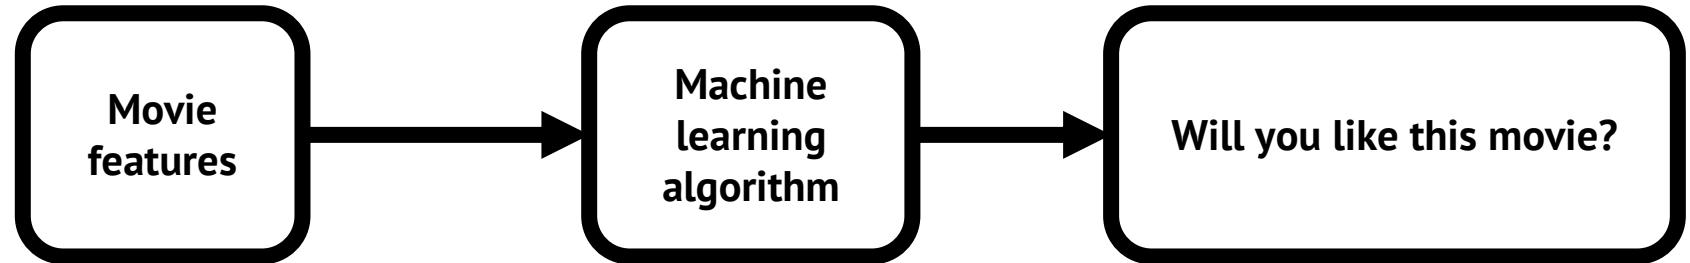
Predicting ratings



Providing suggestions



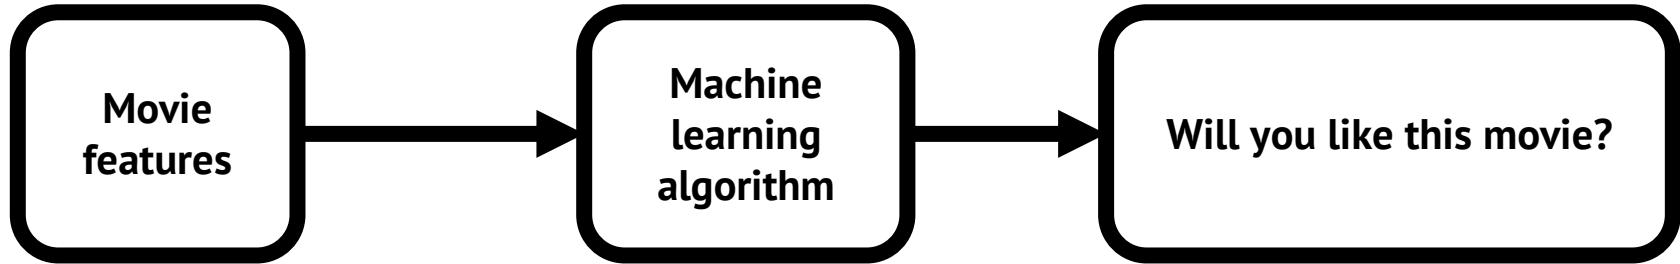
Content-based recommender systems makes recommendations based on features of the content



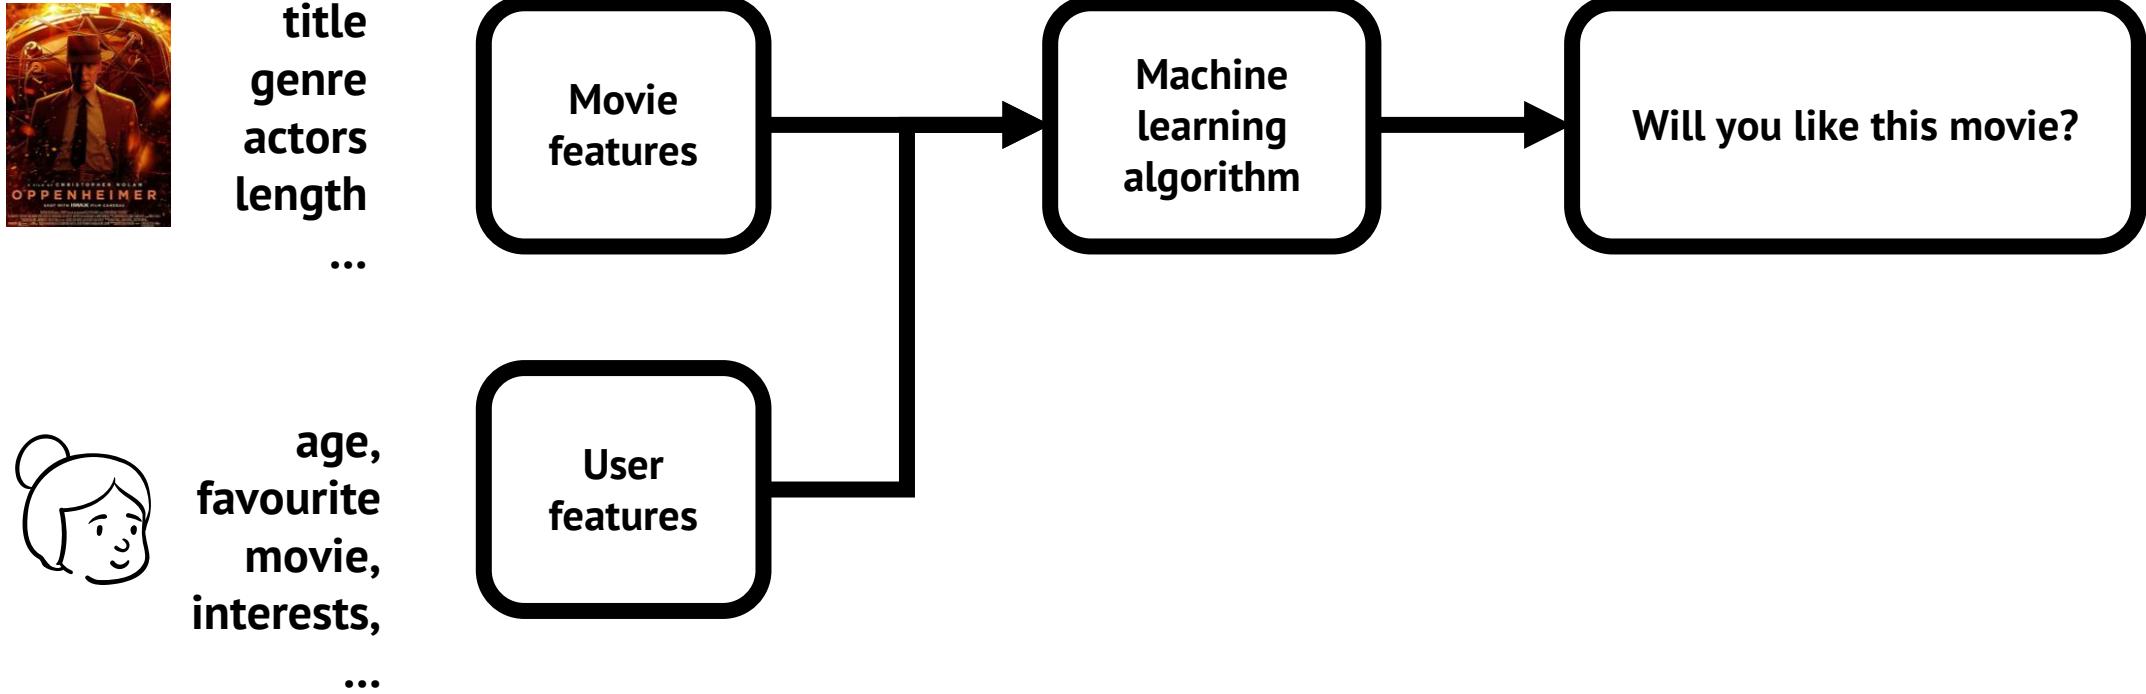
Content-based recommender systems makes recommendations based on features of the content



title
genre
actors
length
...



User features can be combined with item features to provide more personalized recommendations



To find similar movies, we can use the *cosine similarity*

The diagram illustrates the formula for cosine similarity. At the top left, a blue-bordered box contains the text "Cosine similarity". Two blue arrows point from this box down to the variables \mathbf{x}_e and \mathbf{x}_f in the formula below. Below the formula, two blue-bordered boxes provide definitions: the left box says "Movie e's feature vector" and the right box says "Movie f's feature vector".

$$\text{SIM}(\mathbf{x}_e, \mathbf{x}_f) = \frac{\mathbf{x}_e^T \mathbf{x}_f}{\|\mathbf{x}_e\| \|\mathbf{x}_f\|}$$

To find similar movies, we can use the *cosine similarity*

$$\text{SIM}(\mathbf{x}_e, \mathbf{x}_f) = \frac{\sum_i x_{ei}x_{fi}}{\sqrt{\sum_i x_{ei}^2} \sqrt{\sum_i x_{fi}^2}}$$

Cosine similarity

\mathbf{x}_e

\mathbf{x}_f

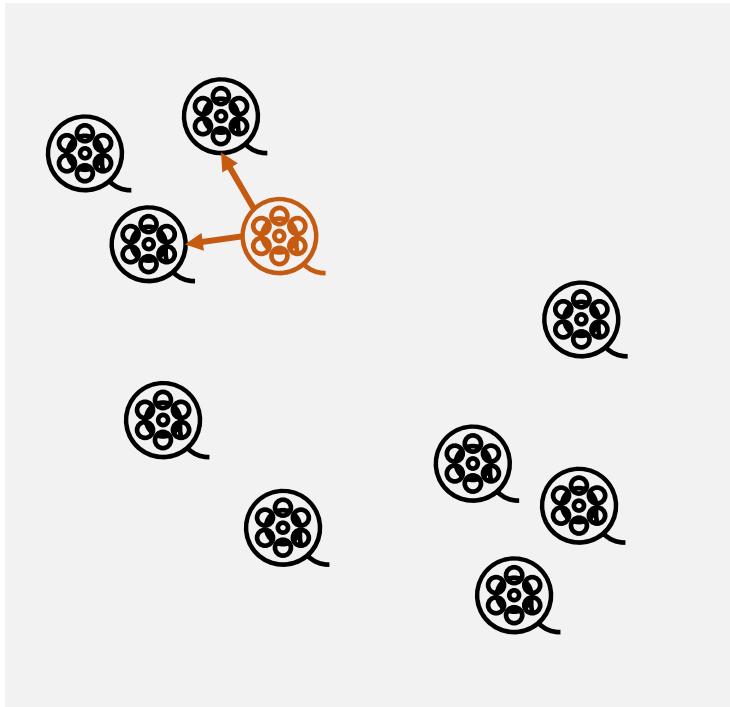
Movie e 's feature vector

Movie f 's feature vector

Feature i of movie e

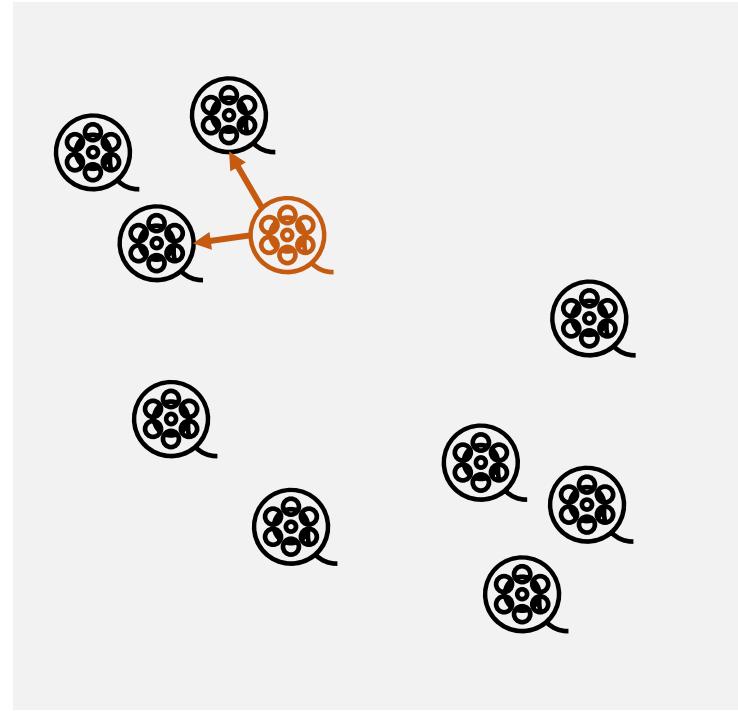
Feature i of movie f

A common algorithm for content based recommendation is KNN



2-nearest neighbours (2NN)

Approximate KNN algorithms is often used when the dataset is very large



2-nearest neighbours (2NN)

≈

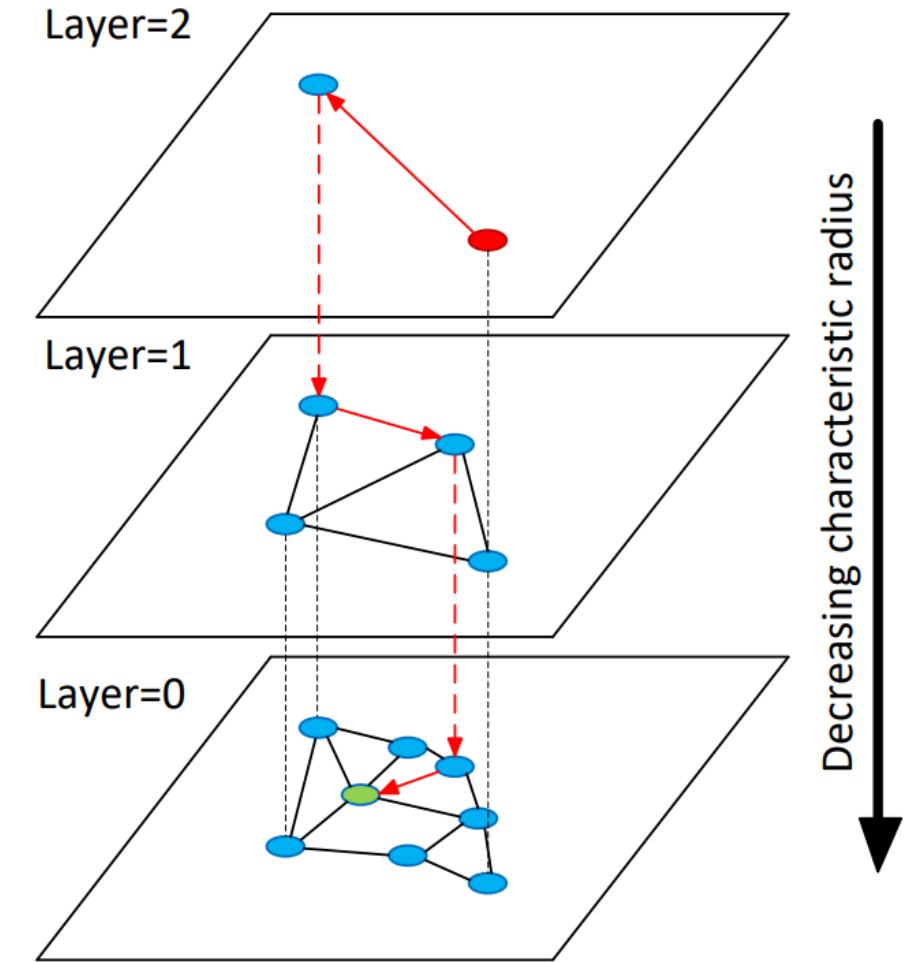


Figure from [Yu. A. Malkov, D. A. Yashunin. IEEE PAMI (2018)]

A key question of content filtering approaches is how to extract features

Similar images

See all



shutterstock®

A key question of content filtering approaches is how to extract features

Similar images

See all



Image embeddings

shutterstock®

A key question of content filtering approaches is how to extract features

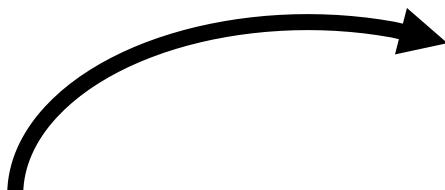


Image and document
embeddings

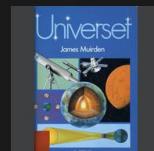
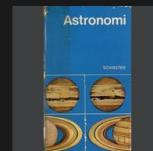
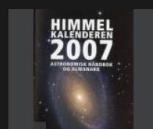
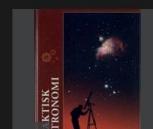
Solen, stjernene og planetene

Stacy, Tom | 1991 | Mer informasjon



Lignende bøker:

Vis flere utgaver av samme bok Vis kun tilgjengelige bøker

 <p>Planetene Couper, Heather 1992 Likhet: O O O O O</p>	 <p>Universet Muirden, James 1990 Likhet: O O O O O</p>	 <p>Astronomi Selje, B.E. 1971 Likhet: O O O O O</p>
 <p>Himmelkalenderen 2007 Ovaldsen, Jan-Erik 2006 Likhet: O O O O O</p>	 <p>Den store boken om astronomi Newth, Eirik Brekke, Pål 2009 Likhet: O O O O O</p>	 <p>Praktisk astronomi Fjørtoft, Magnar 2001 Likhet: O O O O O</p>

 Slik fungerer I



Nasjonalbiblioteket
National Library of Norway

For some applications simple and interpretable features are ideal

Påstand 1 av 23

Det er fjernet for mange parkeringsplasser i Oslo sentrum. De må komme tilbake.

Uenig **Delvis uenig** Vet ikke Delvis enig Enig

Bakgrunn: I Oslo er det siden 2016 fjernet om lag 760 gateparkeringsplasser i sentrumssonene innenfor Ring 1. I hele byen er det fjernet 6000 parkeringsplasser i samme periode, ifølge Rundtakstataren. Formålet har vært å avvise forflytning til utdeler, månde om man ikke i sentrum.

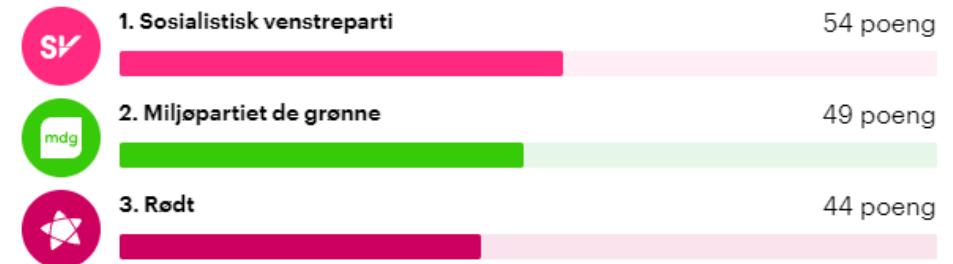
Påstand 2 av 23

Private kommersielle selskaper bør få driftet flere barnehager, sykehjem og barnevernsinstitusjoner enn i dag.

Uenig Delvis uenig Vet ikke Delvis enig Enig

Bakgrunn: Dagens rødgrønne flertall i Oslo sier nei til sykehjem drevet av private kommersielle selskaper. De sier også nei til nye private kommersielle barnehager. Argumentet er at private selskaper ikke skal tjene penger på velferd. De merke selskaperne mener private kommersielle

Du er mest enig med disse partiene:



Du og **Sosialistisk venstreparti** er blant annet enig om:

- **Begge mener** at bompenge kan økes for å finansiere utbygging av kollektivtransport i Oslo.
- **Begge mener** at eiendomskatt trengs for å betale tjenester som sykehjem og barnehageplasser.
- **Begge mener** at det bør innføres kjøttfrie dager i kommunens kantiner og institusjoner.

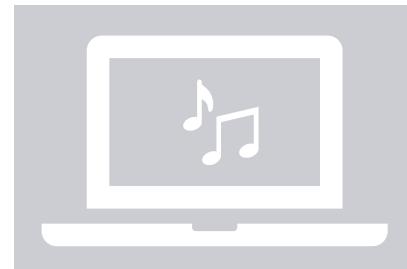
Se hele oversikten under.



Content-based filtering

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	?	3	10	10
2	2	3	2	5	6

Collaborative filtering

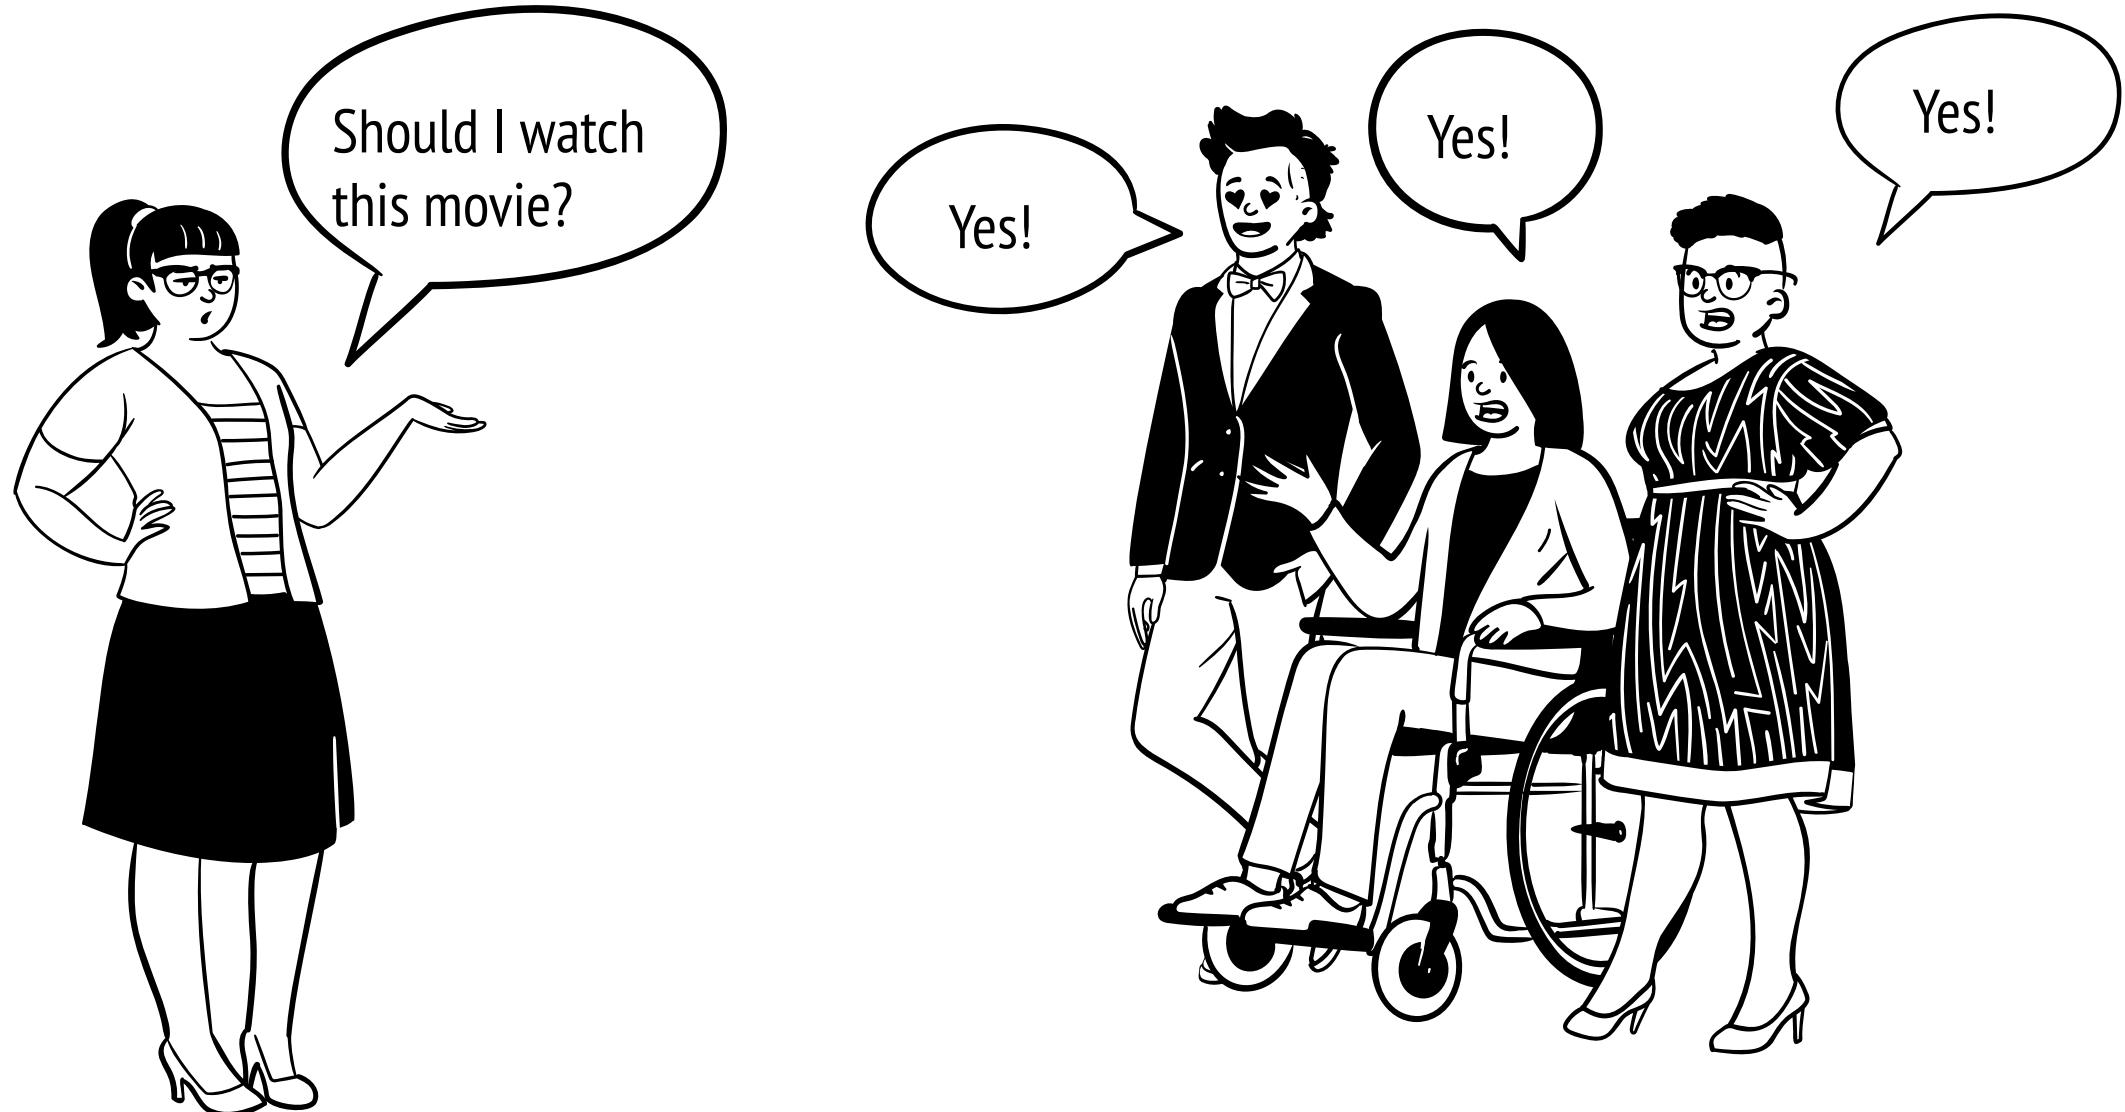


Session-based filtering



Concluding remarks

Collaborative filtering uses similar users' ratings to predict if you will enjoy a movie





Raise your hand if you have seen
“Kill Bill” and **really liked it!**

A close-up photograph of several hands raised against a dark green background. One hand is prominently featured in the foreground, palm facing forward, while others are visible in the background.

Raise your hand if you have seen
“Kill Bill” and **didn’t like it!**

I can use the average rating as an estimate for how much I will like a movie

$$\hat{x}_{ei} = \frac{1}{U} \sum_{u=1}^U x_{eu}$$

Estimate of my rating of the movie

User u 's rating of the movie

The equation shows the calculation of an estimated rating \hat{x}_{ei} for a movie i . It is the average of all ratings x_{eu} given by users u from 1 to U . Two annotations are present: a box labeled 'Estimate of my rating of the movie' with two arrows pointing to the term x_{eu} , and another box labeled 'User u 's rating of the movie' with one arrow pointing to the same term.

I can weight other users' ratings based on how similar their taste is to mine

$$\hat{x}_{ei} = \frac{1}{N} \sum_{u=1}^U x_{eu} \times \text{SIM}(\mathbf{x}_{:,v}, \mathbf{x}_{:,u})$$

Estimate of my rating of the movie

User u 's rating of the movie

I can weight other users' ratings based on how similar their taste is to mine

$$\hat{x}_{ei} = \frac{1}{N} \sum_{u=1}^U x_{eu} \times \text{SIM}(\mathbf{x}_{:,v}, \mathbf{x}_{:,u})$$

Diagram illustrating the components of the equation:

- Estimate of my rating of the movie: \hat{x}_{ei}
- User u 's rating of the movie: x_{eu}
- My previous ratings: $\mathbf{x}_{:,v}$
- User u 's previous ratings: $\mathbf{x}_{:,u}$

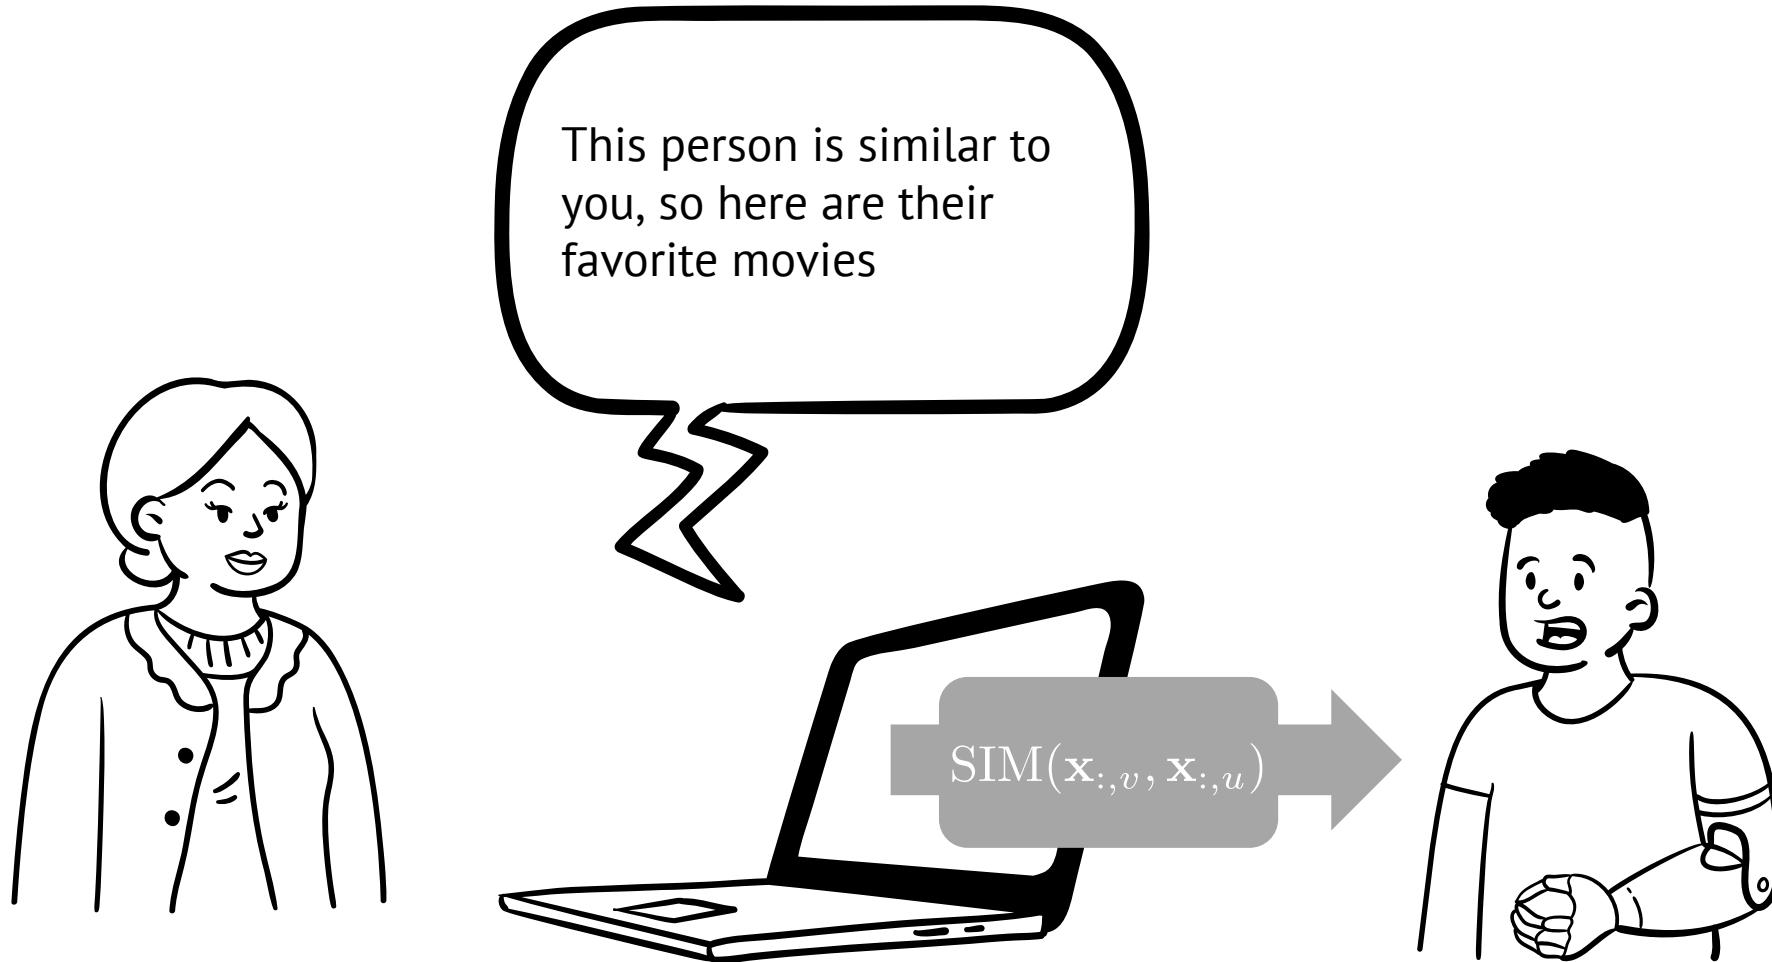
I can weight other users' ratings based on how similar their taste is to mine

$$\hat{x}_{ei} = \frac{1}{N} \sum_{u=1}^U x_{eu} \times \text{SIM}(\mathbf{x}_{:,v}, \mathbf{x}_{:,u})$$

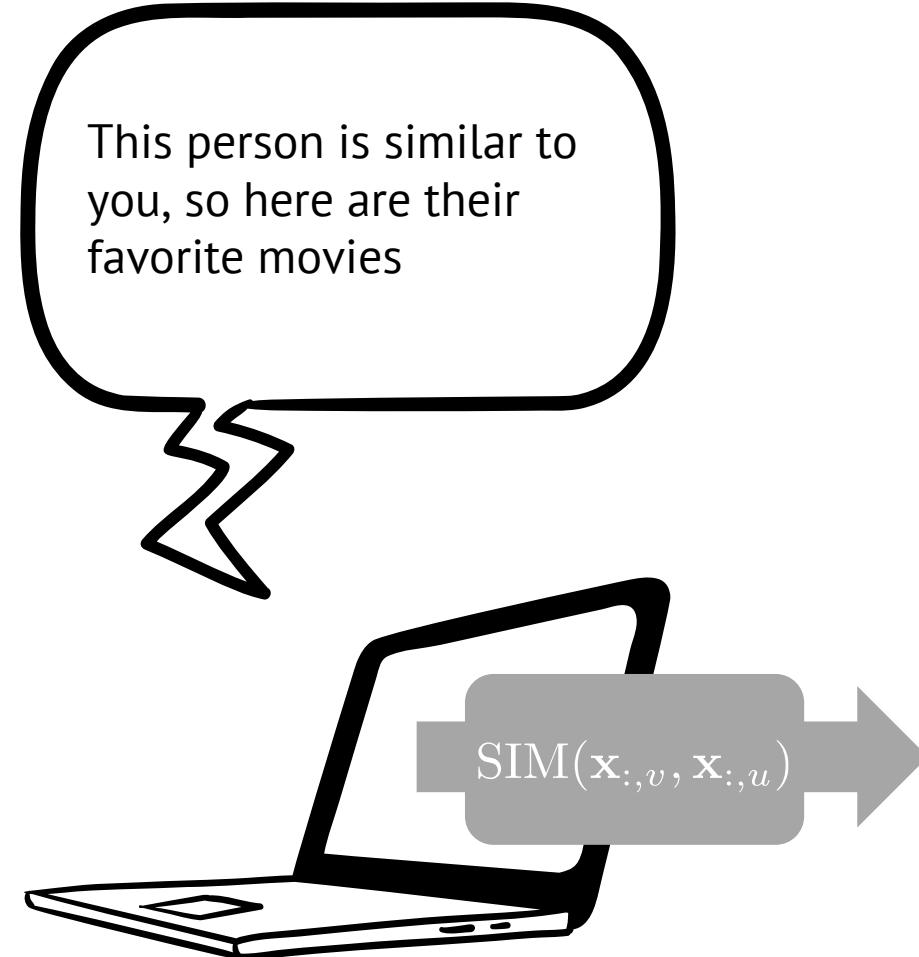
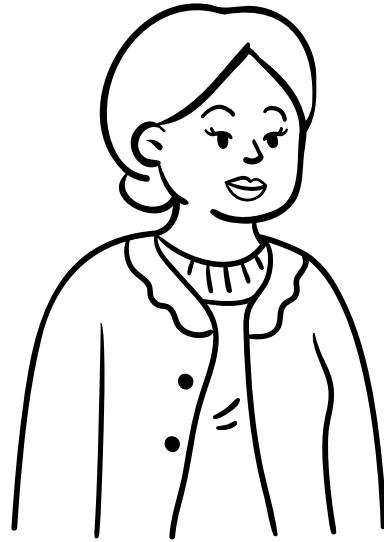
Diagram illustrating the components of the equation:

- Estimate of my rating of the movie**: Points to the term \hat{x}_{ei} .
- User u 's rating of the movie**: Points to the term x_{eu} .
- My previous ratings**: Points to the term $\mathbf{x}_{:,v}$.
- User u 's previous ratings**: Points to the term $\mathbf{x}_{:,u}$.
- Cosine similarity**: Points to the function $\text{SIM}(\mathbf{x}_{:,v}, \mathbf{x}_{:,u})$.
- Sum of all cosine similarities**: Points to the denominator N .

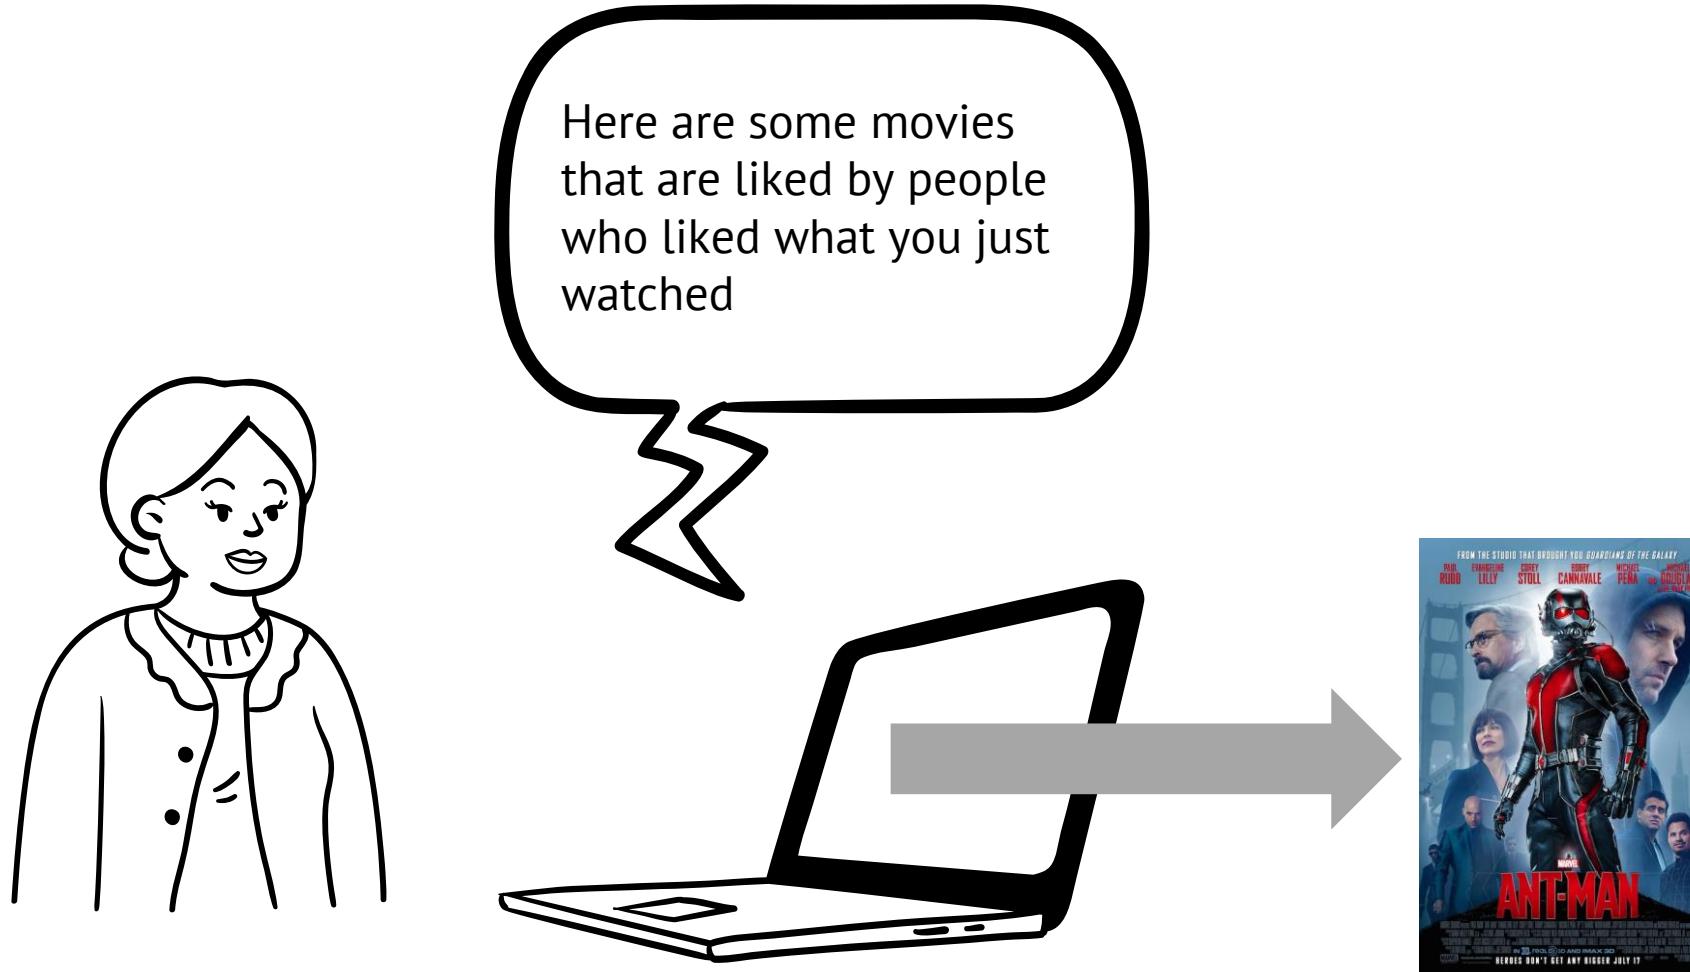
We can also use similarity recommend movies liked by users with similar taste as me



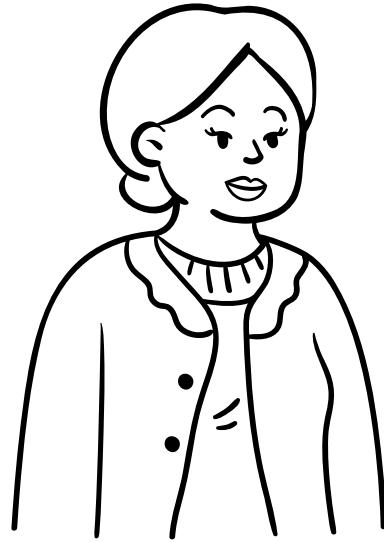
We can also use similarity recommend movies liked by users with similar taste as me



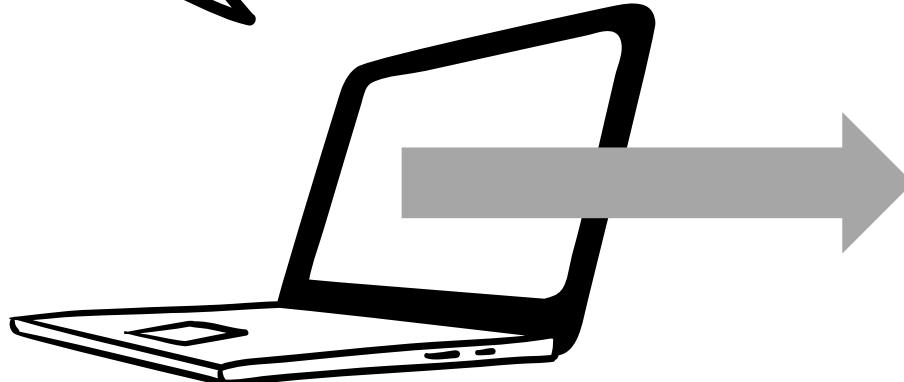
Or use similarity to recommend movies liked by people who also like what you just watched



Or use similarity to recommend movies liked by people who also like what you just watched



Here are some movies
that are liked by people
who liked what you just
watched



$\text{SIM}(\mathbf{x}_e, \mathbf{x}_f)$

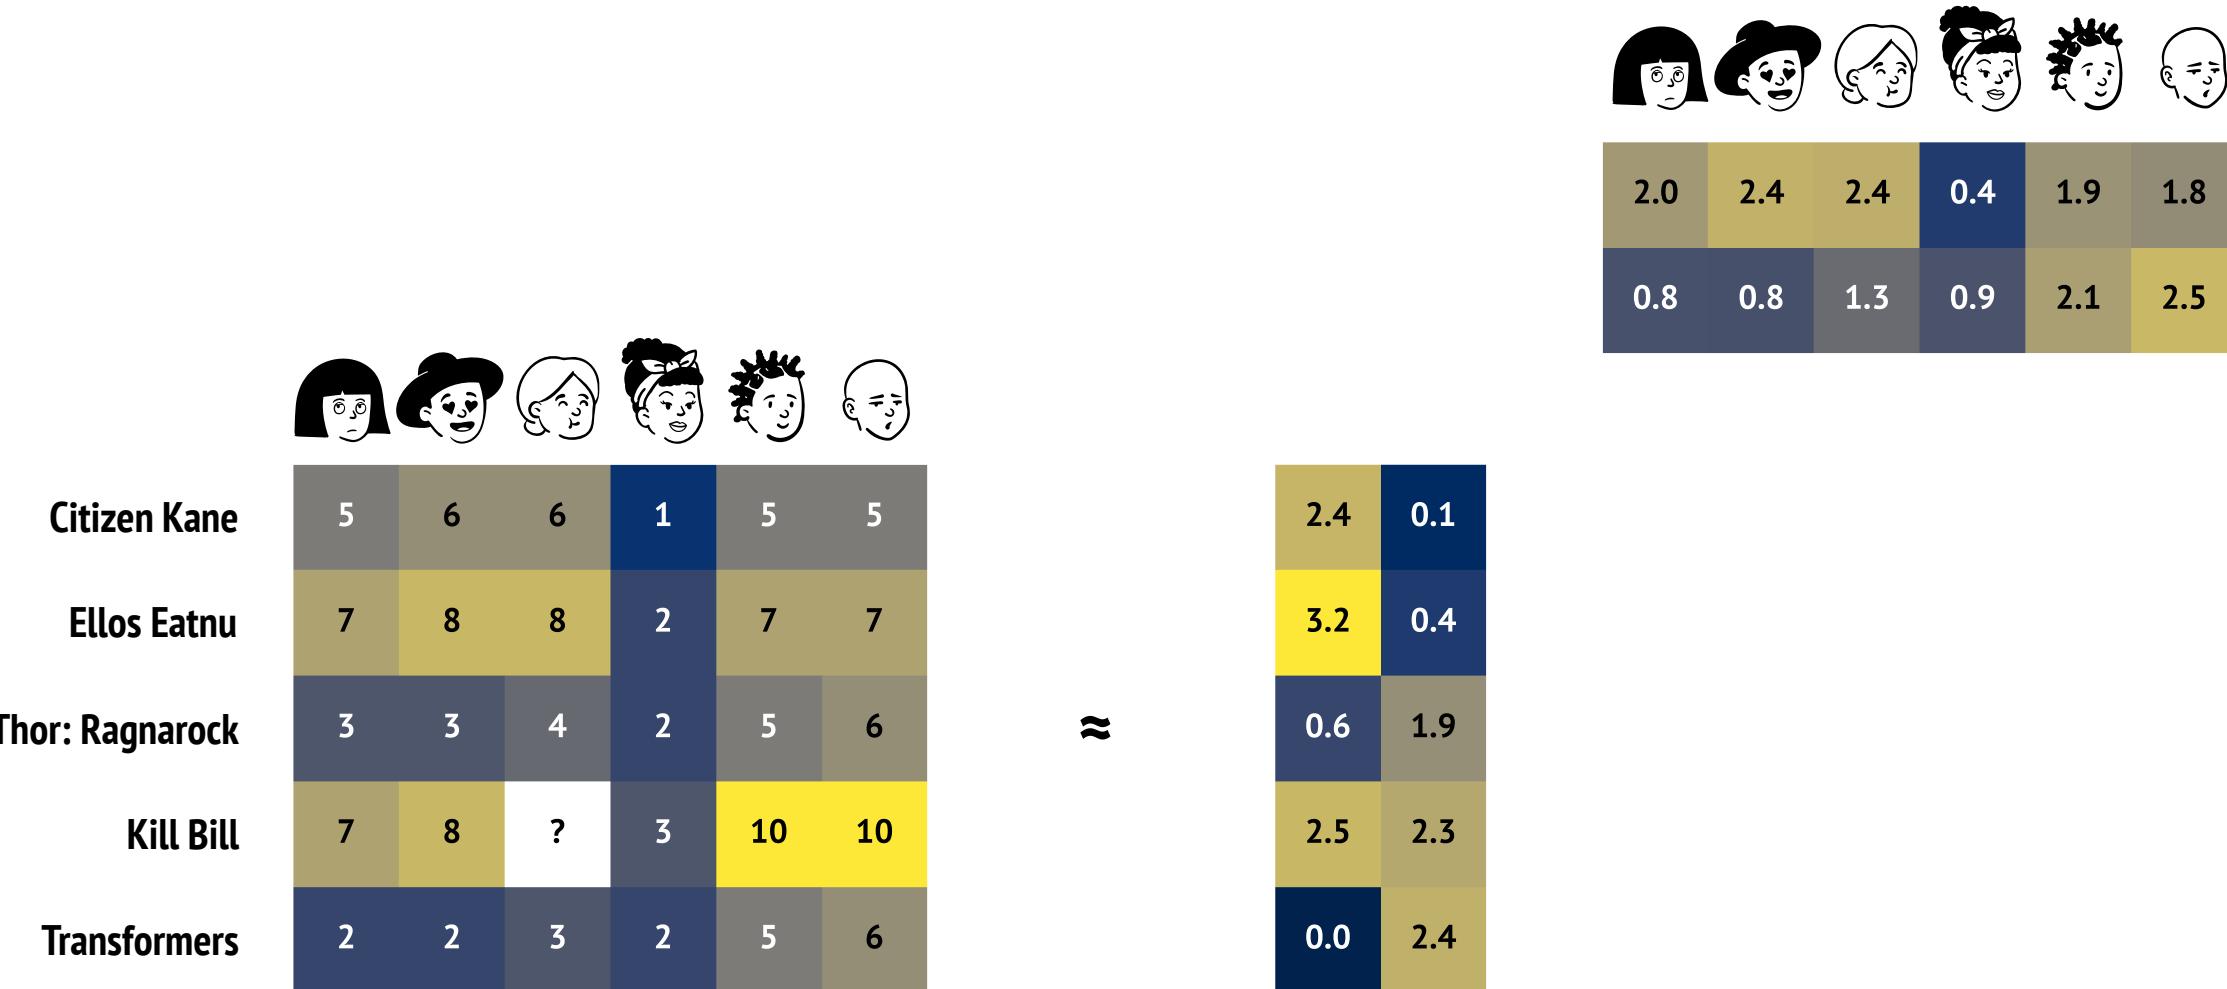
Model based collaborative filtering finds a mathematical model that describe the movie ratings



Matrix factorisations can find underlying factors that describe movie ratings

Citizen Kane	5	6	6	1	5	5
Ellos Eatnu	7	8	8	2	7	7
Thor: Ragnarock	3	3	4	2	5	6
Kill Bill	7	8	?	3	10	10
Transformers	2	2	3	2	5	6

Matrix factorisations can find underlying factors that describe movie ratings

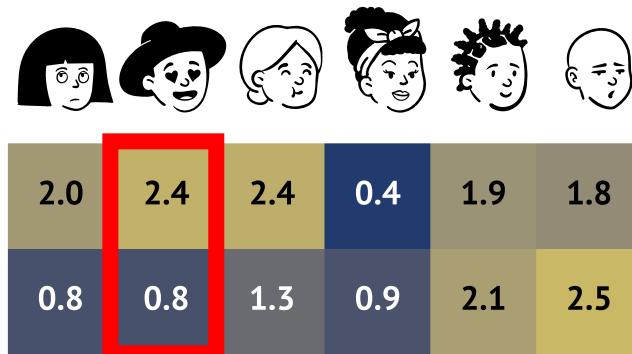


Matrix factorisations can find underlying factors that describe movie ratings

	1	2	3	4	5	6
Citizen Kane	5	6	6	1	5	5
Ellos Eatnu	7	8	8	2	7	7
Thor: Ragnarock	3	3	4	2	5	6
Kill Bill	7	8	?	3	10	10
Transformers	2	2	3	2	5	6

≈

2.4	0.1
3.2	0.4
0.6	1.9
2.5	2.3
0.0	2.4



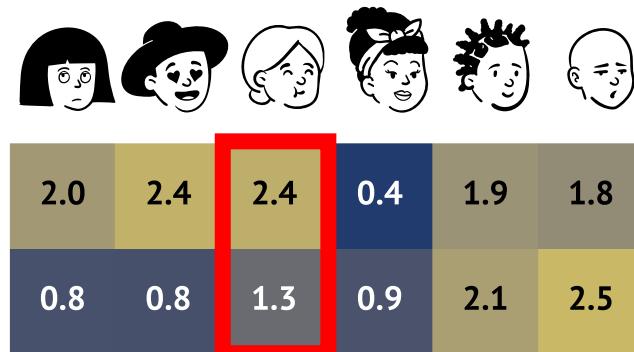
$$3.2 \times 2.4 + 0.4 \times 0.8 = 8$$

Matrix factorisations can find underlying factors that describe movie ratings

	1	2	3	4	5	6
Citizen Kane	5	6	6	1	5	5
Ellos Eatnu	7	8	8	2	7	7
Thor: Ragnarock	3	3	4	2	5	6
Kill Bill	7	8	?	3	10	10
Transformers	2	2	3	2	5	6

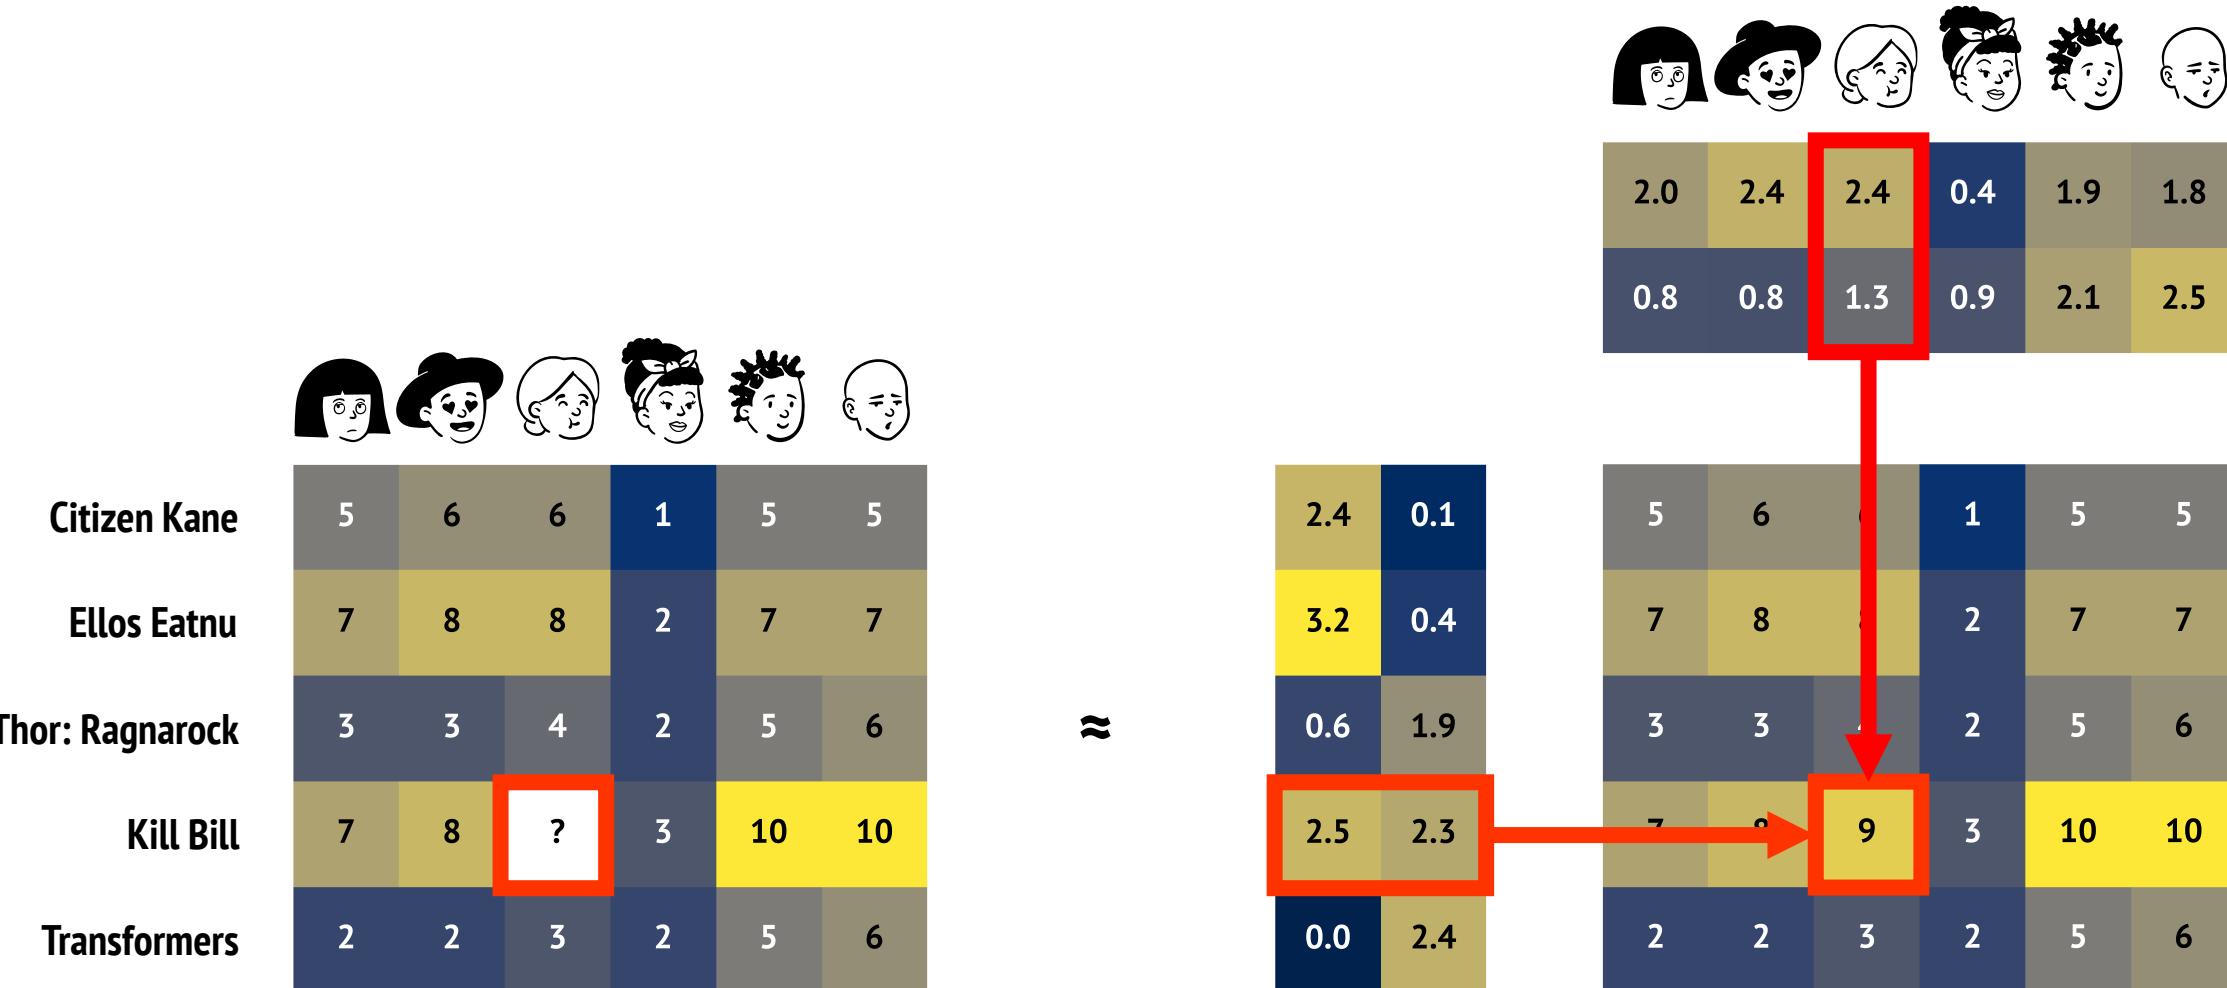
≈

2.4	0.1
3.2	0.4
0.6	1.9
2.5	2.3
0.0	2.4



$$2.4 \times 2.5 + 1.3 \times 2.3 = 8.99$$

Matrix factorisations can find underlying factors that describe movie ratings



$$X \approx EU^T$$

To find the embeddings, we minimise the squared reconstruction error

Minimise error
for known
movie scores

Product of
embedding
vectors

$$\min_{\mathbf{E}, \mathbf{U}} \sum_{e,u} (x_{eu} - \mathbf{e}_{e,:}^T \mathbf{u}_{u,:})^2$$

We often add regularisation to reduce the chance of overfitting

Minimise error
for known
movie scores

Product of
embedding
vectors

Regularisation
to avoid
overfitting

Regularisation
to avoid
overfitting

$$\min_{\mathbf{E}, \mathbf{U}} \sum_{e,u} (x_{eu} - \mathbf{e}_{e,:}^T \mathbf{u}_{u,:})^2 + \gamma_e \|\mathbf{E}\|_F^2 + \gamma_u \|\mathbf{U}\|_F^2$$

We often add regularisation to reduce the chance of overfitting

Minimise error
for known
movie scores

Product of
embedding
vectors

Regularisation
to avoid
overfitting

Regularisation
to avoid
overfitting

$$\min_{\mathbf{E}, \mathbf{U}} \sum_{e,u} (x_{eu} - \mathbf{e}_{e,:}^T \mathbf{u}_{u,:})^2 + \gamma_e \|\mathbf{E}\|_F^2 + \gamma_u \|\mathbf{U}\|_F^2$$

s.t. $\mathbf{E}, \mathbf{U} \geq 0$ elementwise

Nonnegativity to reduce
overfitting and improve
interpretability

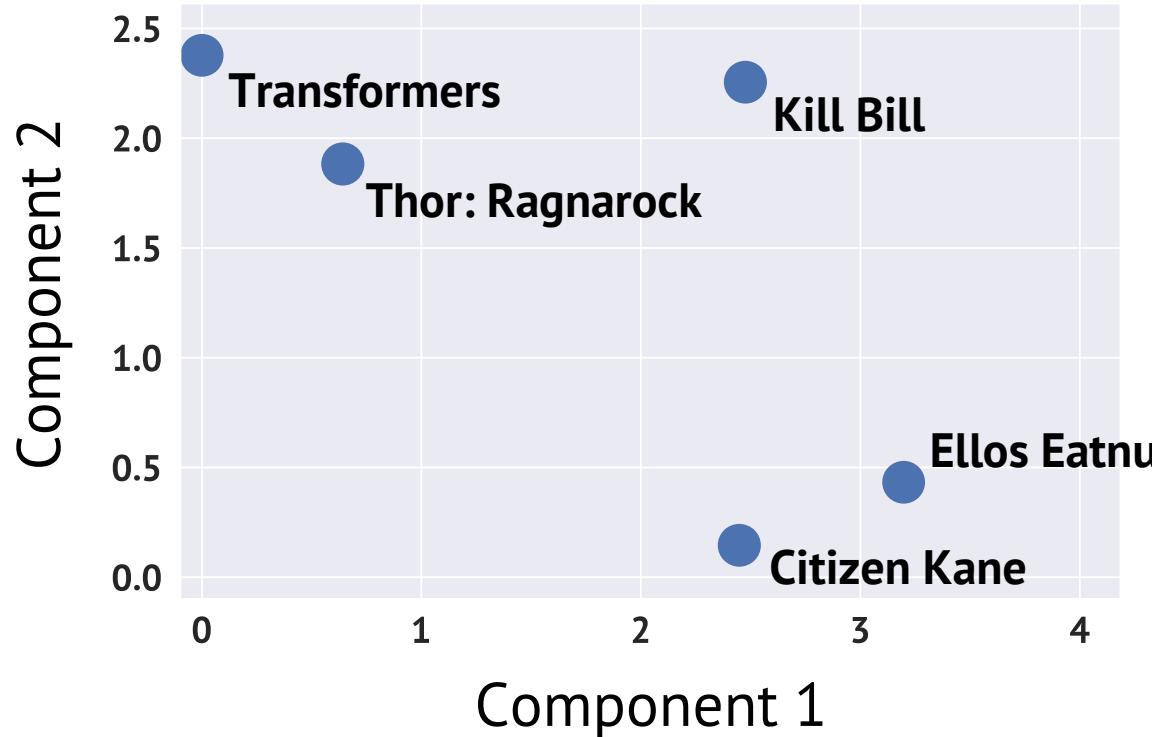
We can also use nonlinear embeddings, e.g., by fitting neural networks to reconstruct our data

$$\min_{\mathbf{E}, \mathbf{U}, \mathcal{W}} \sum_{e,u} (x_{eu} - f(\mathbf{e}_{e,:}, \mathbf{u}_{u,:}; \mathcal{W}))^2 + \gamma_e \|\mathbf{E}\|_F^2 + \gamma_u \|\mathbf{U}\|_F^2$$

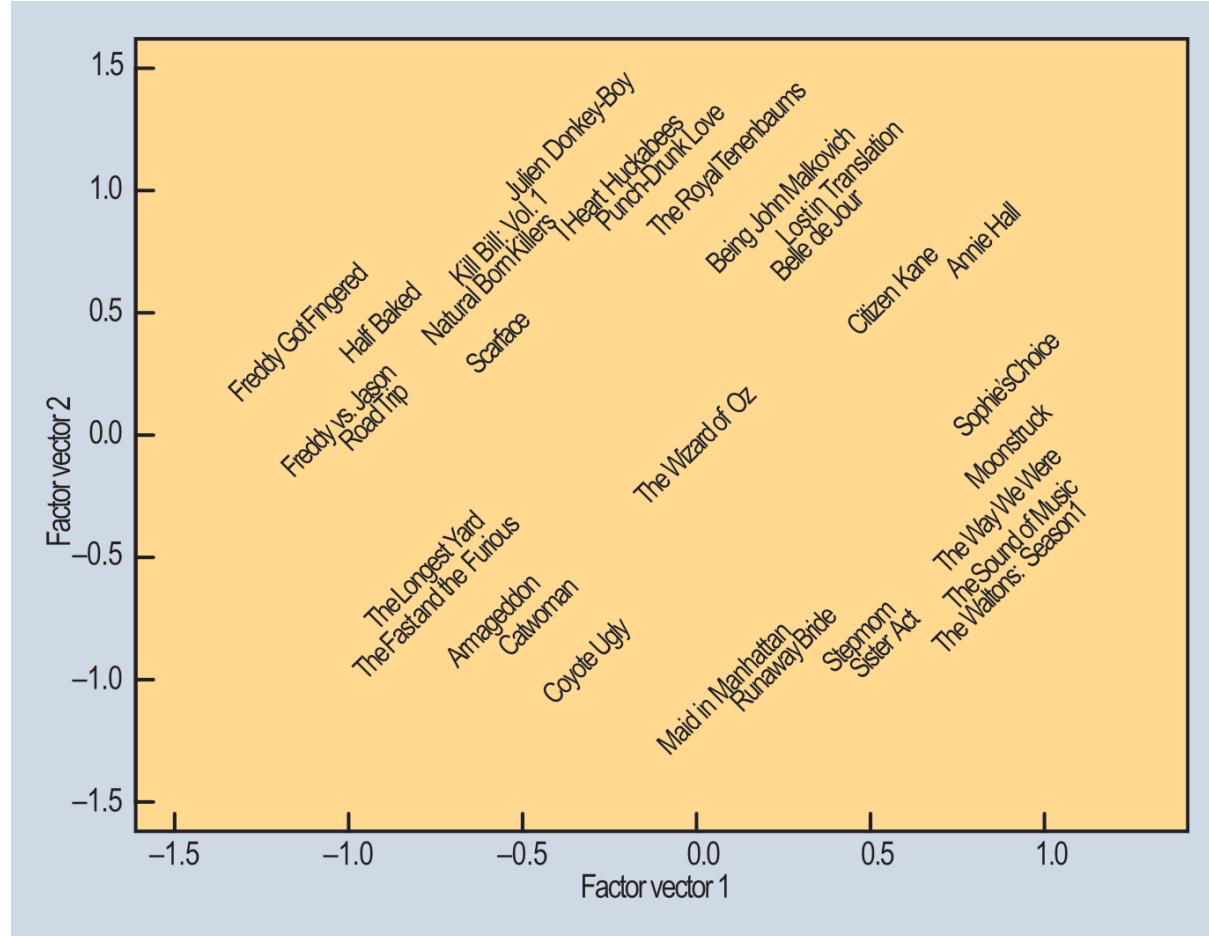
Learned function to
reconstruct the movie
scores from embeddings

Parameters of
the learned
function

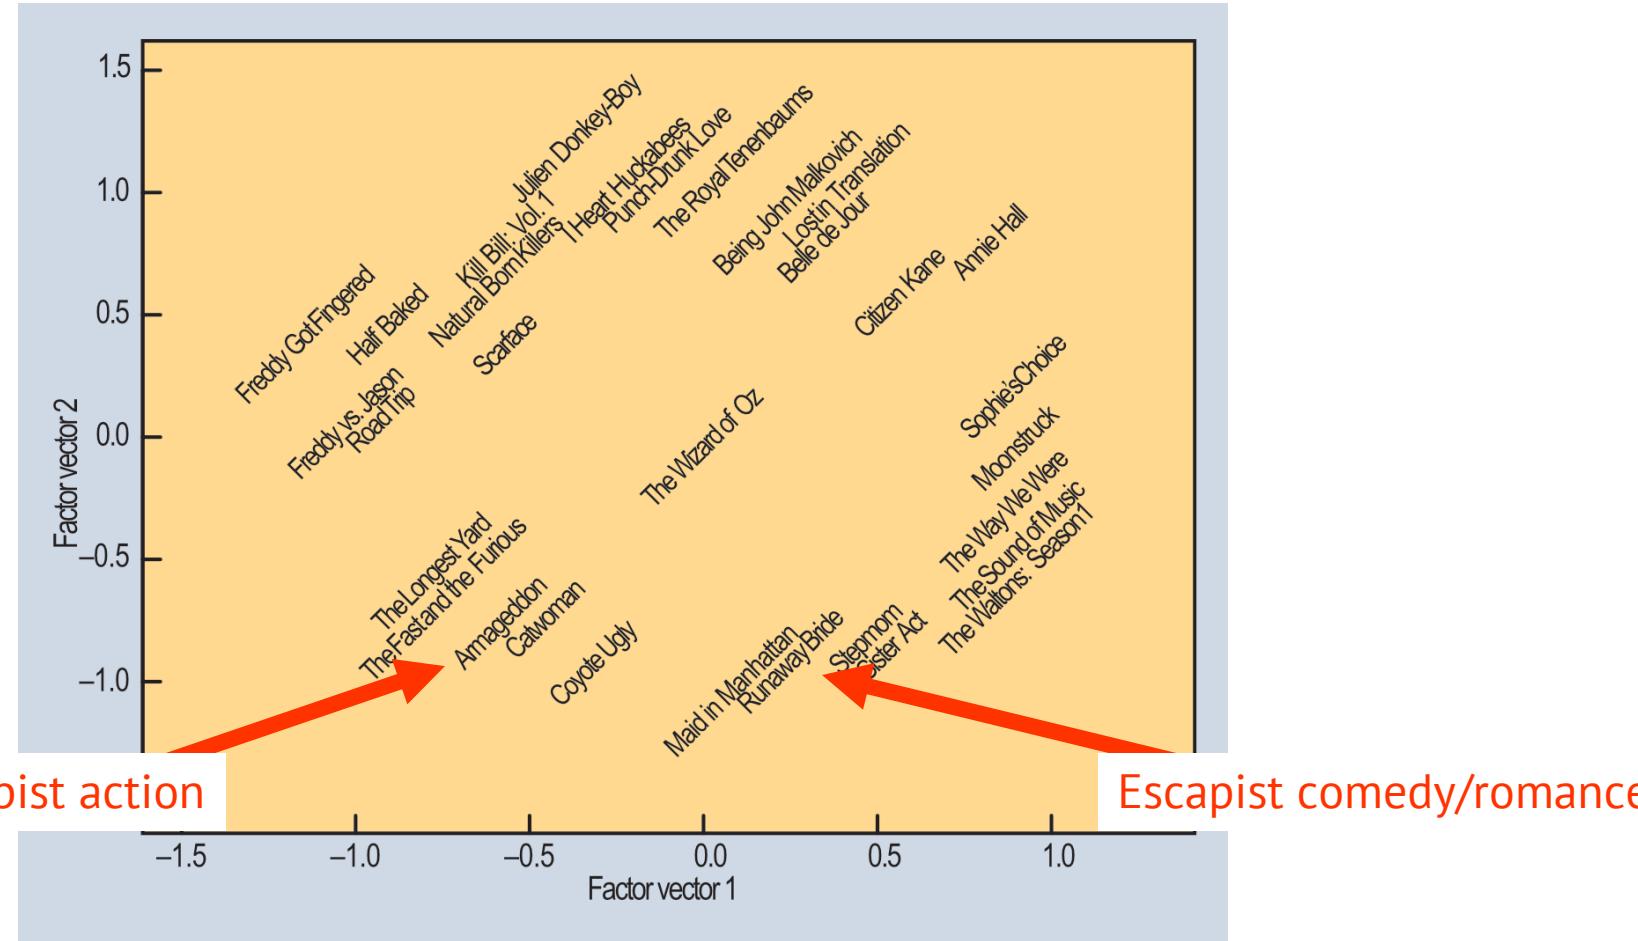
The latent factors embed movies in a space where similar movies are close



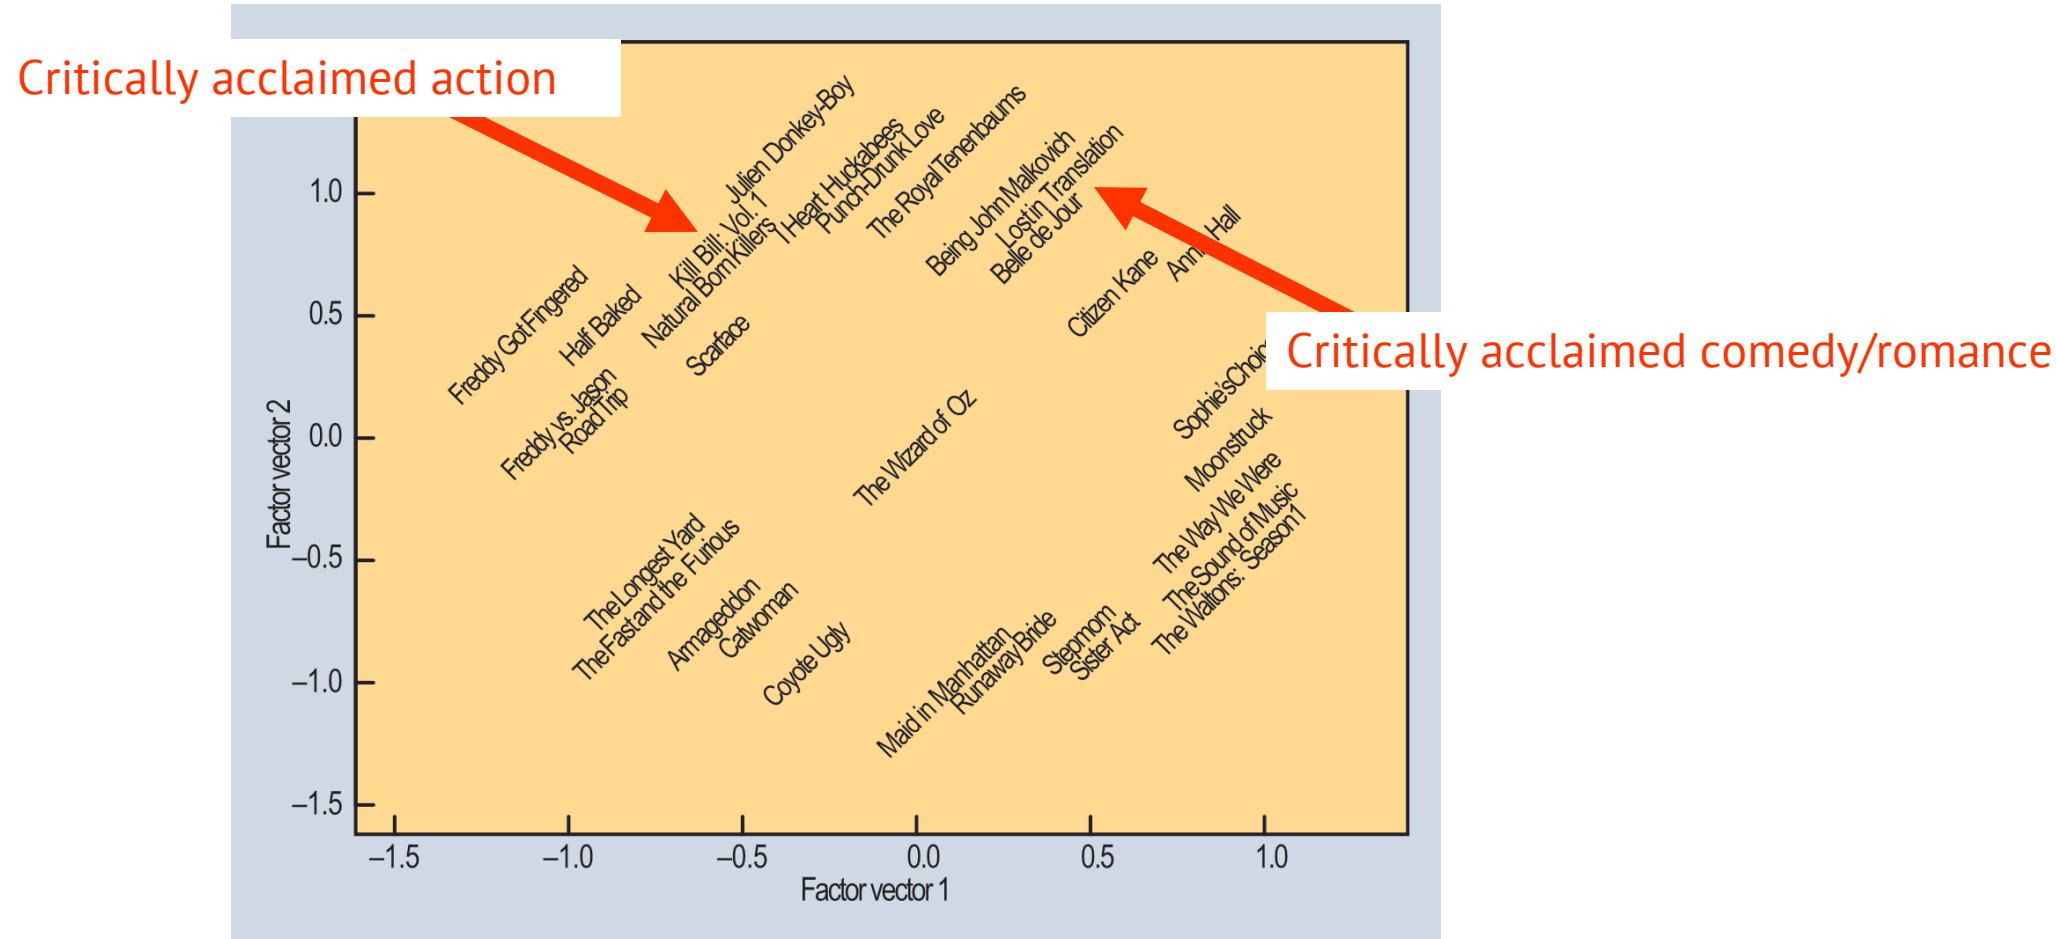
The latent factors embed movies in a space where similar movies are close



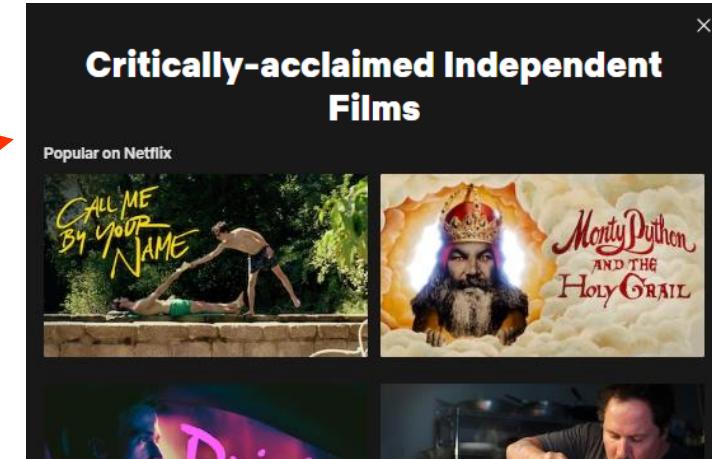
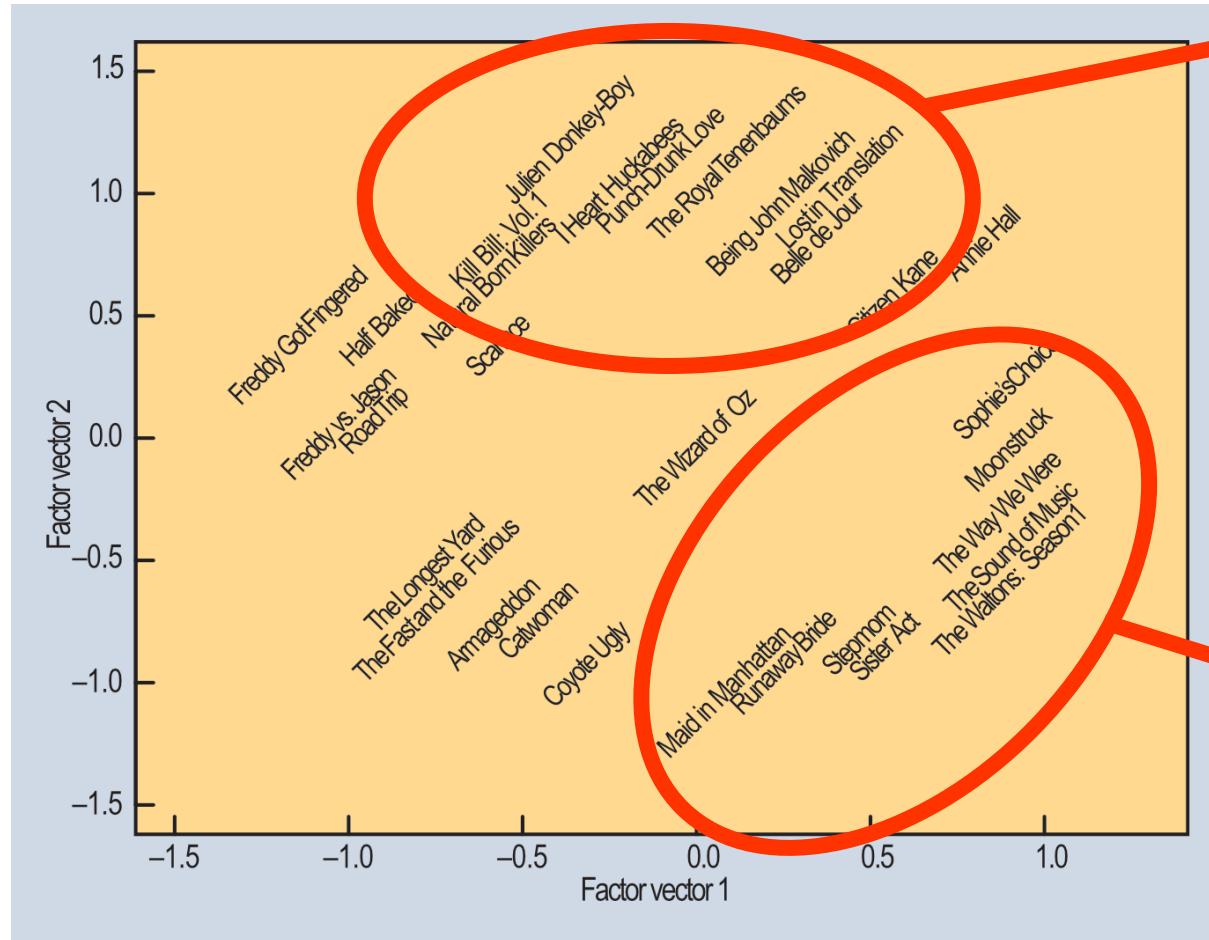
The latent factors embed movies in a space where similar movies are close



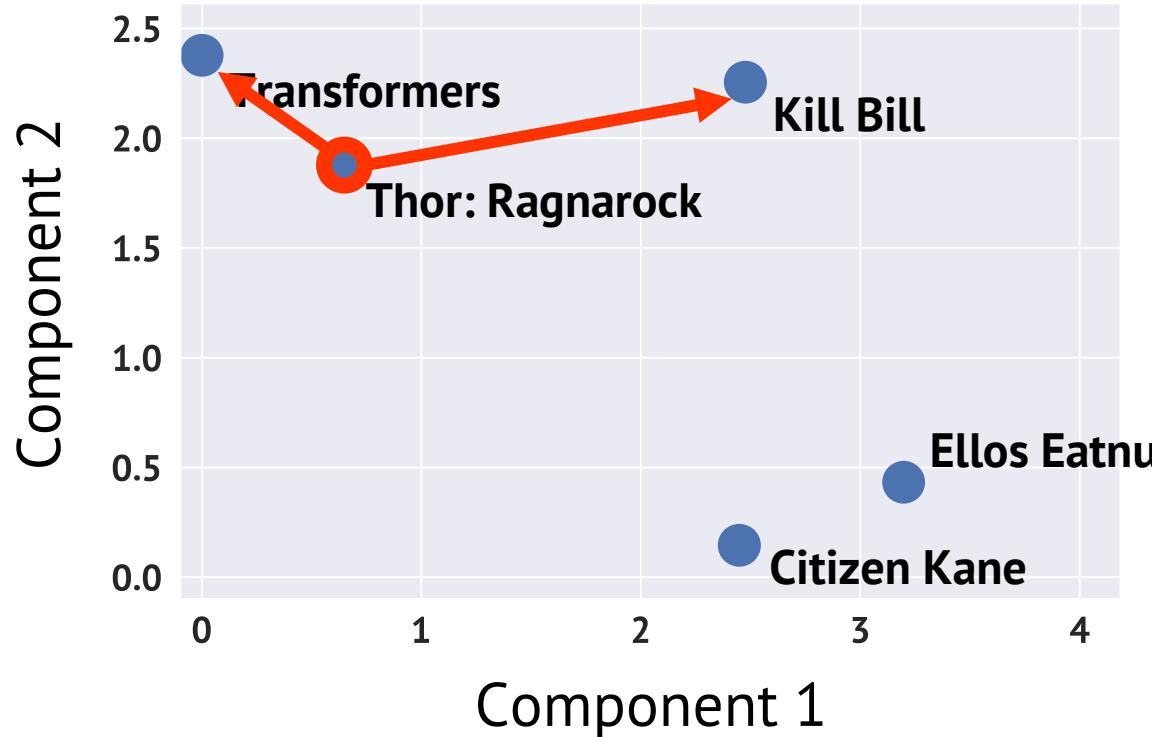
The latent factors embed movies in a space where similar movies are close



The latent factors embed movies in a space where similar movies are close

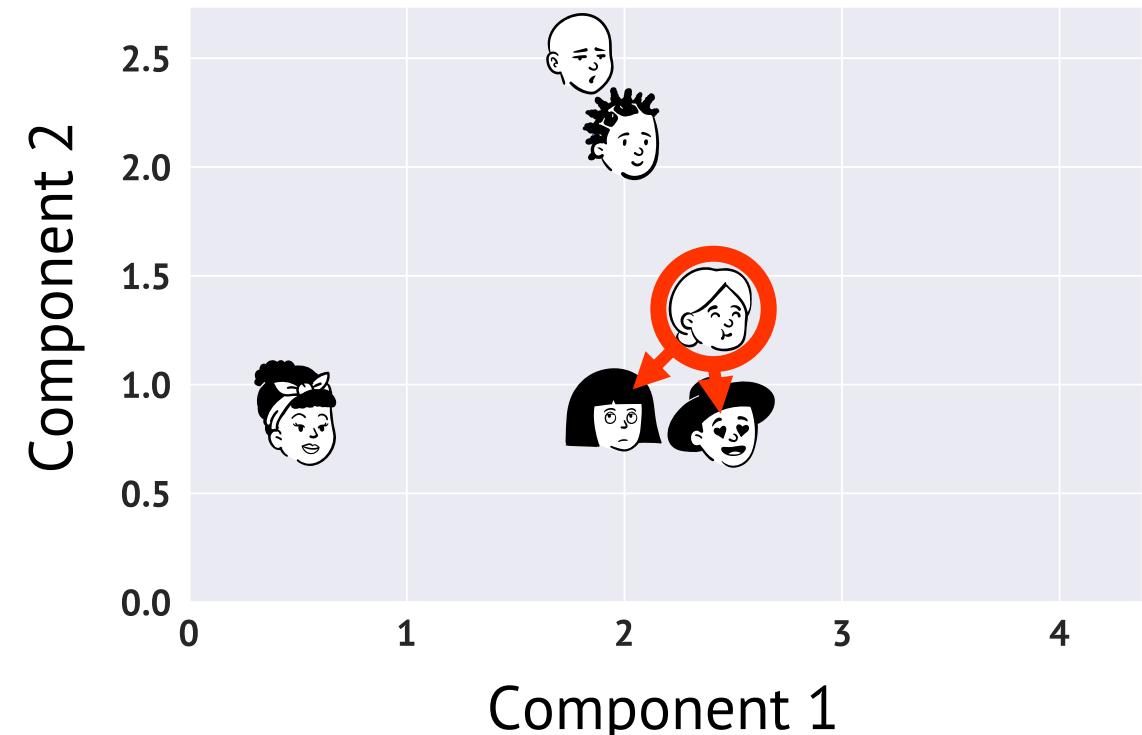
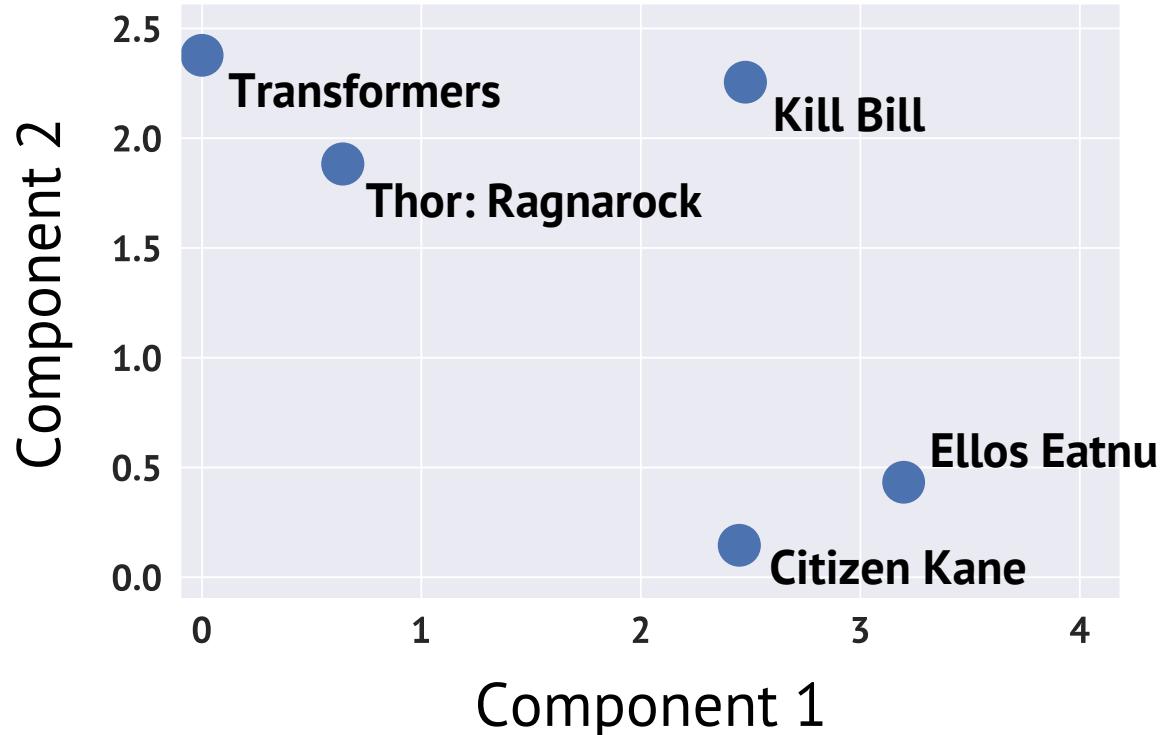


The latent factors embed movies in a space where similar movies are close



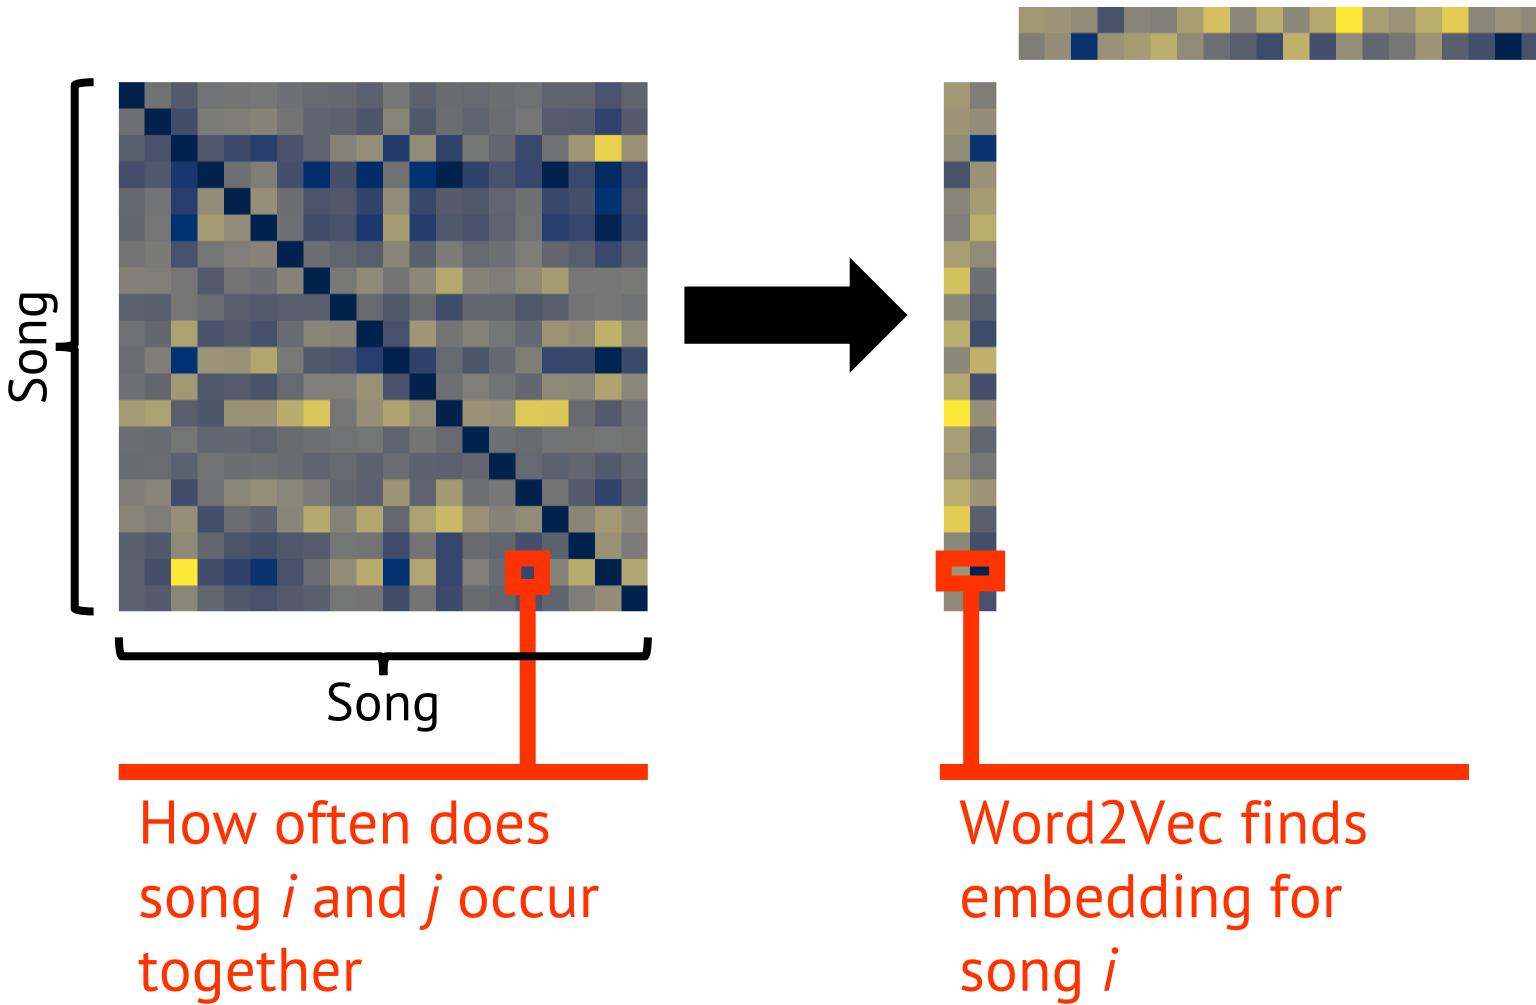
The embeddings can be used to recommend similar movies

The latent factors embed movies in a space where similar movies are close and where similar users are close

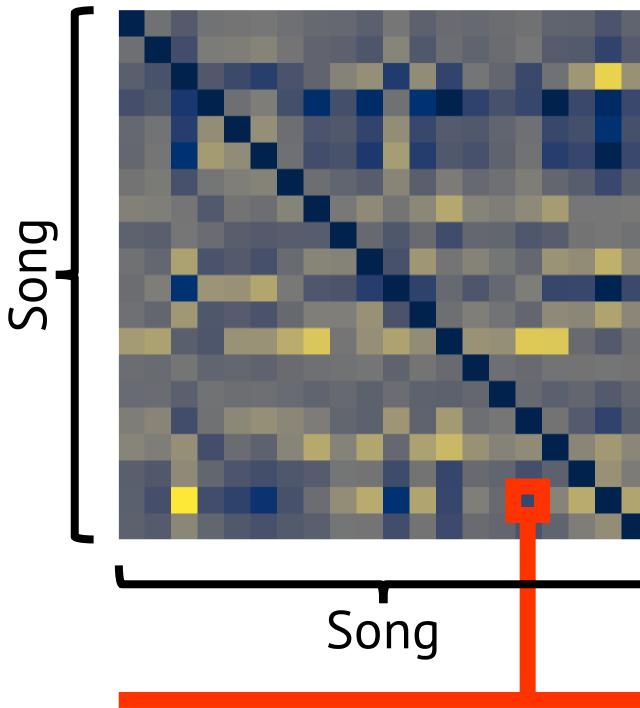


The embeddings can be used to recommend movies liked by similar users

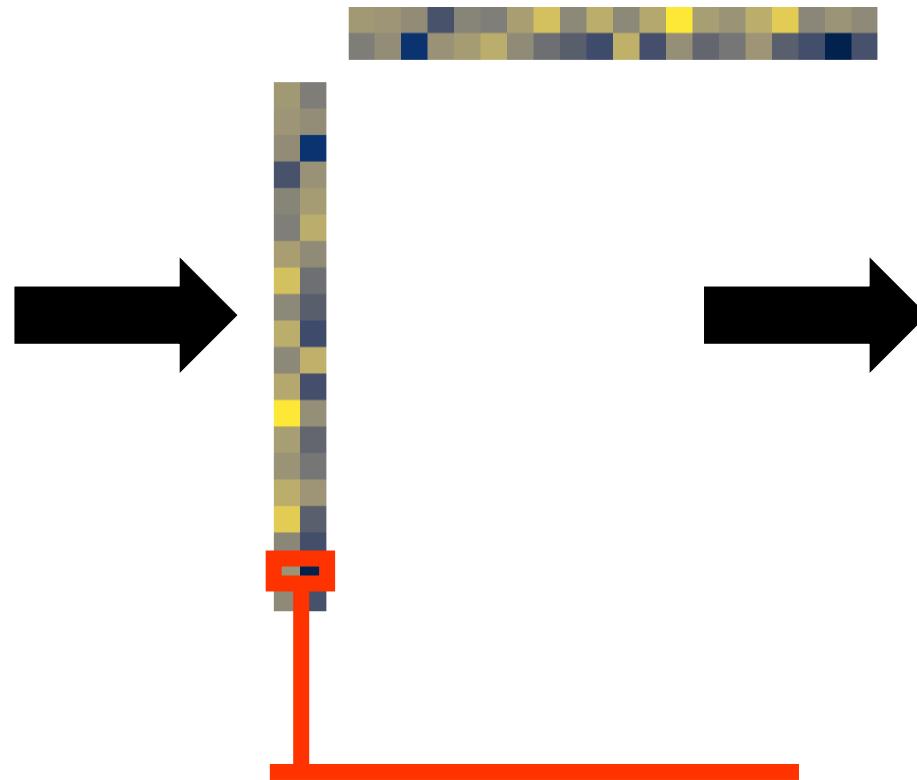
Spotify used nonlinear song embeddings to find similar songs



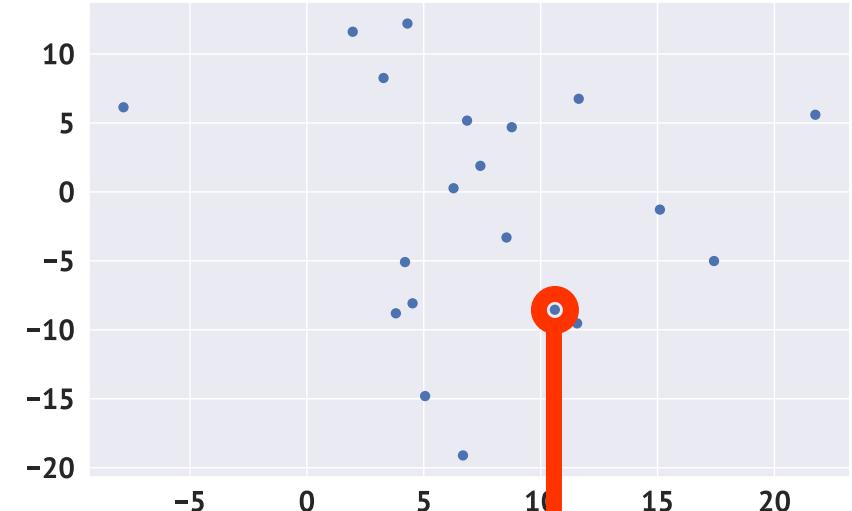
Spotify used nonlinear song embeddings to find similar songs



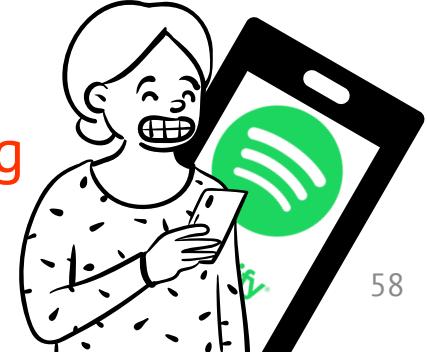
How often does
song i and j occur
together



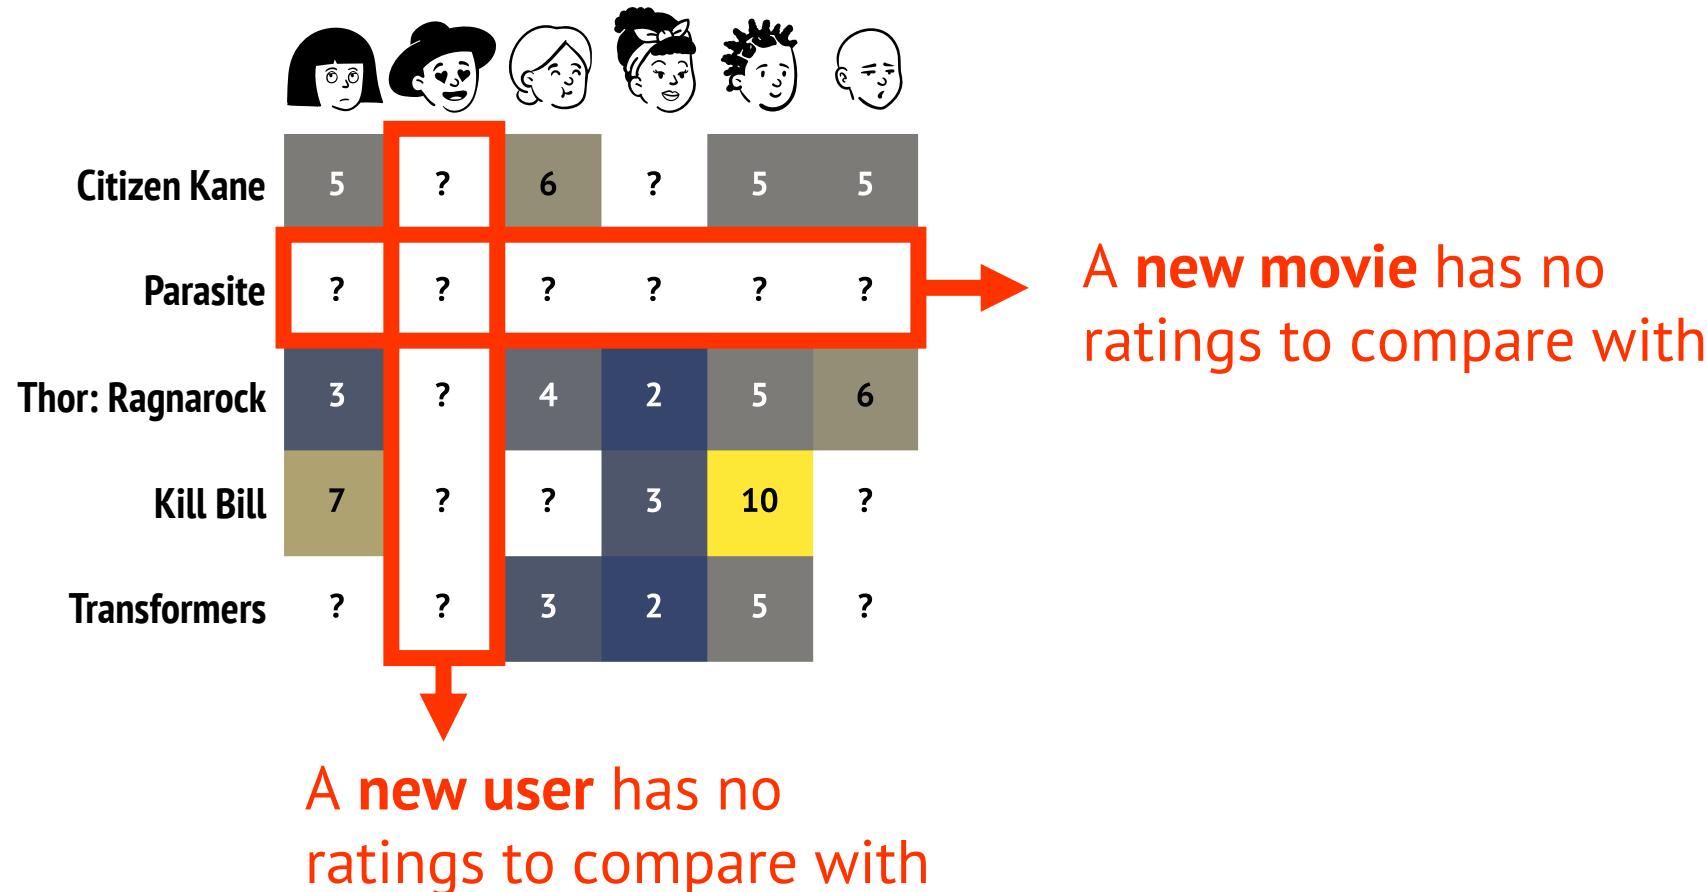
Word2Vec finds
embedding for
song i



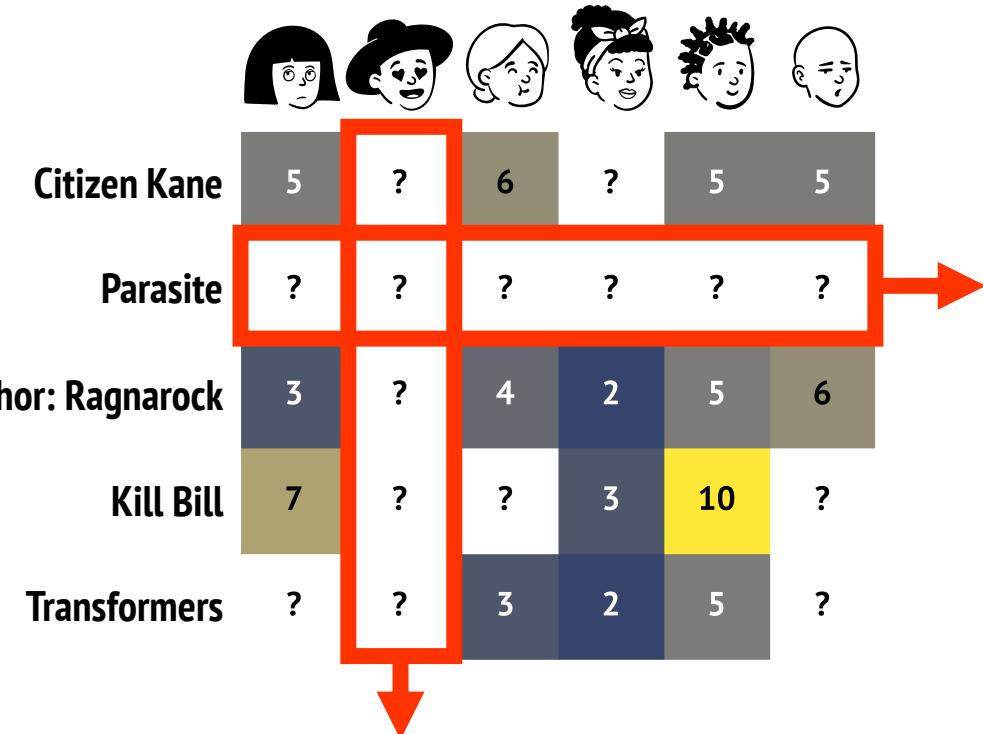
Similar songs
are close in
the embedding
space



Collaborative filtering struggles with the *cold start problem*: it's difficult to provide recommendations for new movies and users



Collaborative filtering struggles with the *cold start problem*: it's difficult to provide recommendations for new movies and users



A new user has no ratings to compare with

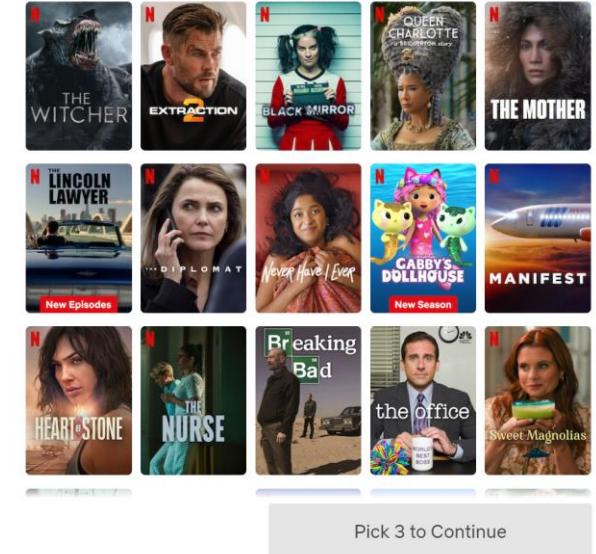
NETFLIX

Step 5 of 5

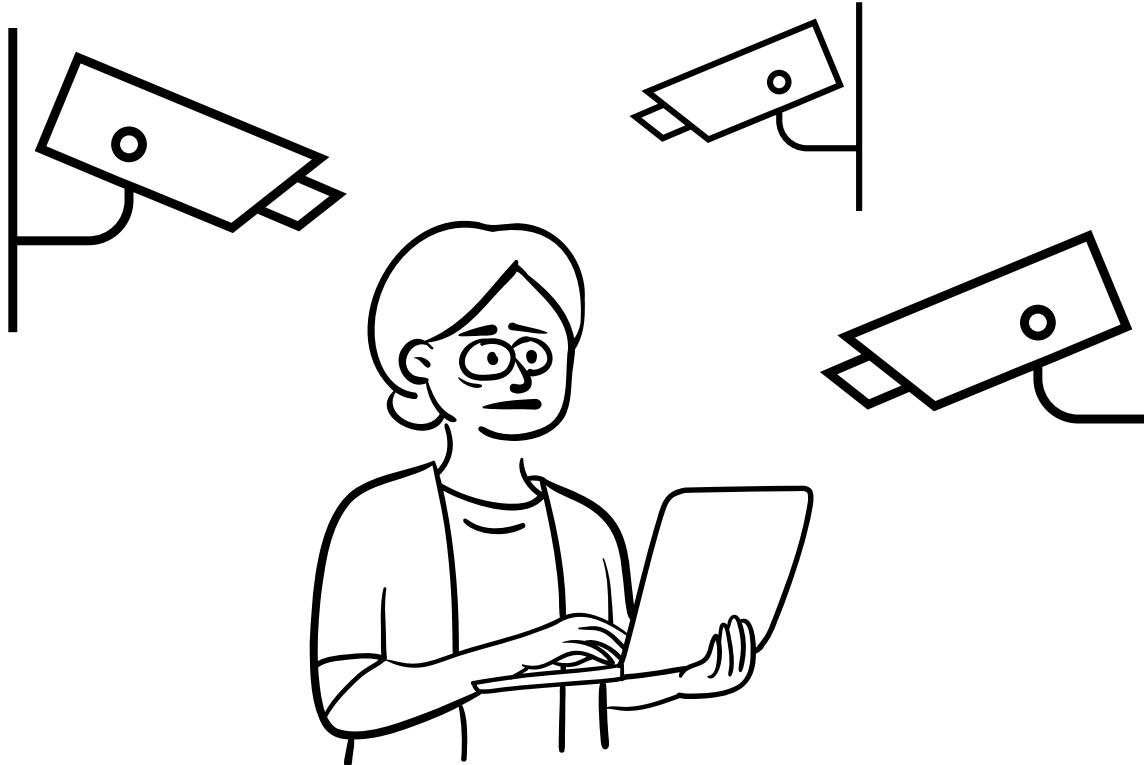
Marie, select 3 you like.

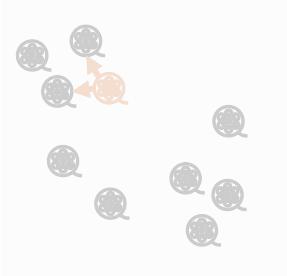
This helps us to find TV programmes and films you'll love. Select the ones you like.

A new movie has no ratings to compare with



Collaborative filtering can be more intrusive than other types of recommender systems





Content-based filtering

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	?	3	10	10
2	2	3	2	5	6

Collaborative filtering



Session-based filtering

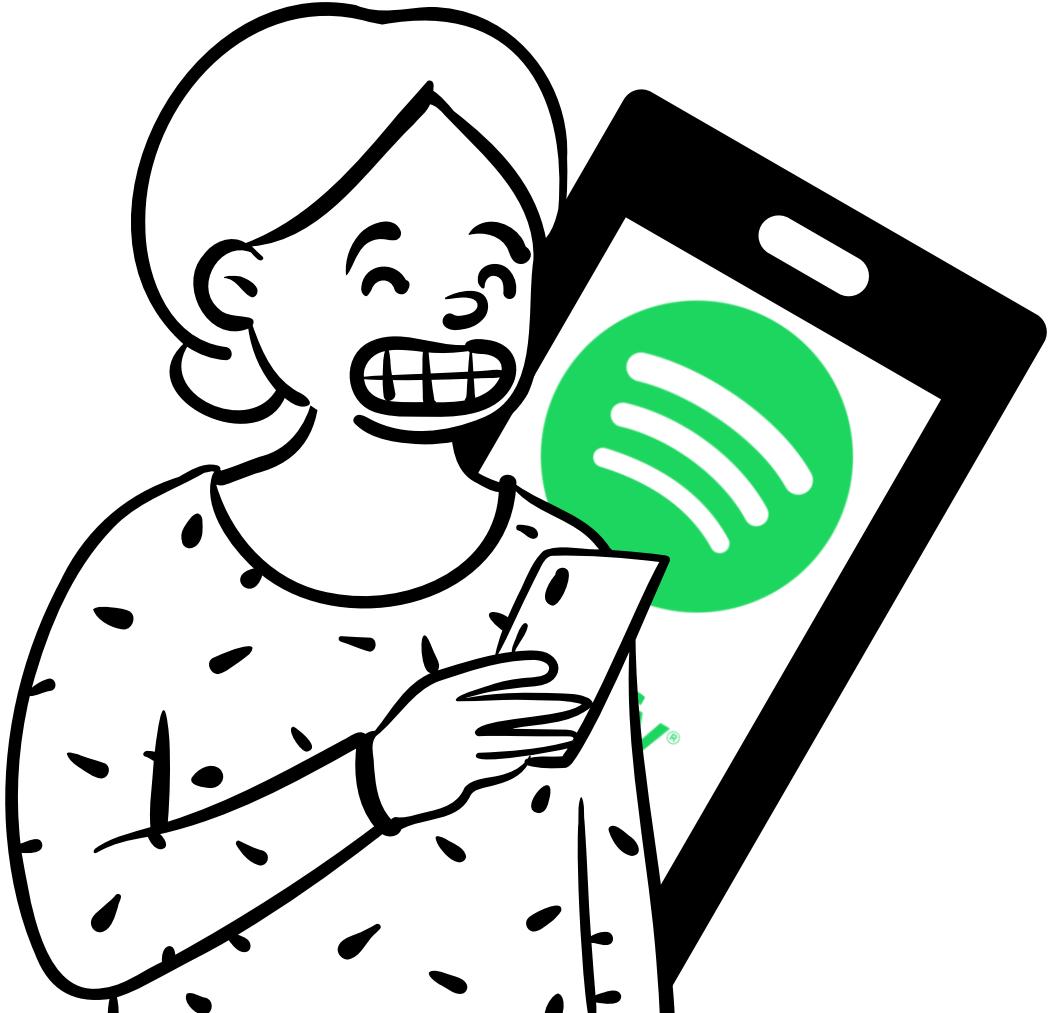


Concluding remarks

Sometimes, we don't want to consider more than the current session when we recommend new items



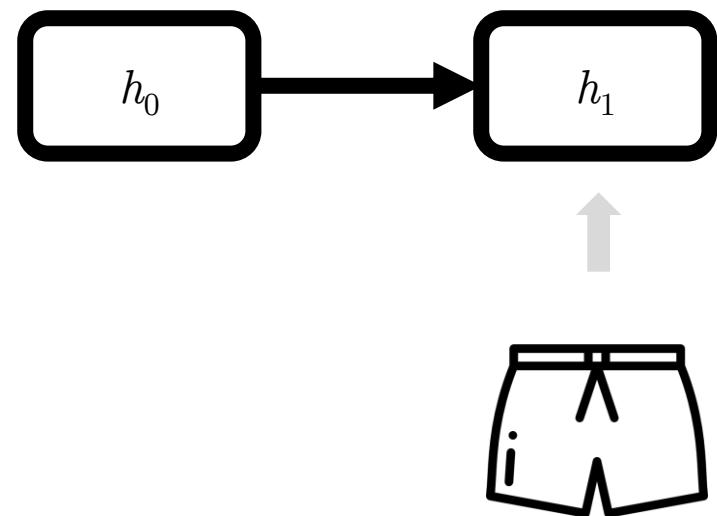
Session-based filtering is beneficial if user behaviour can vary greatly between sessions



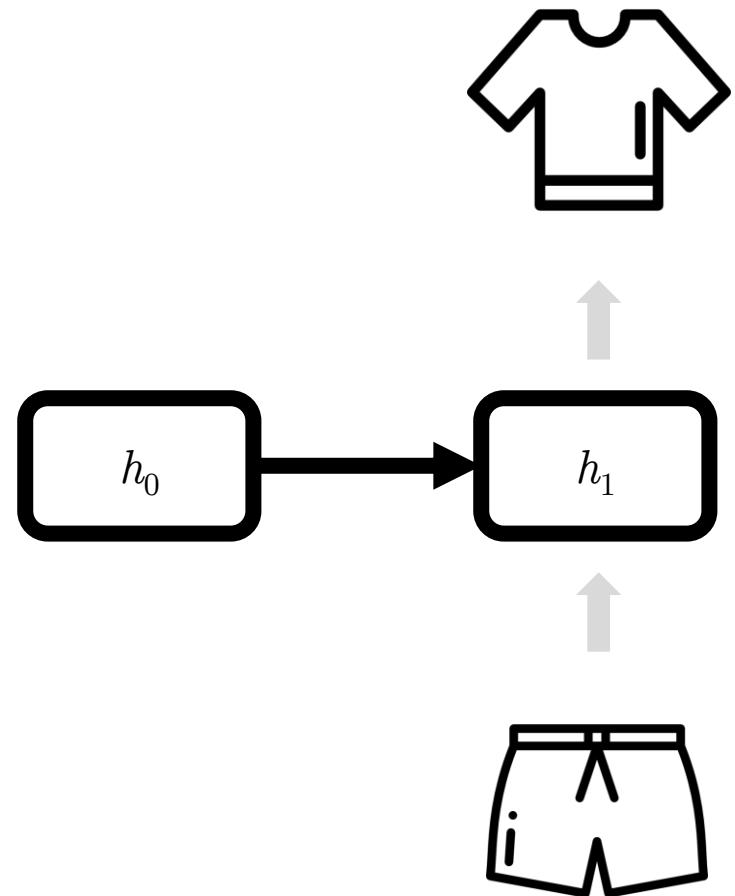
Any collaborative filtering algorithm can be used for session-based filtering

	Session 1	Session 2	Session 3	Session 4	Session 5	Session 6
Citizen Kane	5	6	6	1	5	5
Ellos Eatnu	7	8	8	2	7	7
Thor: Ragnarock	3	3	4	2	5	6
Kill Bill	7	8	9	3	10	10
Transformers	2	2	3	2	5	6

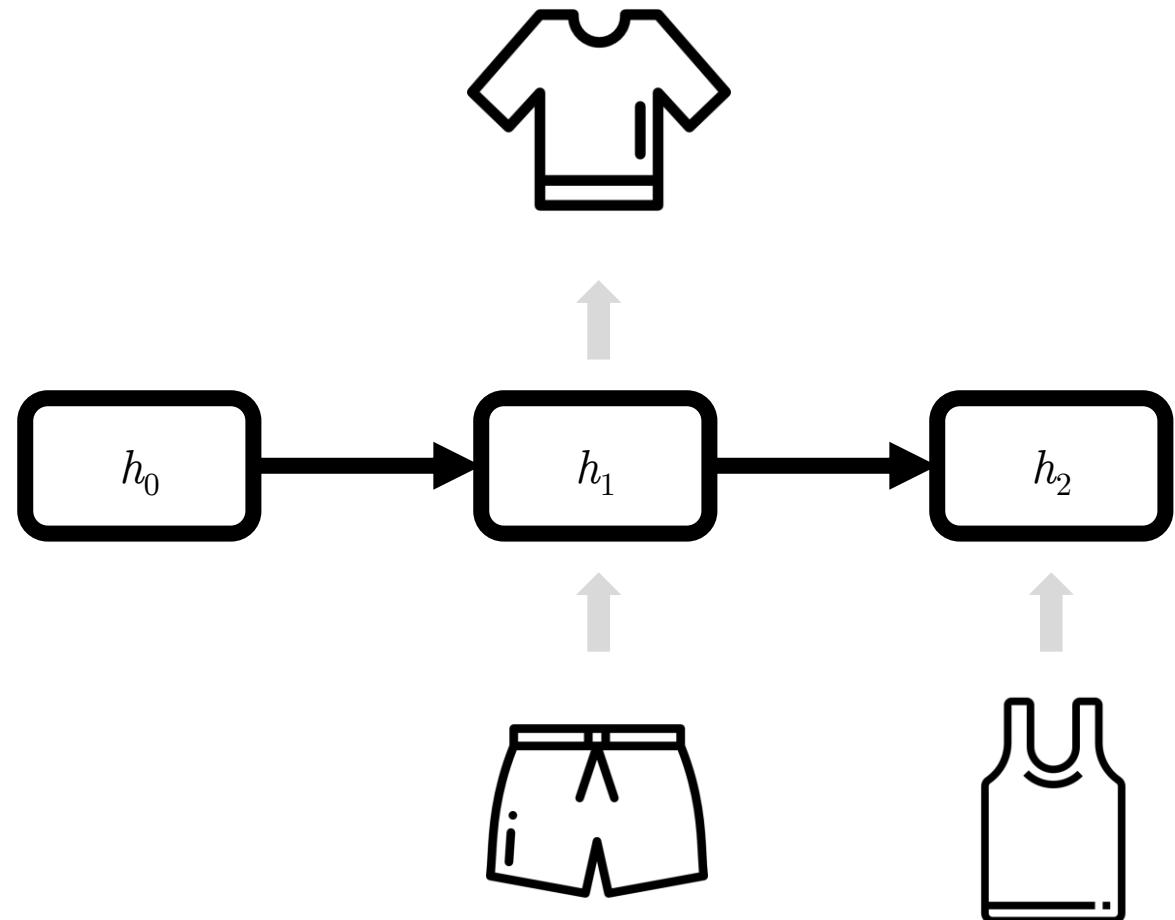
Session-based recommendation system makes them suitable for recurrent neural networks



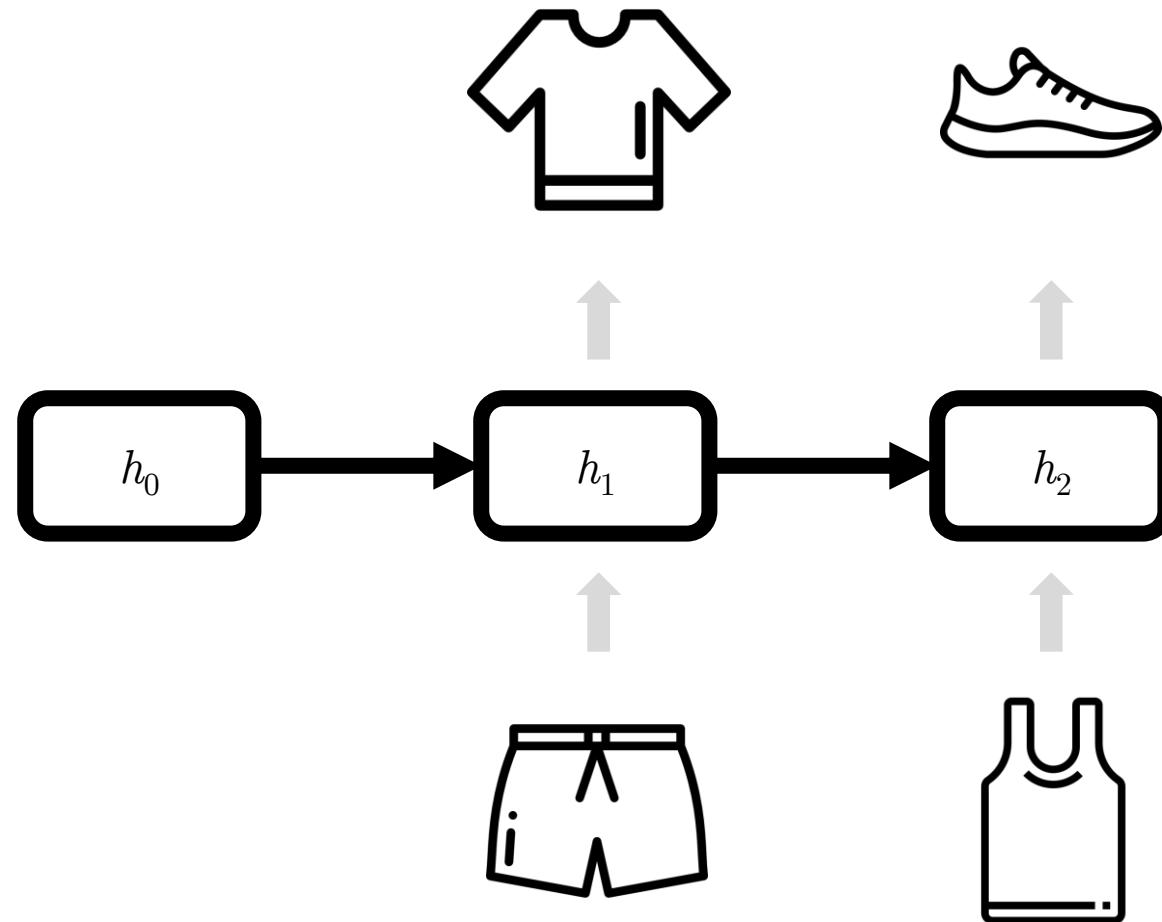
Session-based recommendation system makes them suitable for recurrent neural networks



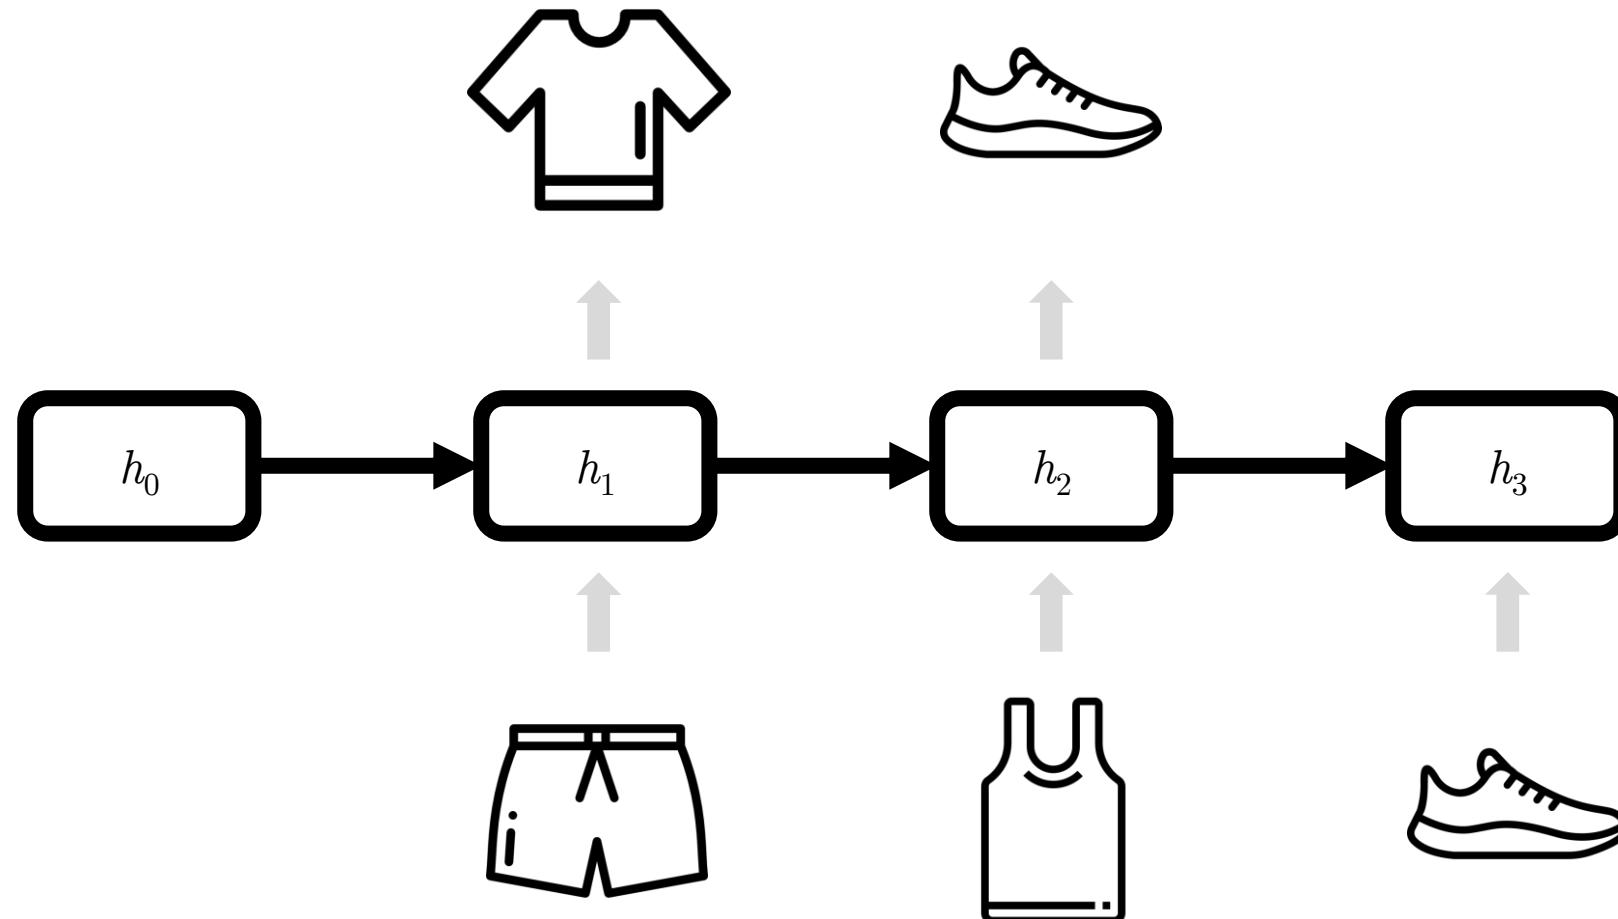
Session-based recommendation system makes them suitable for recurrent neural networks



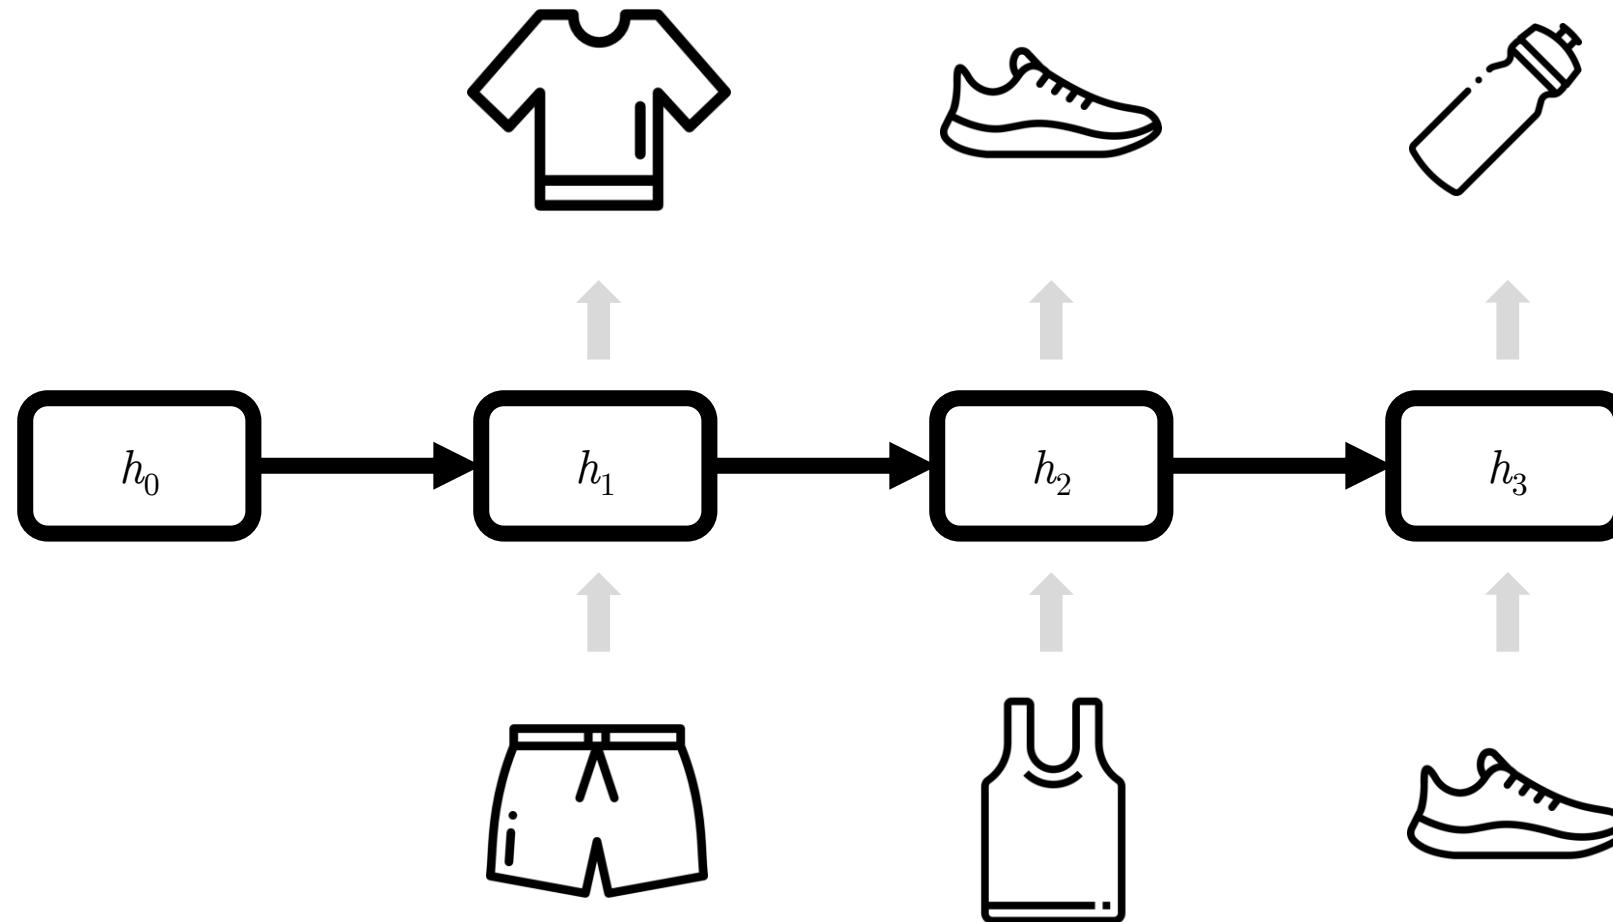
Session-based recommendation system makes them suitable for recurrent neural networks



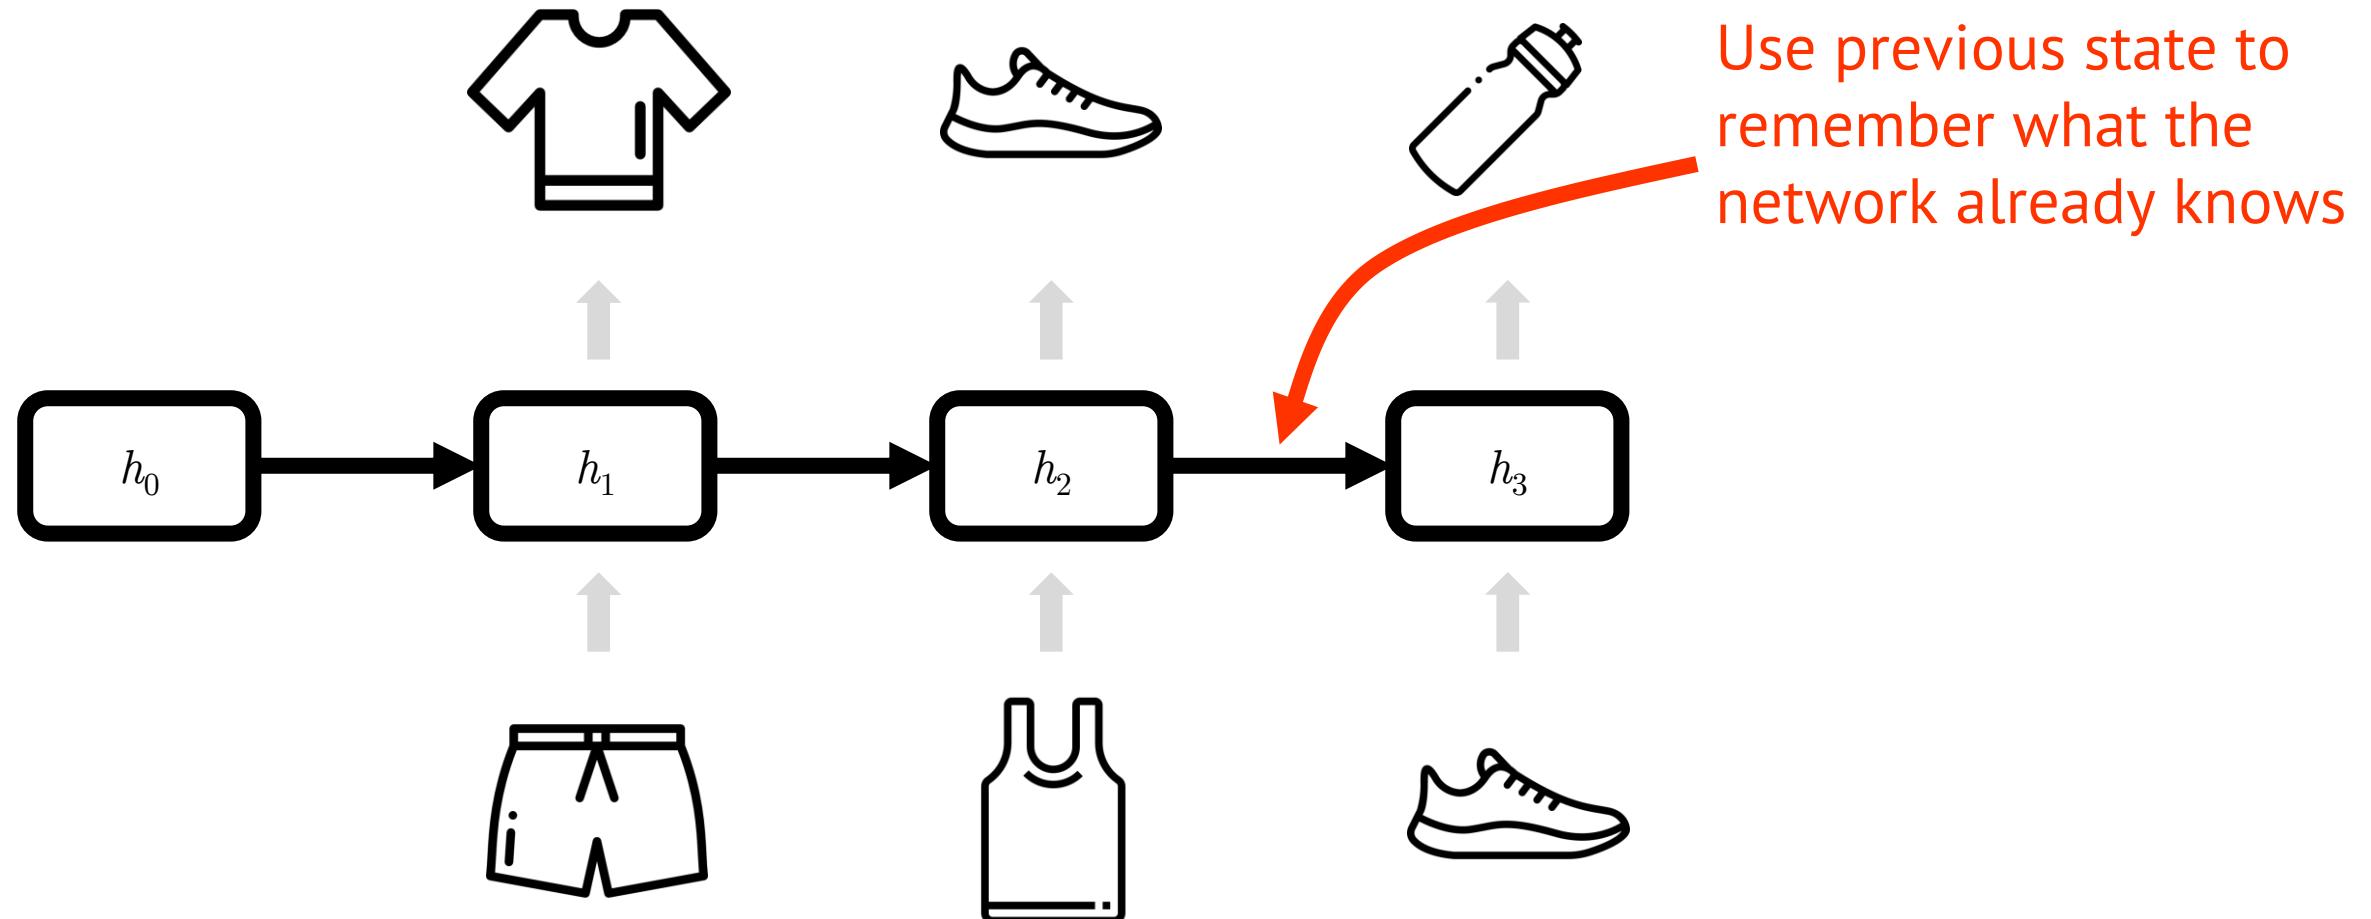
Session-based recommendation system makes them suitable for recurrent neural networks



Session-based recommendation system makes them suitable for recurrent neural networks

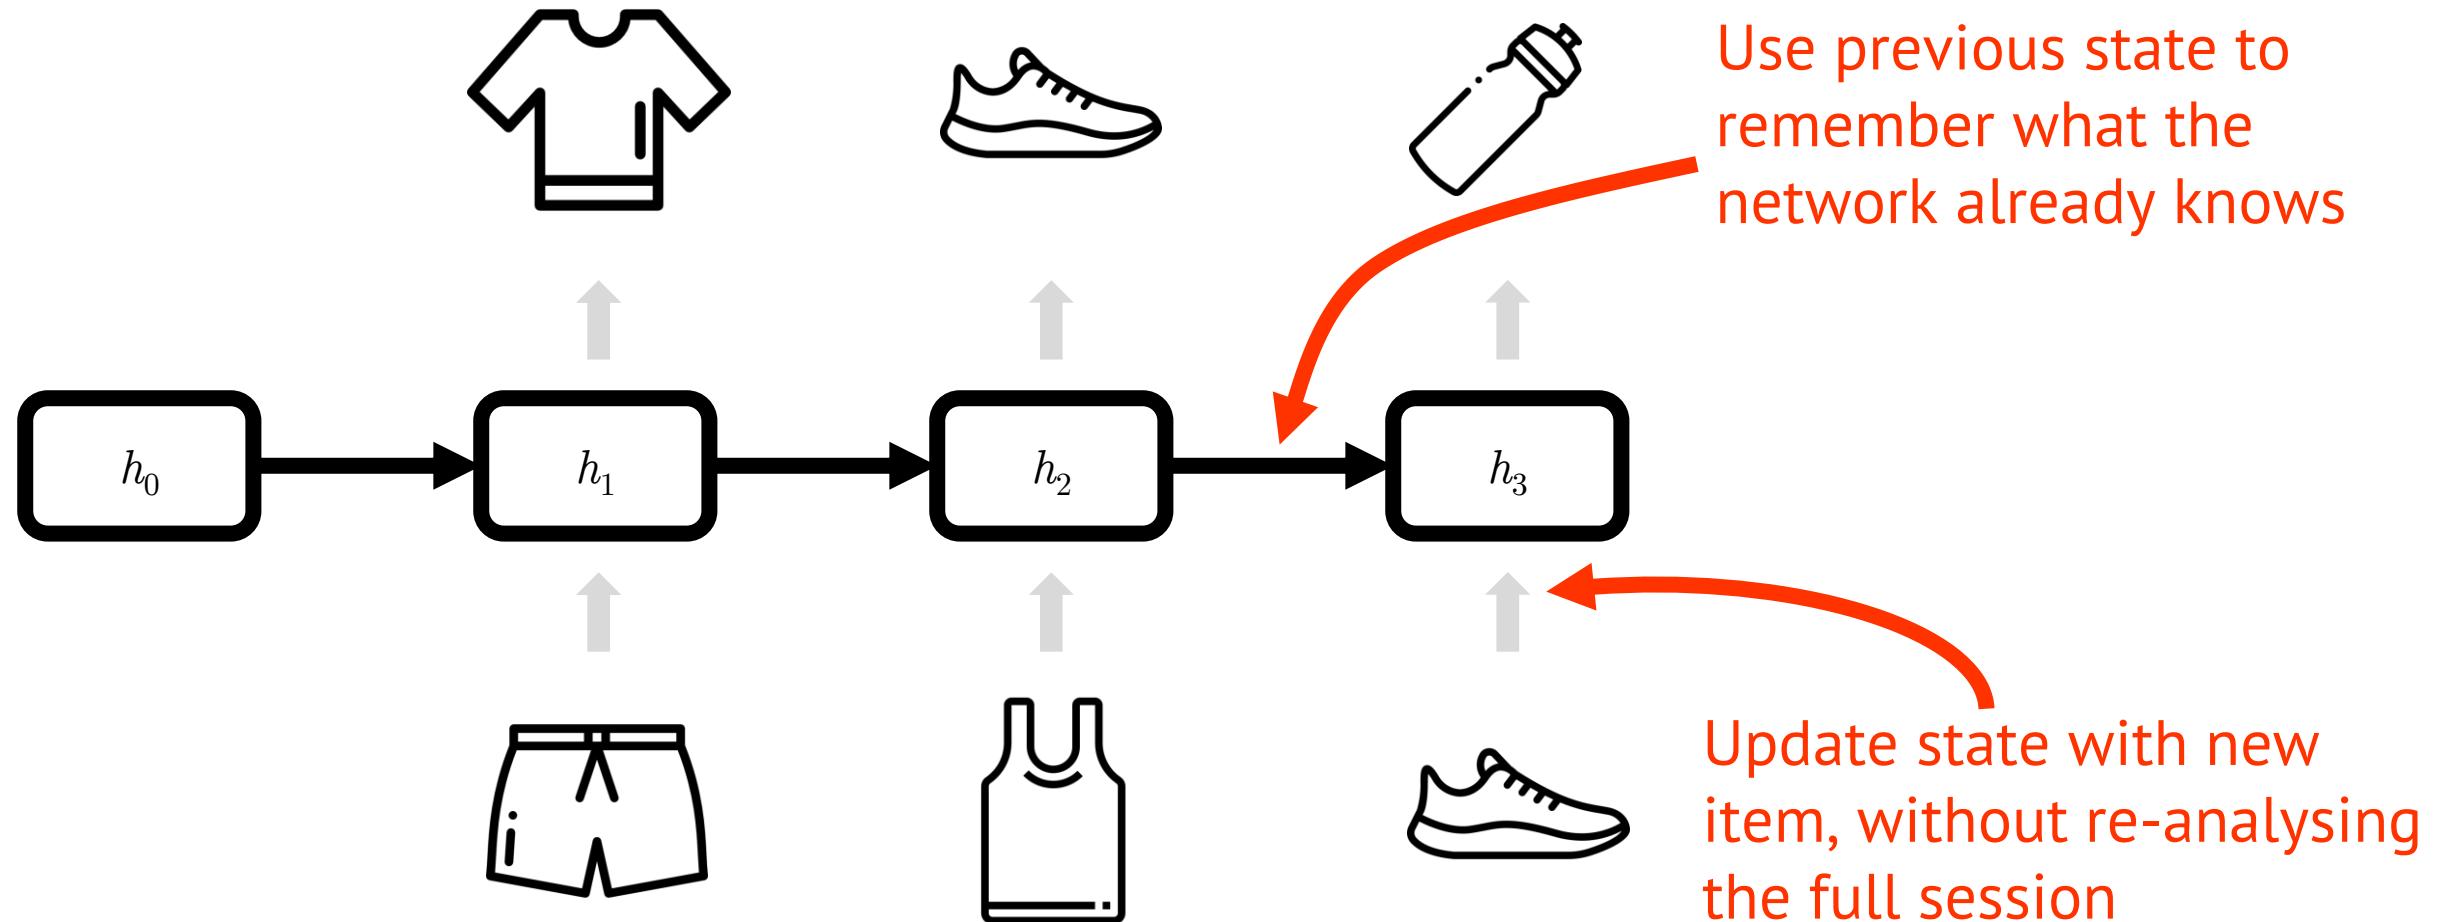


Session-based recommendation system makes them suitable for recurrent neural networks



Use previous state to remember what the network already knows

Session-based recommendation system makes them suitable for recurrent neural networks



Spotify used a recurrent neural network with additional contextual information to give just-in-time recommendations during a session



Contextual and Sequential User Embeddings for Large-Scale Music Recommendation

Casper Hansen^{*}
University of Copenhagen
c.hansen@di.ku.dk

Christian Hansen^{*}
University of Copenhagen
chrh@di.ku.dk

Brian Brost
Spotify
brianbrost@spotify.com

Federico Tomasi
Spotify
federicot@spotify.com

Rishabh Mehrotra
Spotify
rishabhlm@spotify.com

Mounia Lalmas
Spotify
mounia@acm.org

Lucas Maystre
Spotify
lucasmg@spotify.com

Lucas Maystre
Spotify
lucasmg@spotify.com

Rishabh Mehrotra
Spotify
rishabhlm@spotify.com

Mounia Lalmas
Spotify
mounia@acm.org

ACM Reference Format:

Casper Hansen, Christian Hansen, Lucas Maystre, Rishabh Mehrotra, Brian

Brost, Federico Tomasi, and Mounia Lalmas. 2020. Contextual and Sequential

User Embeddings for Large-Scale Music Recommendation. In *Fourteenth*

ACM Conference on Recommender Systems (RecSys '20), September 22–26,

2020, Virtual Event, Brazil. ACM, New York, NY, USA, 16 pages. <https://doi.org/10.1145/3383313.3412248>

INTRODUCTION

Recommender systems are essential for providing an engaging user experience and for helping users navigating the vast amounts of content available in online services. Successful recommender systems have to accurately model each user's individual preferences such that the most relevant content can be presented to the user. In this work, we consider online music streaming services, which have become increasingly popular in the past decade. By letting users access millions of tracks at the click of a button, they are contributing to democratizing access to music. However, in contrast to other well studied domains (such as recommending books, movies or clothes), music recommender systems face distinctive challenges [30]. Tracks are short, and therefore often consumed together with other tracks; we refer to such a set of tracks listened to in short succession as a session. A given session often contains the sequence of sessions captured by a user's recent consumption history [1], suggesting that the sequence of sessions captures essential information about the user's changing preferences. Additionally, the relevance of tracks is highly contextual, and preferences depend, among others, on the time of the day and the current season [24]. We seek to embrace these distinctive characteristics to produce a better, more accurate model of user preferences. We focus on the following problem: for a given user, we are interested in predicting, at the beginning of a session, which tracks the user will listen to during the session.

*this work was done while the first two authors were at Spotify.
Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies that reflect the full extent of the work are not sold in whole or in part. Copyright © 2020, Christian Hansen, Casper Hansen, Lucas Maystre, Rishabh Mehrotra, Brian Brost, Federico Tomasi, and Mounia Lalmas. This work is licensed under a Creative Commons Attribution Non-Commercial-ShareAlike 4.0 International License. For details, see: <https://creativecommons.org/licenses/by-nd/4.0/>. Publication rights licensed to ACM.
© 2020 Copyright held by the owner/author(s). Publication rights licensed to ACM.
<https://doi.org/10.1145/3383313.3412248>

We begin our investigation by exploring a dataset from an online music streaming service, containing detailed information about the tracks we define context as the time of the day (morning, afternoon,

Transformer-based architectures, similar to those used by large language models are also very suitable for session-based recommendation

I

love

fruit,

can

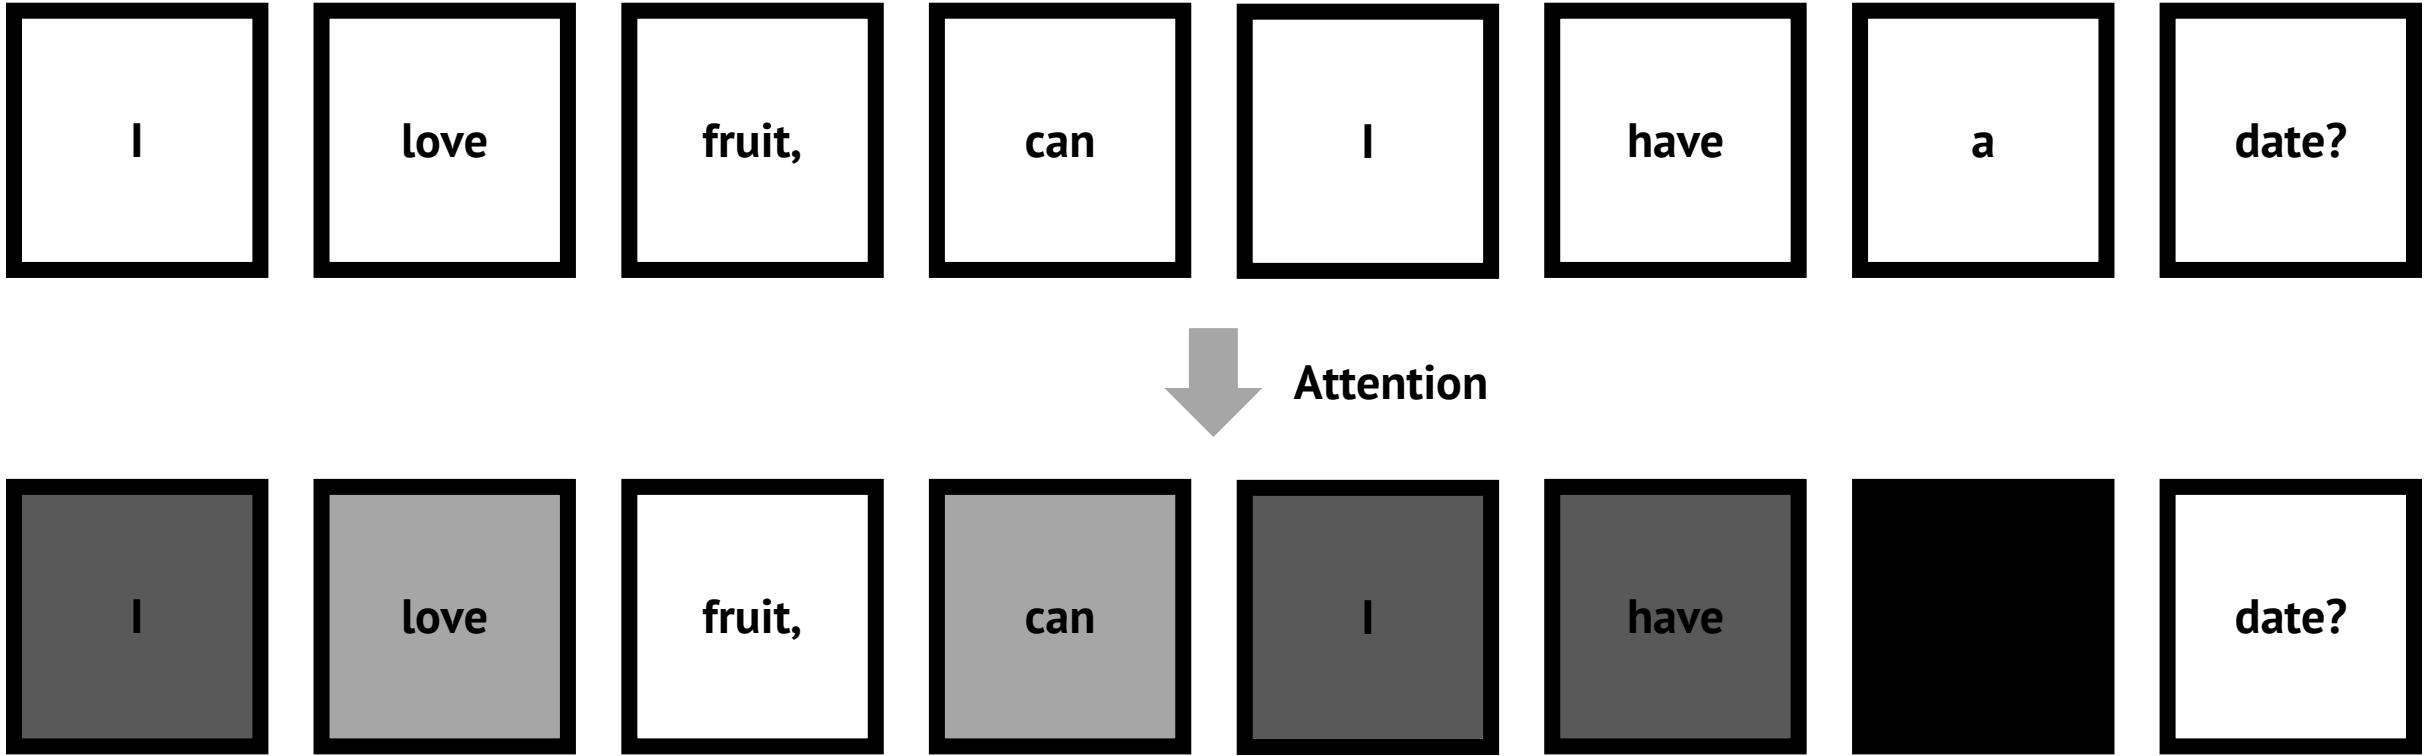
I

have

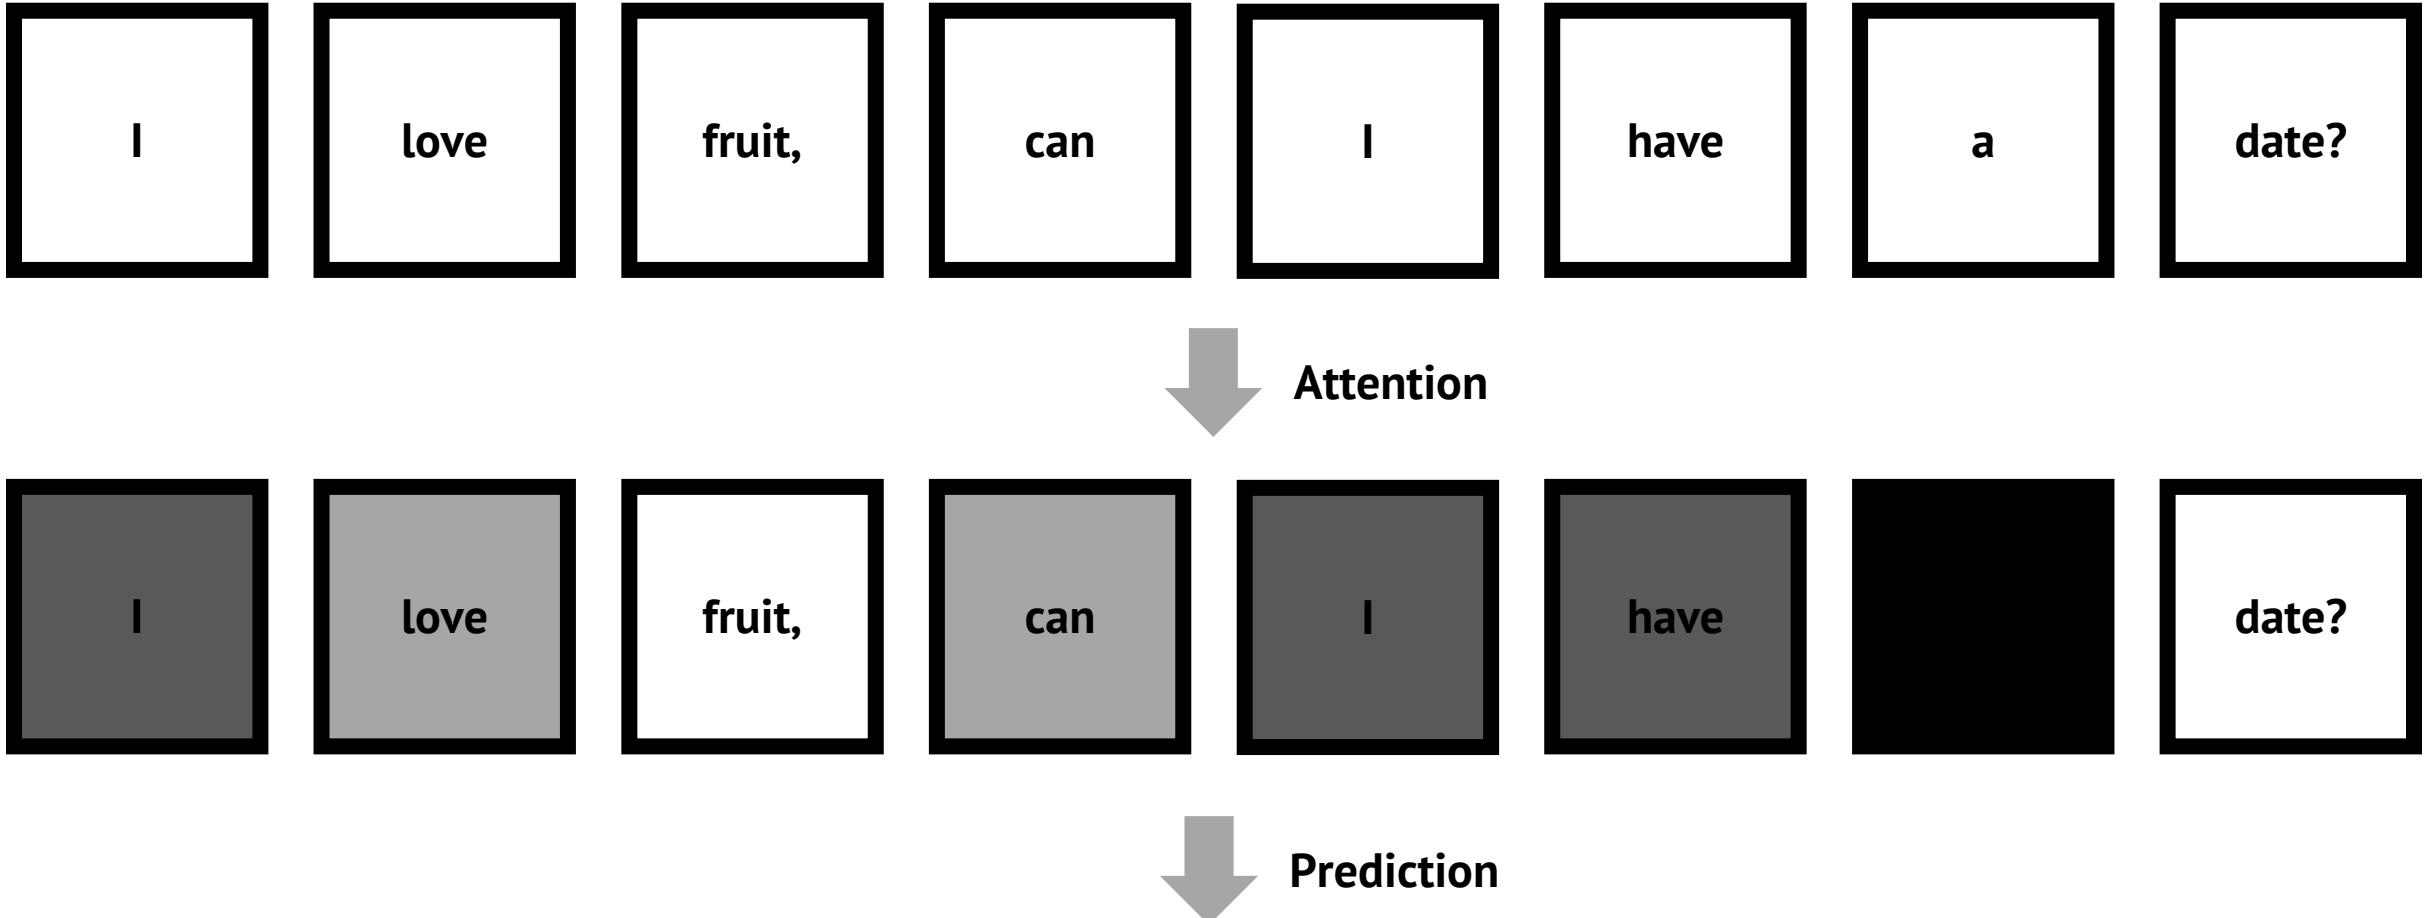
a

date?

Transformer-based architectures, similar to those used by large language models are also very suitable for session-based recommendation

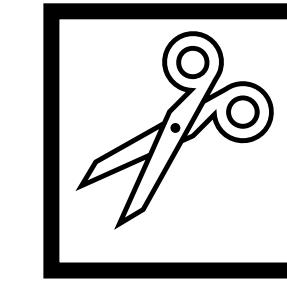
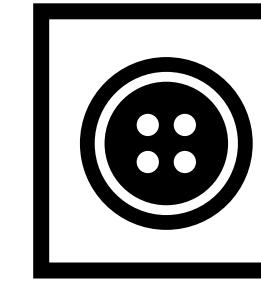
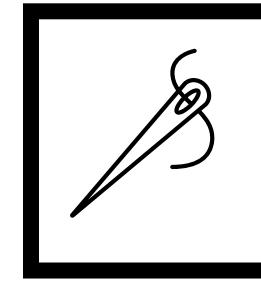
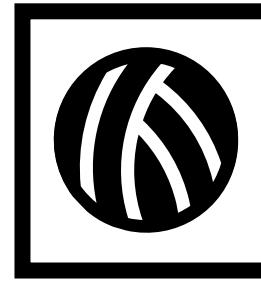
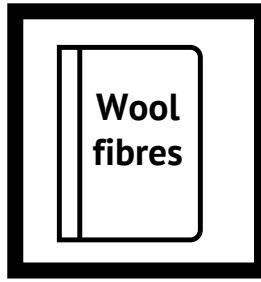
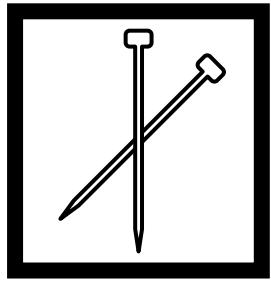


Transformer-based architectures, similar to those used by large language models are also very suitable for session-based recommendation

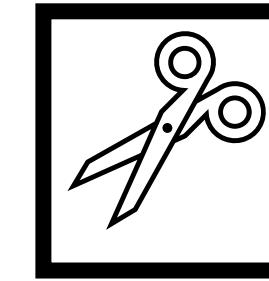
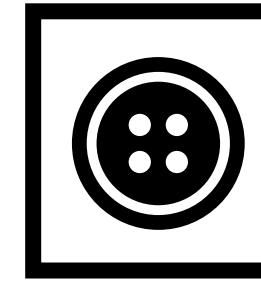
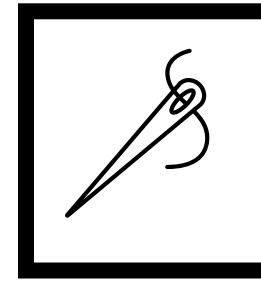
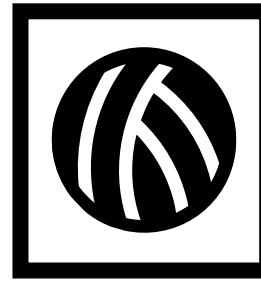
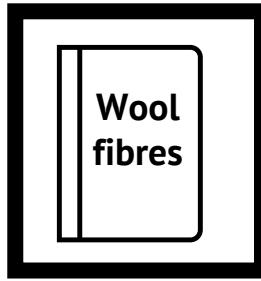
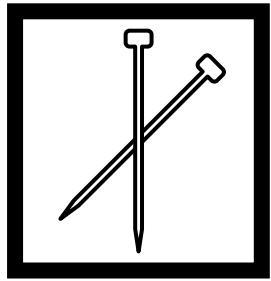


Of course, enjoying fruit is a great way to satisfy your cravings for something sweet and healthy!

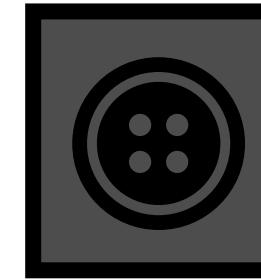
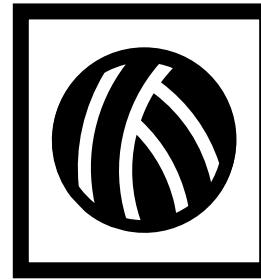
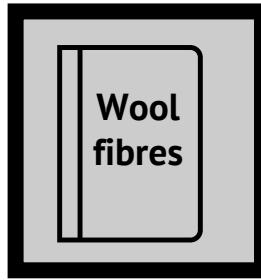
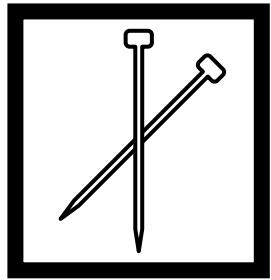
Transformer-based architectures, similar to those used by large language models are also very suitable for session-based recommendation



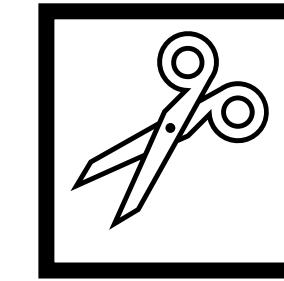
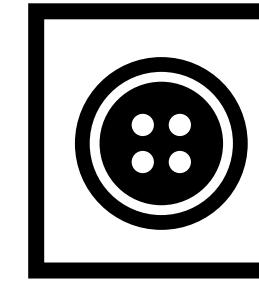
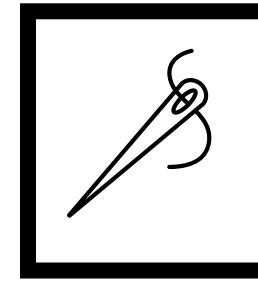
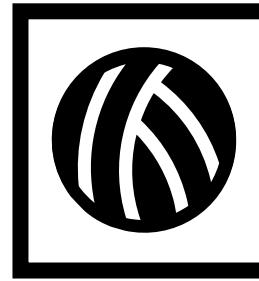
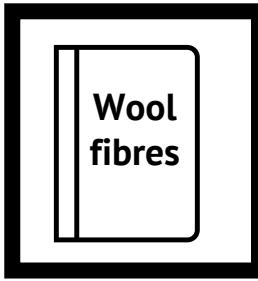
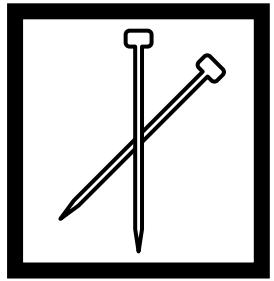
Transformer-based architectures, similar to those used by large language models are also very suitable for session-based recommendation



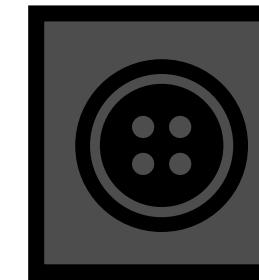
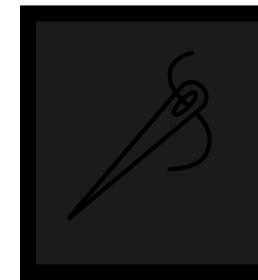
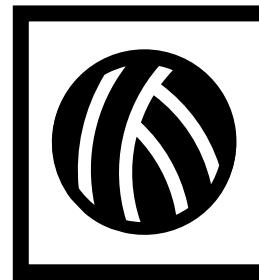
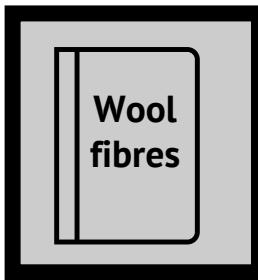
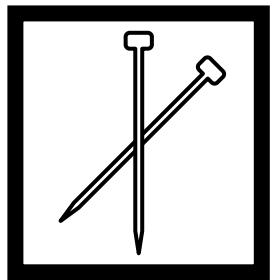
↓ Attention



Transformer-based architectures, similar to those used by large language models are also very suitable for session-based recommendation



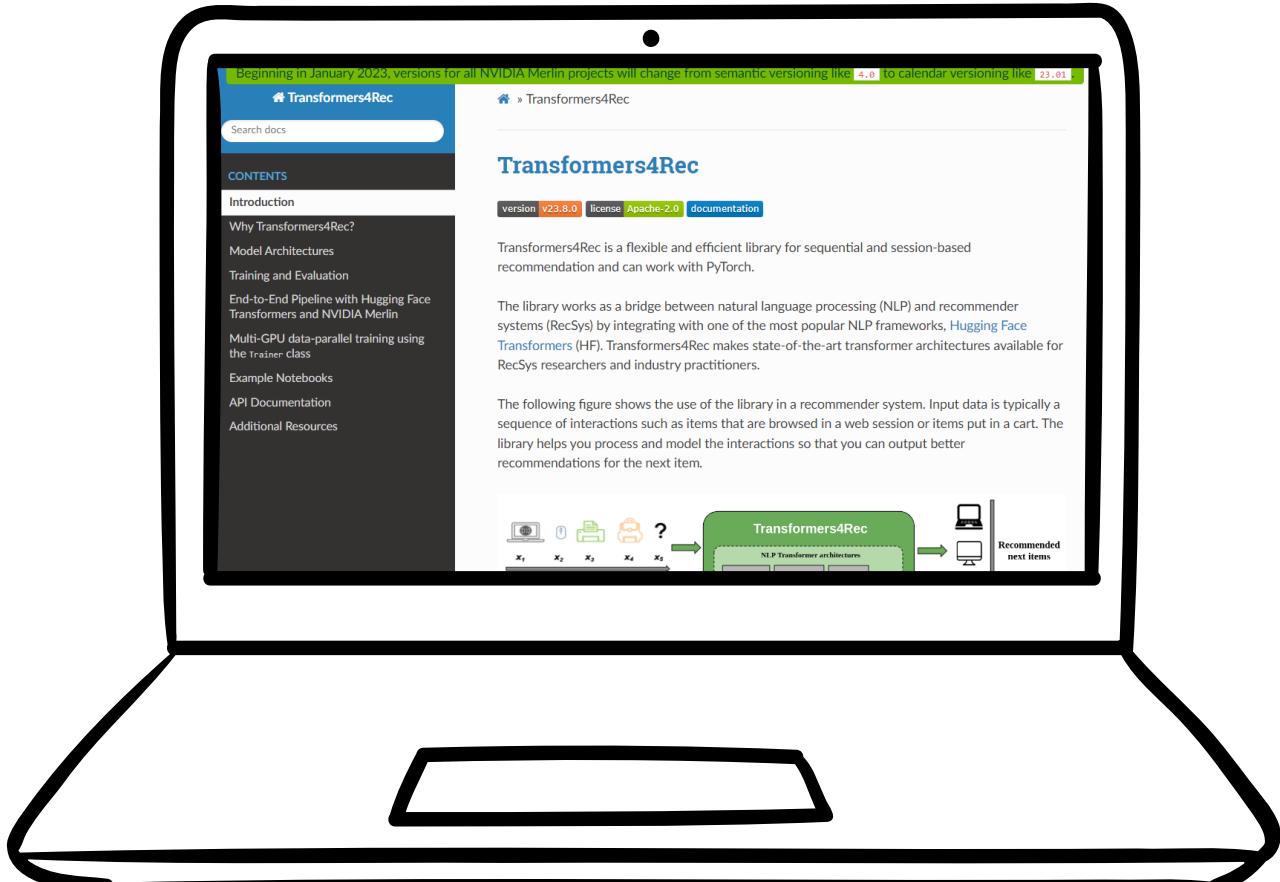
Attention

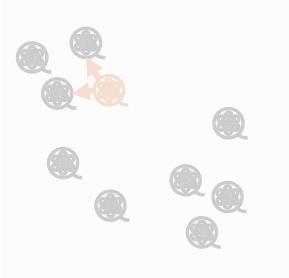


Prediction



Transformers4Rec by Nvidia makes it easy to use transformers-based NLP algorithms for session-based recommendation systems

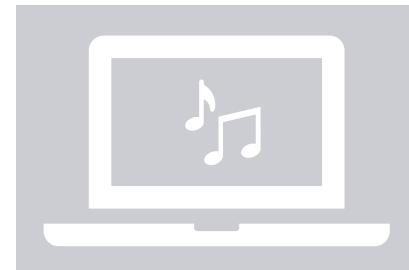




Content-based filtering

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	?	3	10	10
2	2	3	2	5	6

Collaborative filtering



Session-based filtering



Concluding remarks

It's common to combine these paradigms in a hybrid approach

Session embedding

User embedding

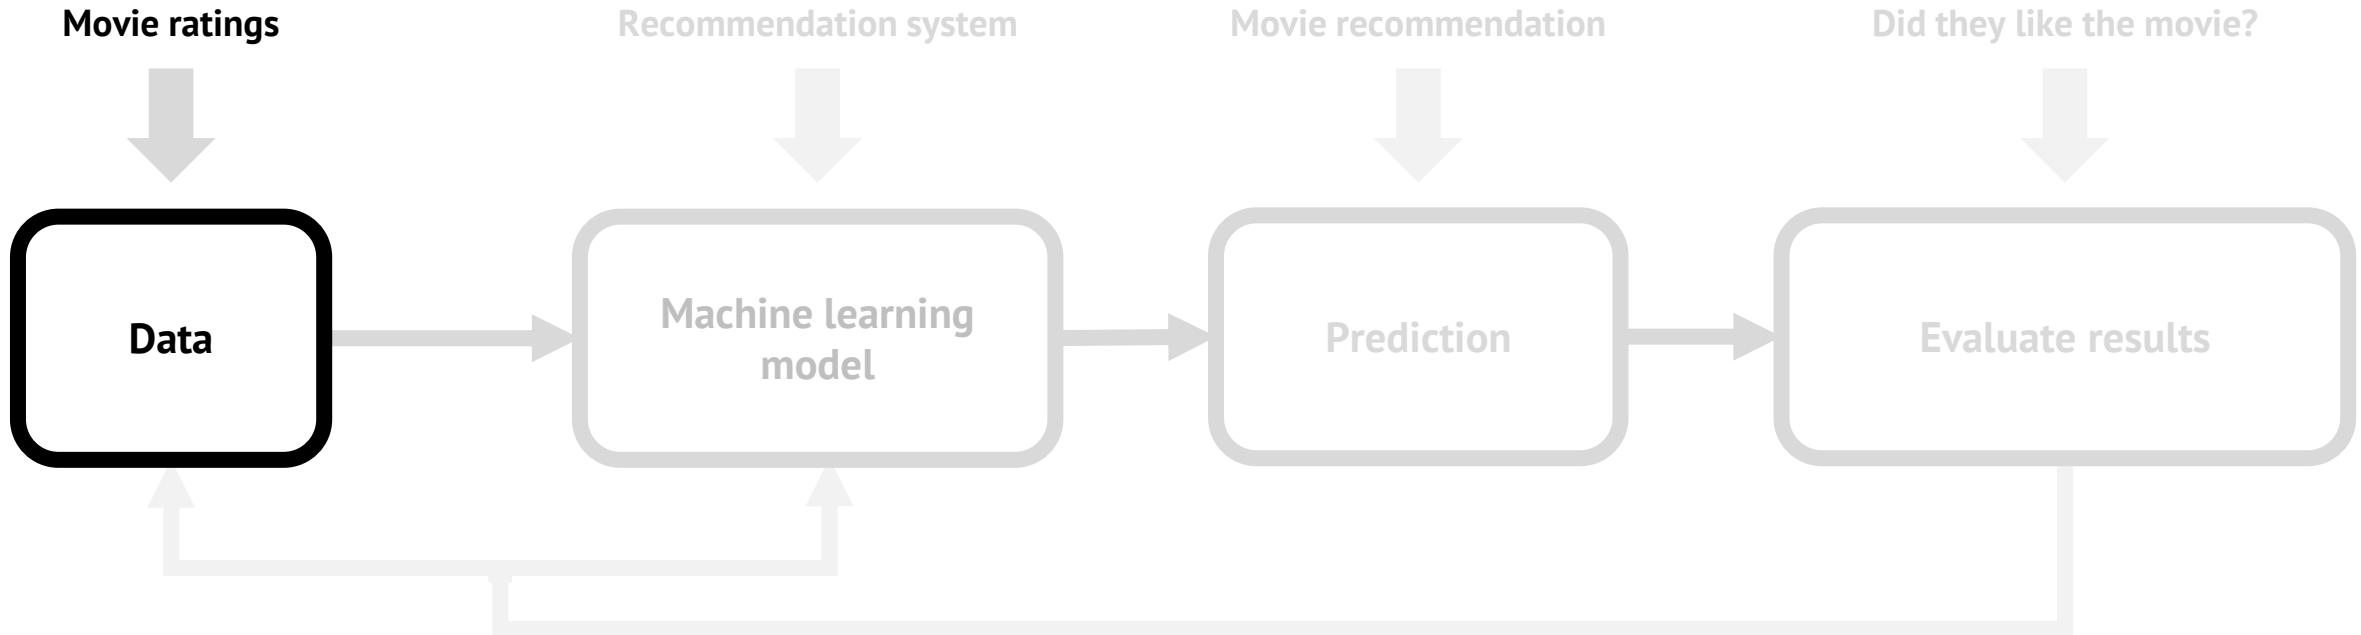
Song embedding

Contextual features

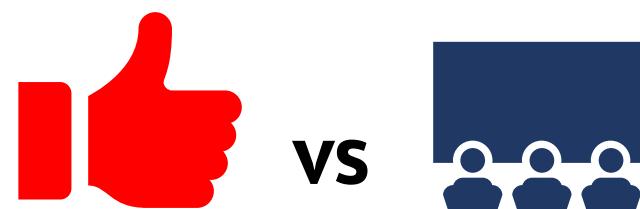
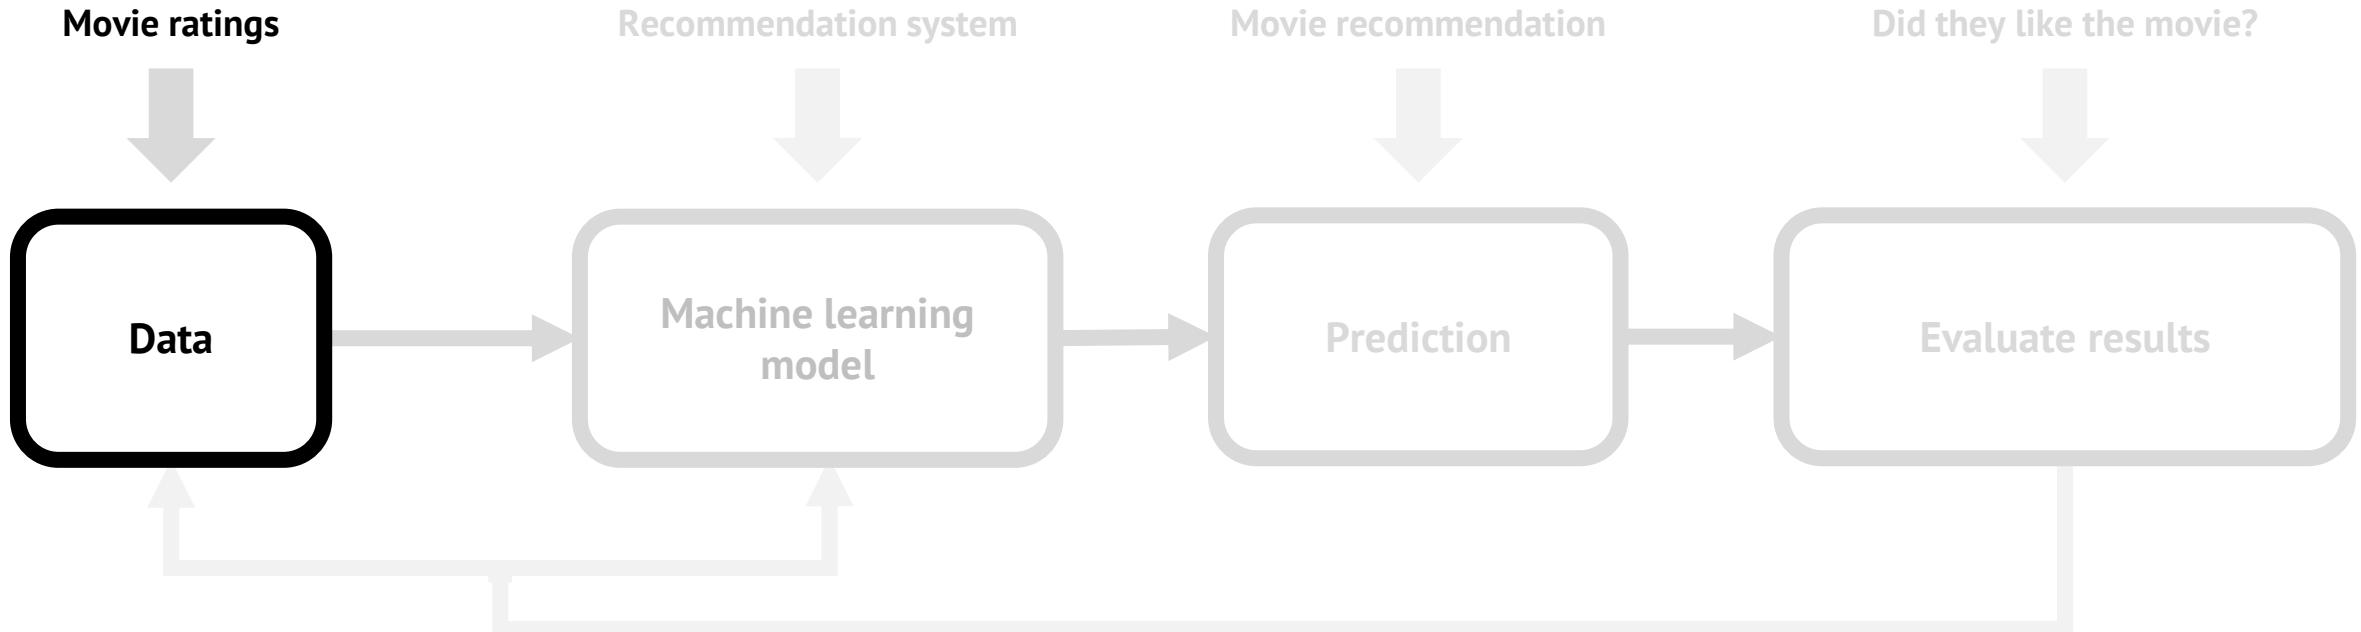


Song recommendation

What data we collect shapes what type of recommendation we perform

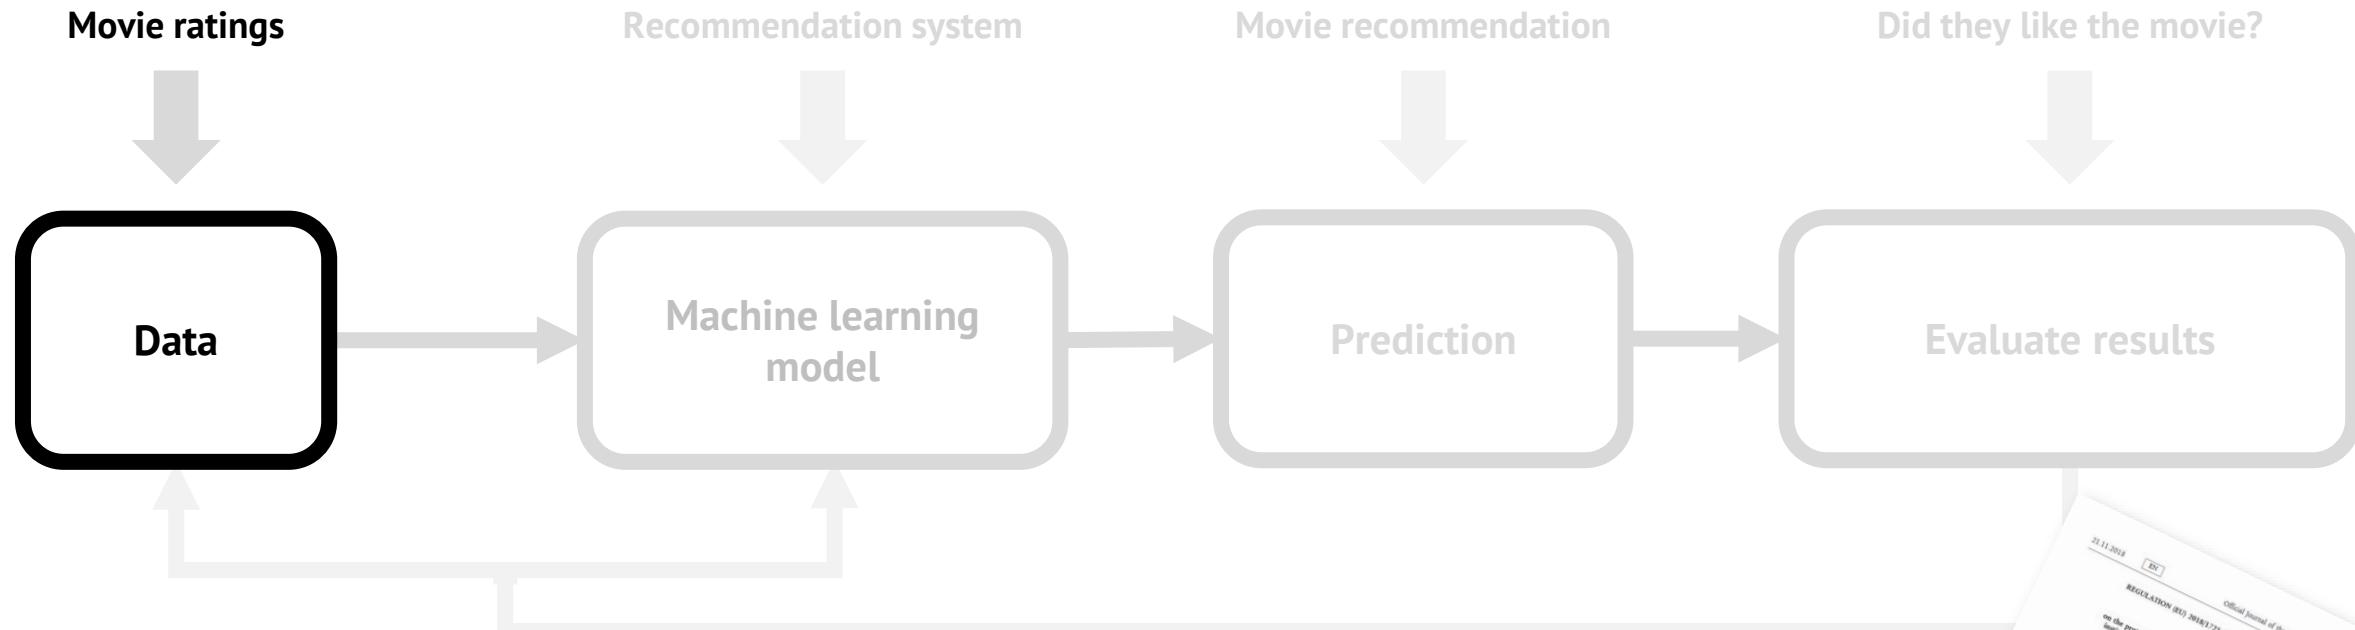


What data we collect shapes what type of recommendation we perform



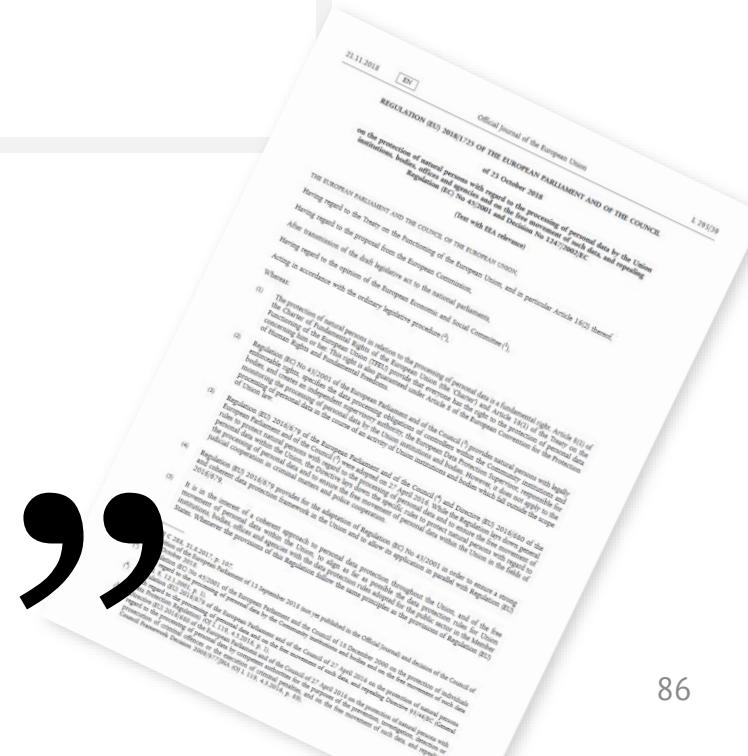
Explicit vs Implicit

What data we collect shapes what type of recommendation we perform

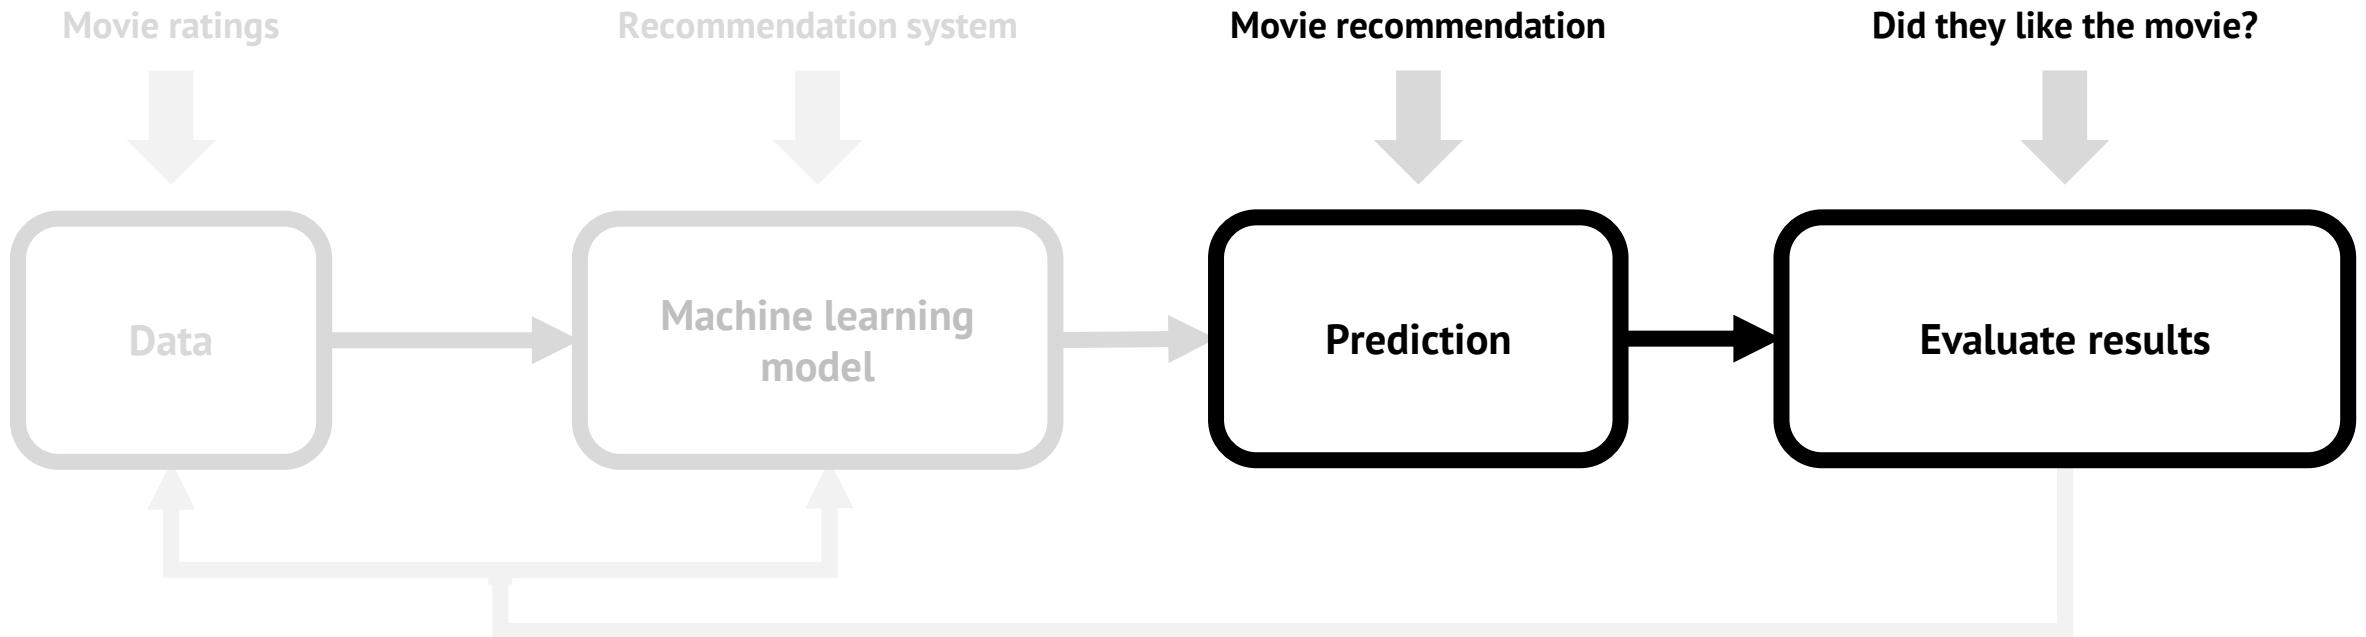


Article 4

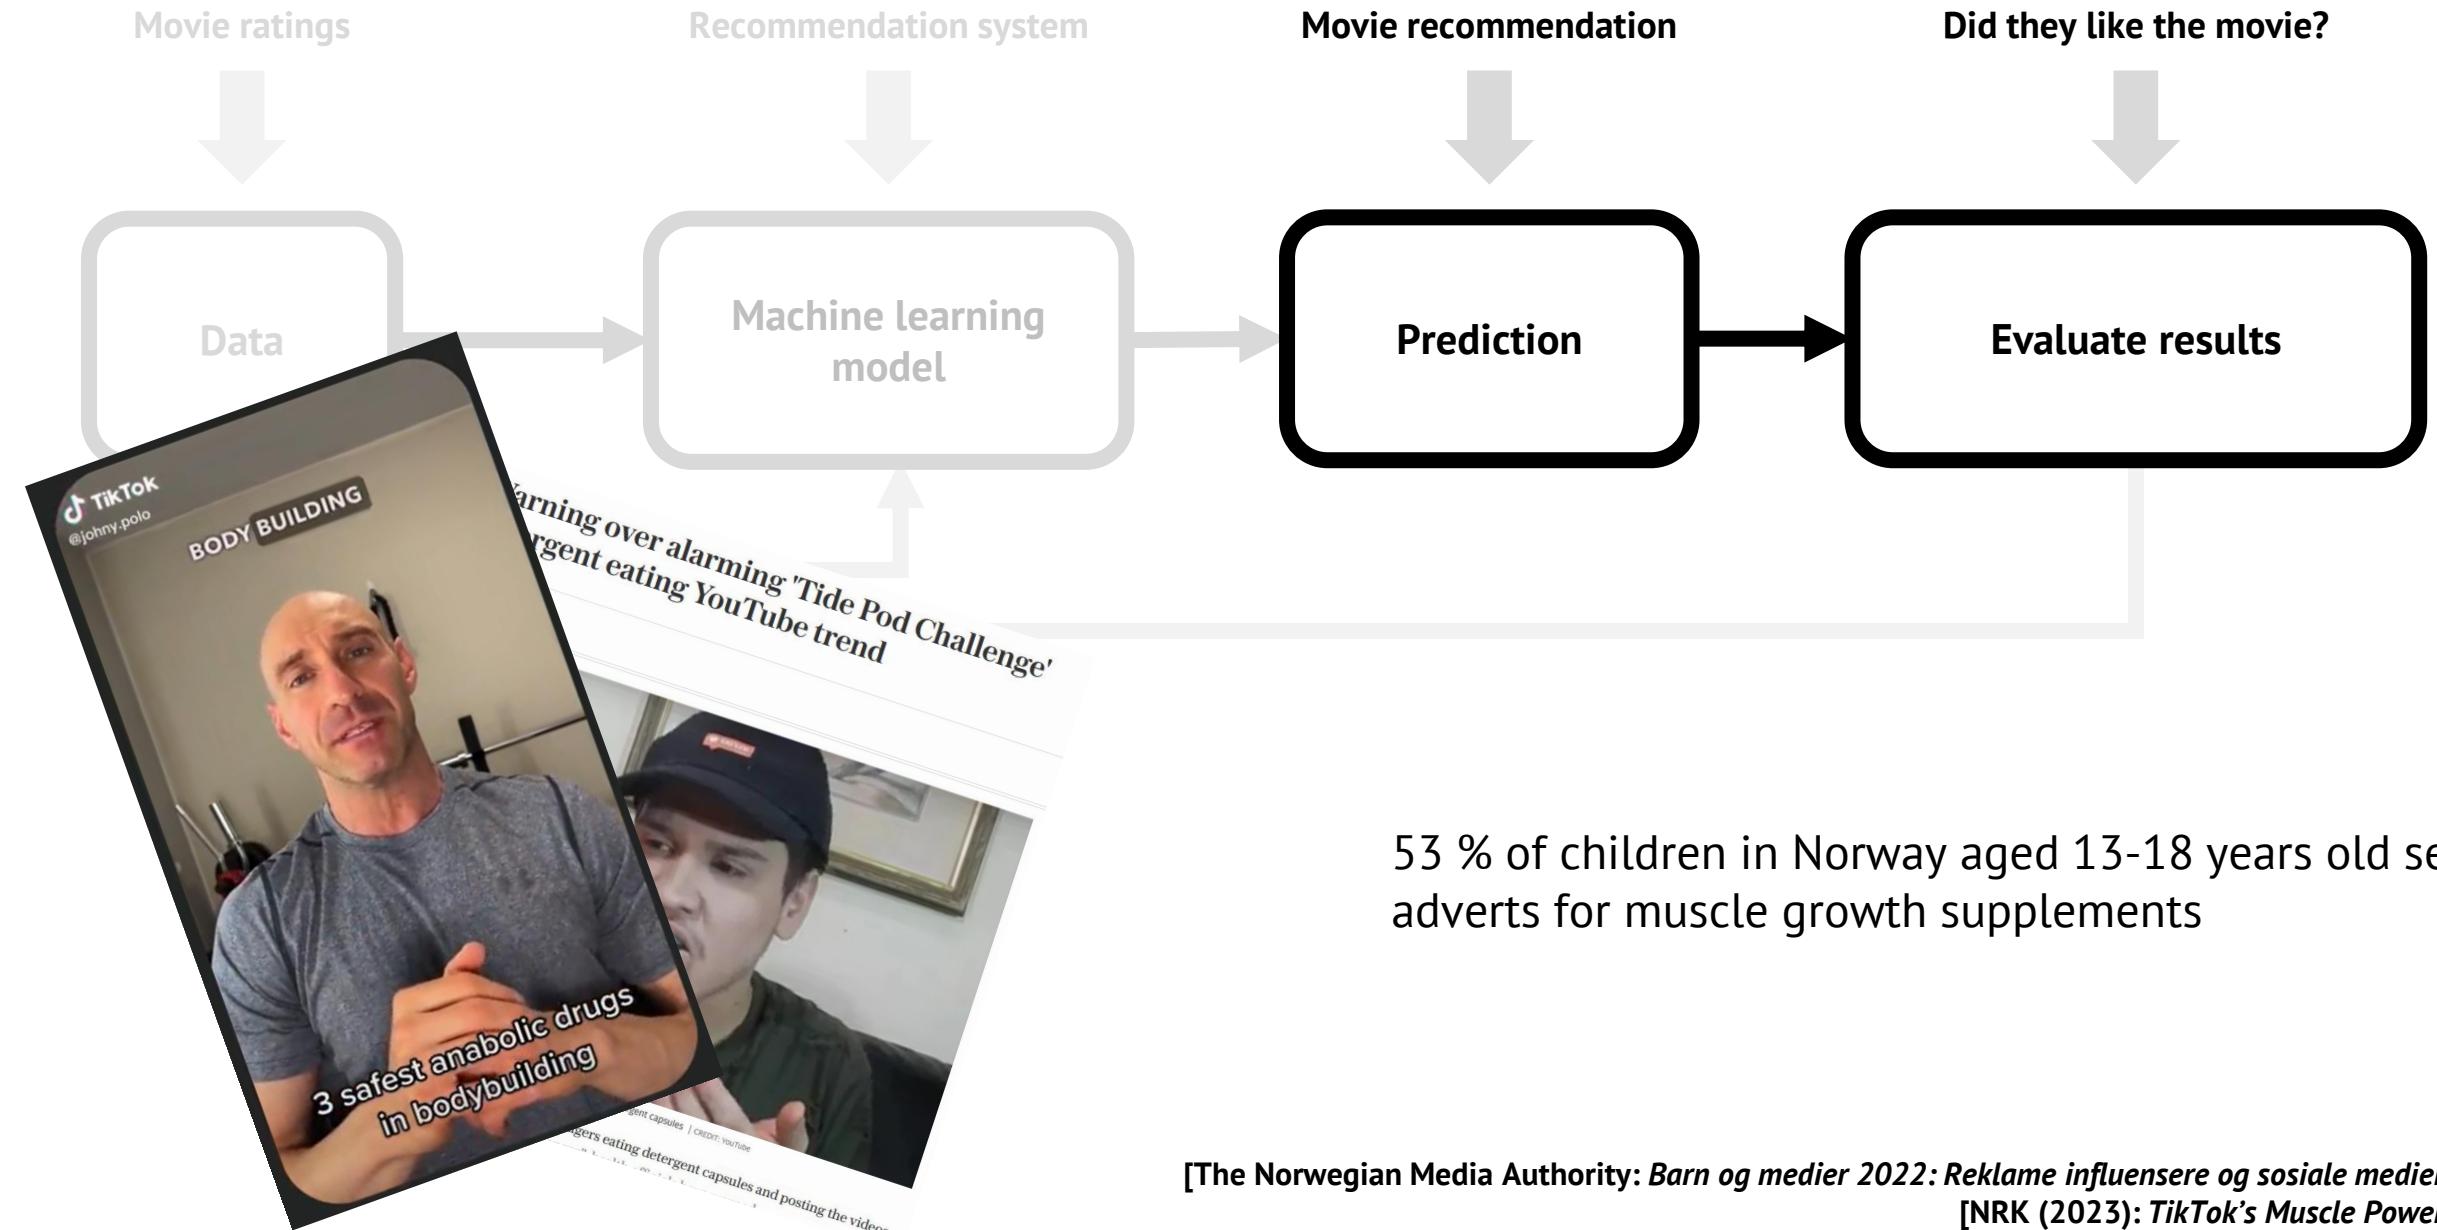
1. Personal data shall be: [...]
 - c) adequate, relevant and limited to what is necessary in relation to the purposes for which they are processed ('data minimisation');



How you evaluate your recommendations affects which model you use



How you evaluate your recommendations affects which model you use



Further reading about recommendation systems:

Websites and blog posts

- Cloudera (2021). *Session-based Recommender Systems*
 - <https://session-based-recommenders.fastforwardlabs.com/>
- Nvidia (2021). *Winning the SIGIR eCommerce Challenge on session-based recommendation with Transformers*
 - <https://medium.com/nvidia-merlin/winning-the-sigir-e-commerce-challenge-on-session-based-recommendation-with-transformers-v2-793f6fac2994>
- Spotify (2021). *Contextual and Sequential User Embeddings for Music Recommendation*
 - <https://research.at spotify.com/2021/04/contextual-and-sequential-user-embeddings-for-music-recommendation/>
- Shutterstock Data Science (2022). *Building Content Personalization*
 - <https://www.shutterstock.com/blog/data-science-building-content-personalization>

Papers

- Johnson, Douze & Jégou. IEEE Trans Big Data (2019). *Billion-Scale Similarity Search with GPUs*
- Vaswan et al. NeurIPS (2017). *Attention is all you need*
- Hansen et al. ACM RecSys (2020). *Contextual and Sequential User Embeddings for Large-Scale Music Recommendation*
- Moreira et al. ACM RecSys (2021). *Transformers4Rec: Bridging the Gap between NLP and Sequential / Session-Based Recommendation*
- Koren, Bell & Volinsky. IEEE Computer (2009). *Matrix Factorization Techniques for Recommender Systems*

Books

- Falk. Manning Publications (2019). *Practical Recommender Systems*

Software resources for making recommendation systems in Python

Software packages

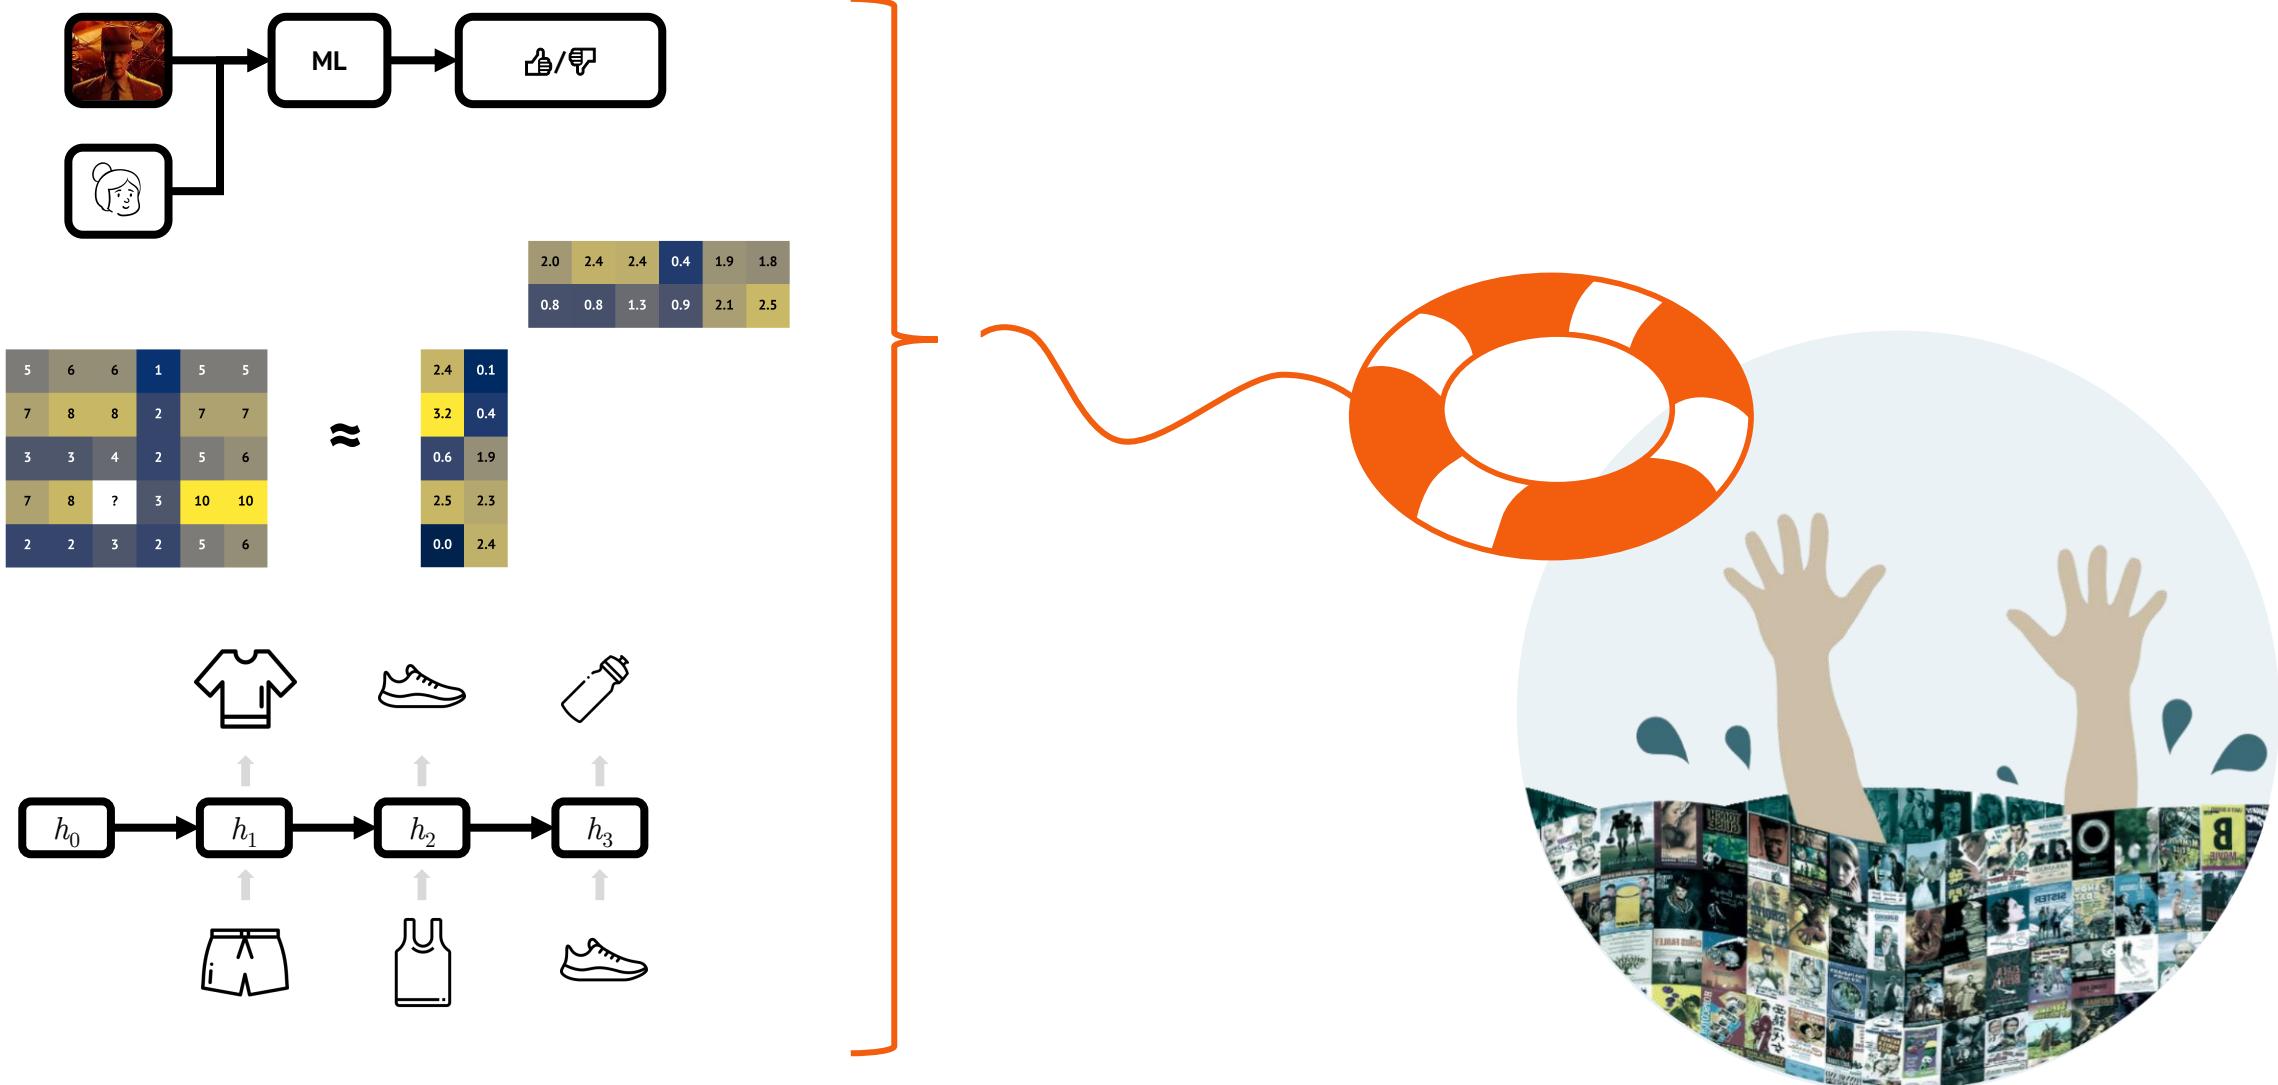
- **Surprise:** <https://surpriselib.com/>
 - Simple but effective methods for collaborative filtering:
 - Builtin support for model evaluation
- **Implicit:** <https://benfred.github.io/implicit/>
 - Collaborative filtering for implicit ratings
- **Gensim:** <https://radimrehurek.com/gensim/>
 - Word2vec
- **Transformers4Rec:** <https://nvidia-merlin.github.io/Transformers4Rec/main/>
 - Transformers for recommendation systems
- **FAISS:** <https://faiss.ai/>
 - Fast approximate KNN searches supporting several metrics

Code example

- **Example notebook:** <https://github.com/MarieRoald/phd-defence-trial-lecture>
 - Covering content-based, memory-based and model-based collaborative filtering with matrix factorisation and approximate nearest neighbour searches
- **Workshop from TMLS 2020:** <https://github.com/topspinj/tmls-2020-recommender-workshop>
 - More in-depth notebook showing explorative analysis and how to create various types of recommender systems

The screenshot shows a Jupyter Notebook interface with the title "Movie recommendation system". The code in the notebook includes imports for numpy, scipy, and pandas, along with various functions for data processing and matrix factorization. It also includes sections for "Download data", "Dummy encoding of the genres", and "Mapping from movie and user ID to row and column index". The code is written in Python and uses standard libraries for data science.

User features can be combined with item features to provide more personalized recommendations



Co-clustering utilise clusters of similar scores among both users and movies simultaneously

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	9	3	10	10
2	2	3	2	5	6

User-movie matrix



5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2

Overall mean

-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8

Co-cluster offset

Co-clustering utilise clusters of similar scores among both users and movies simultaneously

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	9	3	10	10
2	2	3	2	5	6

User-movie matrix



5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2

Overall mean

-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8

Co-cluster offset

4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8

User mean

4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6

User cluster mean

Co-clustering utilise clusters of similar scores among both users and movies simultaneously

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	9	3	10	10
2	2	3	2	5	6

User-movie matrix



5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2

Overall mean

-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8

Co-cluster offset

4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8

User mean

4.7	4.7	4.7	4.7	4.7	4.7
6.5	6.5	6.5	6.5	6.5	6.5
3.8	3.8	3.8	3.8	3.8	3.8
7.8	7.8	7.8	7.8	7.8	7.8
3.3	3.3	3.3	3.3	3.3	3.3

Movie mean

4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6

User cluster mean

5.6	5.6	5.6	5.6	5.6	5.6
5.6	5.6	5.6	5.6	5.6	5.6
5.0	5.0	5.0	5.0	5.0	5.0
5.0	5.0	5.0	5.0	5.0	5.0
5.0	5.0	5.0	5.0	5.0	5.0

Movie cluster mean

Co-clustering utilise clusters of similar scores among both users and movies simultaneously

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	9	3	10	10
2	2	3	2	5	6

User-movie matrix



5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2

Overall mean

-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8

Co-cluster offset

4.7	5.3	5.9	1.9	4.9	5.3
6.5	7.1	7.7	3.7	6.7	7.1
3.1	3.7	4.3	0.3	5.6	6.0
7.1	7.7	8.3	4.3	9.6	10.0
2.6	3.2	3.8	-0.2	5.1	5.5

Estimated ratings



4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8

User mean

4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6

User cluster mean



4.7	4.7	4.7	4.7	4.7	4.7
6.5	6.5	6.5	6.5	6.5	6.5
3.8	3.8	3.8	3.8	3.8	3.8
7.8	7.8	7.8	7.8	7.8	7.8
3.3	3.3	3.3	3.3	3.3	3.3

Movie mean

5.6	5.6	5.6	5.6	5.6	5.6
5.6	5.6	5.6	5.6	5.6	5.6
5.0	5.0	5.0	5.0	5.0	5.0
5.0	5.0	5.0	5.0	5.0	5.0
5.0	5.0	5.0	5.0	5.0	5.0

Movie cluster mean

Co-clustering utilise clusters of similar scores among both users and movies simultaneously

5	6	6	1	5	5
7	8	8	2	7	7
3	3	4	2	5	6
7	8	9	3	10	10
2	2	3	2	5	6

User-movie matrix



5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2
5.2	5.2	5.2	5.2	5.2	5.2

Overall mean

-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
-0.1	-0.1	-0.1	-0.1	-0.8	-0.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8
1.2	1.2	1.2	1.2	-1.8	-1.8

Co-cluster offset

4.7	5.3	5.9	1.9	4.9	5.3
6.5	7.1	7.7	3.7	6.7	7.1
3.1	3.7	4.3	0.3	5.6	6.0
7.1	7.7	8.3	4.3	9.6	10.0
2.6	3.2	3.8	-0.2	5.1	5.5

Estimated ratings



4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8
4.8	5.4	6.0	2.0	6.4	6.8

User mean



4.7	4.7	4.7	4.7	4.7	4.7
6.5	6.5	6.5	6.5	6.5	6.5
3.8	3.8	3.8	3.8	3.8	3.8
7.8	7.8	7.8	7.8	7.8	7.8
3.3	3.3	3.3	3.3	3.3	3.3

Movie mean

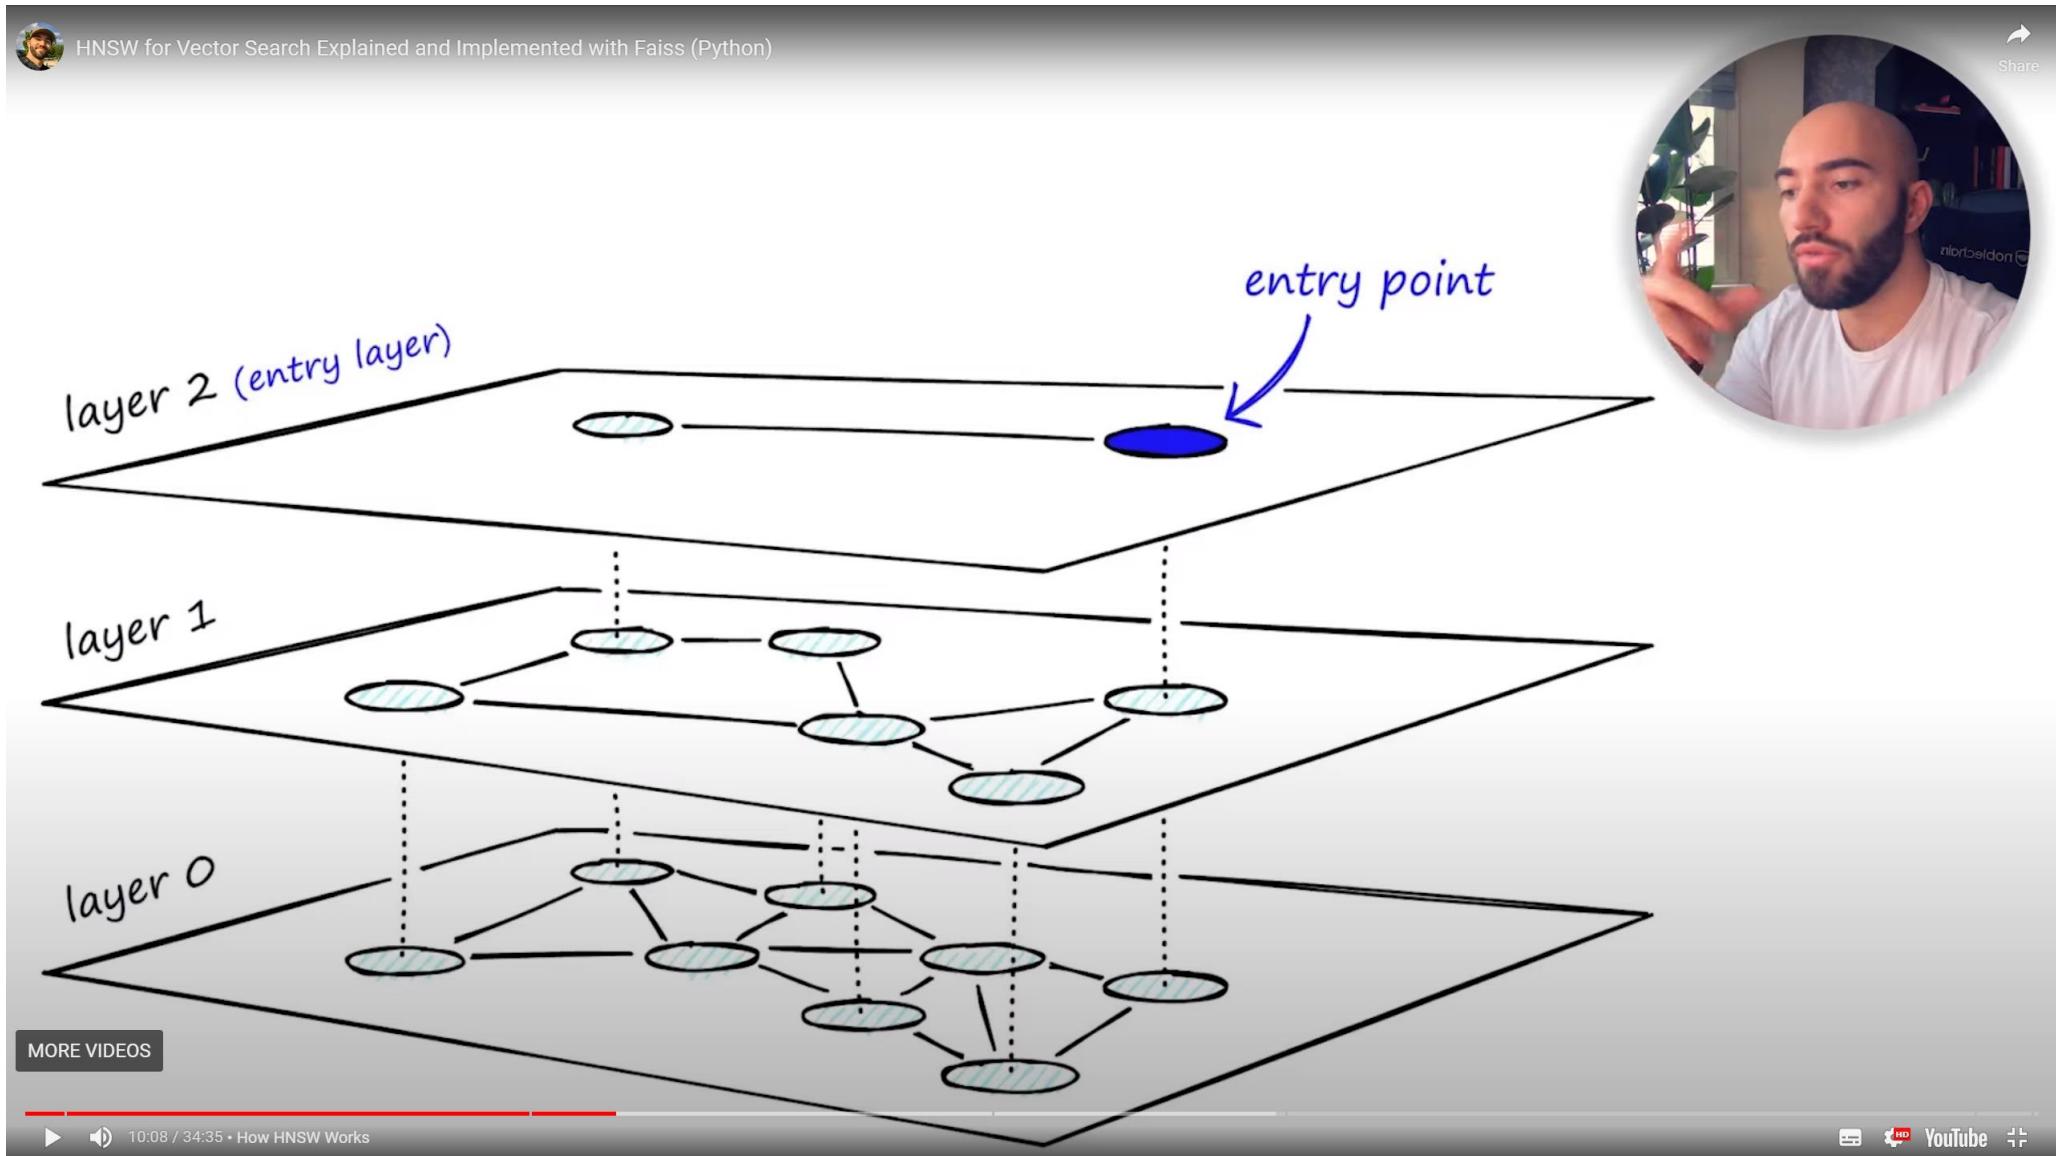
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6
4.5	4.5	4.5	4.5	6.6	6.6

User cluster mean

5.6	5.6	5.6	5.6	5.6	5.6
5.6	5.6	5.6	5.6	5.6	5.6
5.0	5.0	5.0	5.0	5.0	5.0
5.0	5.0	5.0	5.0	5.0	5.0
5.0	5.0	5.0	5.0	5.0	5.0

Movie cluster mean

Find clusters that minimise the difference between the estimated and true ratings





Booth A

242

104 (40 %)



Booth B

Walked by

260

Stopped to taste

145 (60 %)



Booth A

242

(40 %) 104

(30 %) 31



Booth B

260

145 (60 %)

4 (3 %)

Walked by

Stopped to taste

Bought jam